

HealthHub: A Healthcare Data Management System

Paolo Palumbo, Francesco Panattoni, Nedal Hadam

June 19, 2025

Contents

1	Introduction	5
2	Dataset and Web Scraping	7
2.1	Web Scraping	7
2.2	Synthetic Data Generation	7
2.3	Dataset Characteristics: Velocity and Variety	8
2.3.1	Velocity	8
2.3.2	Variety	8
2.4	Resulting Dataset	8
3	Design	11
3.1	Actors	11
3.2	Functional Requirements	11
3.2.1	Guest Users	11
3.2.2	Patients	12
3.2.3	Doctors	12
3.3	Non-Functional Requirements	12
3.3.1	Performance and Scalability	12
3.3.2	Availability and Reliability	13
3.3.3	Security and Privacy	13
3.3.4	Usability	13
3.3.5	Portability and Flexibility	13
3.3.6	Maintainability	13
3.4	UML Class Diagrams	14
3.4.1	Relationships between classes	14
3.4.2	Design of the classes	15
4	Data Model	19
4.1	Document Database - MongoDB	19
4.1.1	Collections	19
4.1.2	Document Structure	22
4.2	Graph Database - Neo4j	25
4.2.1	Node Types	26
4.2.2	Relationships	26

5	Implementation	27
5.1	Front-End	27
5.2	Front-End Structure	27
5.2.1	Static	27
5.2.2	Templates	28
5.3	Back-End	28
5.3.1	Spring Boot	28
5.4	Back-End Structure	29
5.4.1	package it.unipi.healthhub.config	29
5.4.2	package it.unipi.healthhub.controller	29
5.4.3	package it.unipi.healthhub.controller.api	30
5.4.4	package it.unipi.healthhub.dto	36
5.4.5	package it.unipi.healthhub.exception	36
5.4.6	package it.unipi.healthhub.filter	37
5.4.7	package it.unipi.healthhub.model	37
5.4.8	package it.unipi.healthhub.model.mongo	38
5.4.9	package it.unipi.healthhub.model.neo4j	51
5.4.10	package it.unipi.healthhub.projection	53
5.4.11	package it.unipi.healthhub.repository	53
5.4.12	package it.unipi.healthhub.service	53
5.4.13	package it.unipi.healthhub.util	55
5.4.14	Scheduler Component	56
6	Queries	57
6.1	Mongo DB	57
6.1.1	Search Query	57
6.1.2	Analytics	59

Chapter 1

Introduction

Healthhub is a web-based platform designed to simplify and centralize the management of medical appointments. It offers dedicated features for both patients and healthcare professionals. Regular users can search for doctors, book visits, and leave reviews, while medical professionals can manage the services they offer and track their appointments.

- Allow users to search for medical professionals and retrieve information on their services
- Allow users to book appointments and review doctors after a visit
- Allow users to manage their appointment history and cancel or modify bookings
- Give users recommendations on doctors based on their location and past activity
- Allow doctors to specify their specializations, availability, and offered services
- Allow doctors to view and manage their upcoming appointments

Chapter 2

Dataset and Web Scraping

To populate the dataset for our application, we used a combination of web scraping, synthetic data generation, and logical inference based on real-world patterns of user behavior.

2.1 Web Scraping

The initial data was scraped from a platform containing medical reviews, MioDottore¹. This platform provides a comprehensive database of medical professionals, including their specializations, locations, and user reviews. The scraping process involved extracting the following information:

- Doctor names and specializations
- Locations (cities and regions)
- User reviews and ratings
- Contact information (where available)
- Service offerings (e.g., types of consultations, treatments)

2.2 Synthetic Data Generation

To enhance the dataset, we generated synthetic data to simulate a more realistic user base and appointment history. This involved several steps:

- **Unique Users:** We extracted all unique usernames from the scraped reviews. For each user, we generated a detailed profile including demographic and geographical information.
- **Appointments:** For every review, we created a corresponding medical appointment. These were dated in the weeks preceding the review, to simulate a realistic flow where users leave feedback shortly after being treated.

¹<https://www.miodottore.it/>

- **Likes:** To simulate engagement features such as likes on reviews, we analyzed the provinces of the doctors each user had interacted with. We then randomly selected other doctors in the same provinces and associated a number of likes comparable to the number of reviews written by each user.

2.3 Dataset Characteristics: Velocity and Variety

2.3.1 Velocity

The application context is inherently dynamic, and our dataset reflects this through the following characteristics:

- **High review frequency:** The dataset reflects an average of at least 100 new reviews per day, simulating the continuous flow of patient feedback that a live system would experience.
- **Growing doctor base:** Based on trends from medical boards, we estimate around 400 to 500 new doctors would register on the platform annually, contributing to the system's dynamism and evolving content.

2.3.2 Variety

The dataset incorporates diverse data types, which contribute to its heterogeneity:

- **Multi-format records:** Structured user and doctor profiles, unstructured textual reviews, and timestamped interaction logs (likes, appointments) are all represented.
- **Diverse entities:** The inclusion of users, doctors, reviews, likes, and appointments enables multi-relational analysis and feature richness.

2.4 Resulting Dataset

The original scraped dataset, stored in `scraped.json` (265 MB), contains the raw information collected from the web. This dataset served as the foundation for the data generation process, providing:

- 699,987 reviews
- 214,682 unique reviewers
- 87,632 doctors

Building upon this, the final dataset is stored in JSON format, totaling approximately 960 MB, and includes both real and synthetic data entries. Its main components are:

- **Doctors** (`doctors.json`, 289 MB) Contains 87,632 healthcare professionals, each with profile data and linked to 699,987 reviews. Reviews include timestamps and patient feedback.
- **Users** (`users.json`, 66 MB) Comprises 214,682 unique users extracted from the original dataset and extended with synthetic profiles.
- **Appointments** (`appointments.json`, 422 MB) Each review is connected to a synthetic appointment, scheduled in the weeks preceding the review date. Appointments simulate realistic scheduling and clinic visit patterns.
- **Templates** (`templates.json`, 162 MB) Stores template structures for appointment scheduling logic, such as available time slots, weekdays, and timing constraints.
- **User Likes** (`user_likes.json`, 22 MB) Represents user interactions with doctors, generated according to the geographic distribution of the doctors reviewed by each user.

Overall, this dataset offers a realistic simulation of user and doctor activity on a healthcare review platform, reflecting both volume and behavioral complexity.

Chapter 3

Design

3.1 Actors

The application involves three primary categories of actors, each with distinct roles and permissions:

- *Non-Authenticated User (Guest User)*: refers to an anonymous individual who accesses the application without logging in. This actor is allowed to register, authenticate using existing credentials, or explore the platform by searching for doctors and viewing their public profiles.
- *Patient (User)*: represents the end-user of the service. Once authenticated, the patient can book appointments with doctors, confirm or cancel them, and subsequently provide feedback through reviews.
- *Doctor*: a professional figure who offers medical appointments. The doctor can manage their availability schedule, making it accessible to patients, and consult an analytics dashboard displaying aggregated data such as earnings, reviews, and the number of visits performed.

3.2 Functional Requirements

The following section outlines the functional requirements.

3.2.1 Guest Users

Guest users are allowed to:

- Register for an account within the application;
- Log in using their credentials;
- Initiate the password recovery process in case of forgotten credentials;
- Search for doctors and access their public profiles.

3.2.2 Patients

Authenticated users identified as **patients** are granted the following capabilities:

- Manage and update their personal profiles;
- Search for doctors and access their public profiles;
- Book appointments and optionally include a note for the doctor;
- Endorse a doctor as a form of support or recommendation;
- Submit reviews regarding their medical experience;
- View their appointments, including past, current, and upcoming ones;
- Search for doctors through the recommendation feature provided by the system.

3.2.3 Doctors

Authenticated users identified as **doctors** are provided with the following functionalities:

- Manage and update their personal and professional profiles;
- Search for other doctors and access their public profiles;
- Define and configure the types of visits offered, along with the associated pricing;
- Oversee appointments by identifying scheduled patients and appointments for the current day;
- Configure and manage appointment availability through customizable scheduling templates;
- Monitor and respond to patient reviews;
- Access analytics, including visual representations of revenue, newly acquired patients, recent reviews, and a summary of completed visits.

3.3 Non-Functional Requirements

The following non-functional requirements address the quality attributes of the system, ensuring its performance, reliability, security, scalability, and maintainability.

3.3.1 Performance and Scalability

- The system shall maintain **acceptable response times** for common operations such as viewing medical records and scheduling appointments;
- The system shall efficiently **handle increased workload** during peak usage periods without degradation of service.

3.3.2 Availability and Reliability

- The application shall be **available** to all users **24/7**, minimizing downtime through redundancy and failover mechanisms;
- The system shall implement **backup and recovery procedures** to preserve data integrity and support rapid restoration after failures.
- The system shall tolerate occasional data staleness in non-critical views, ensuring high availability even under degraded conditions.

3.3.3 Security and Privacy

- The system shall enforce **secure, authenticated access for all users**, employing strong password policies and session management;
- All sensitive data shall be encrypted both in transit and at rest;
- The system provides defenses against injections.

3.3.4 Usability

- The user interface shall be intuitive and **user-friendly**, enabling users to perform tasks with minimal learning curve.
- The application shall exhibit **low latency** in user interactions to maintain a responsive experience.

3.3.5 Portability and Flexibility

- The application shall be **deployable on multiple operating systems** (e.g., Windows, macOS, Linux) without requiring behavioral changes.
- The system's architecture shall support the addition of new attributes or modules (e.g., new appointment types) with minimal code modification.

3.3.6 Maintainability

- The codebase shall follow **Object-Oriented design** principles, promoting modularity and ease of comprehension.
- The system shall **minimize single points of failure** through component decoupling and redundancy.
- The application shall include comprehensive documentation and code comments to facilitate future enhancements and debugging.

3.4 UML Class Diagrams

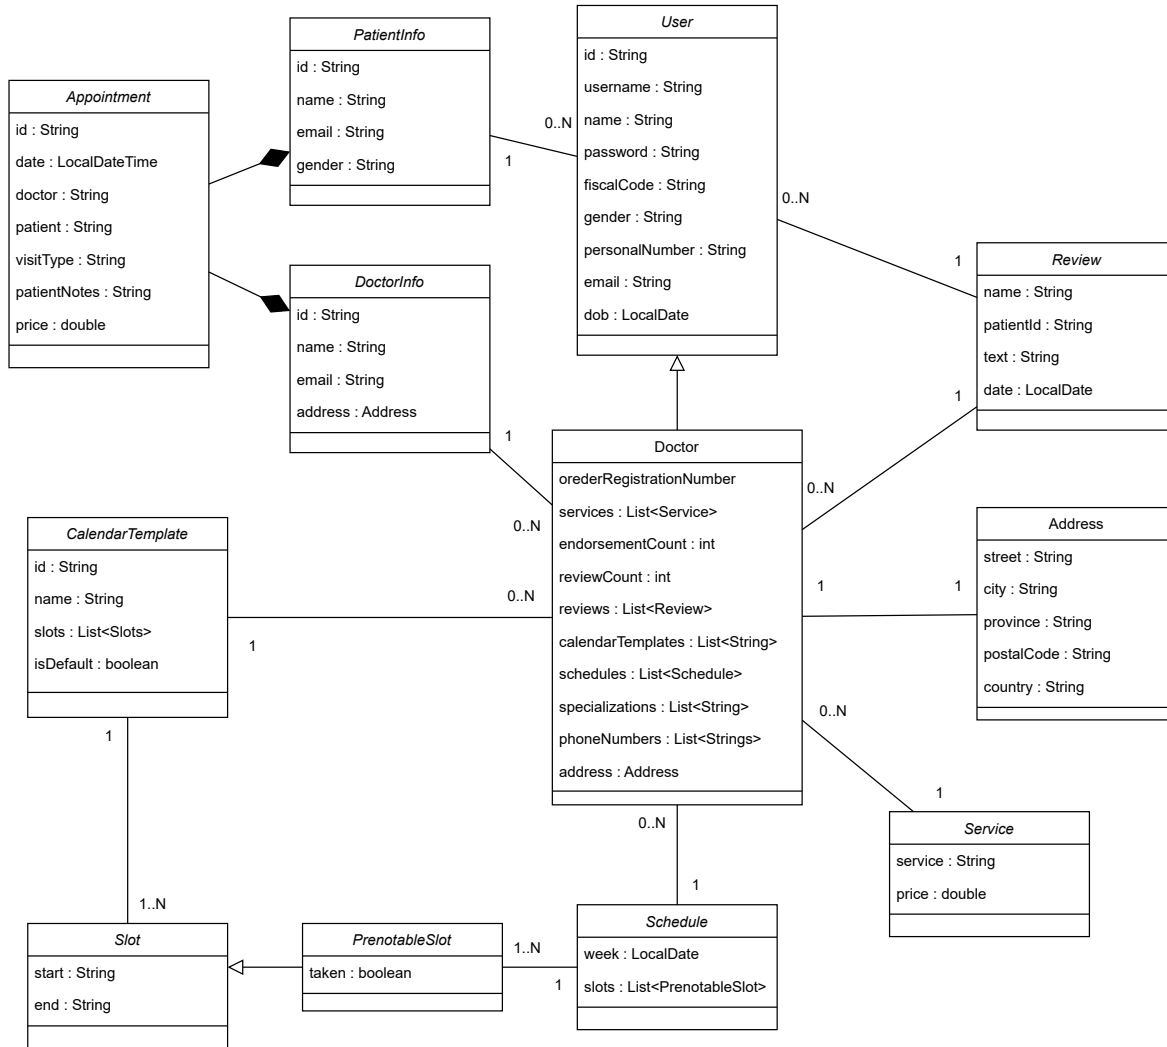


Figure 3.1: GraphDB structure for HealthHub

3.4.1 Relationships between classes

Here we briefly summarise the relation between each class.

- **Doctor** is a specialization of **User**;
- A **Doctor** may offer zero or more **Service** instances and each **Service** is provided by exactly one **Doctor**;
- A **Doctor** may have zero or more **Review** entries and each **Review** refers to exactly one **Doctor**;

- A **Doctor** may define zero or more **CalendarTemplates** and each **CalendarTemplate** belongs to exactly one **Doctor**;
- A **Doctor** may publish zero or more **Schedules** and each **Schedule** is associated with exactly one **Doctor**;
- Each **Schedule** comprises zero or more **PrenotableSlots** and each **PrenotableSlot** is part of exactly one **Schedule**;
- Each **CalendarTemplate** comprises zero or more **Slots** and each **Slot** belongs to exactly one **CalendarTemplate**;
- A **Doctor** has exactly one **Address** and each **Address** is linked to exactly one **Doctor** (and similarly to **DoctorInfo**);
- An **Appointment** is associated with exactly one **DoctorInfo** and one **PatientInfo** and each **DoctorInfo** or **PatientInfo** may have zero or more **Appointments**.

3.4.2 Design of the classes

Here we briefly summarise the primary attributes of each entity.

User

- `id: String` – unique identifier
- `username: String`
- `name: String`
- `password: String` – stored as hash
- `fiscalCode: String`
- `gender: String`
- `personalNumber: String`
- `email: String`
- `dob: LocalDate`

Doctor (specialization of User)

- `orderRegistrationNumber: String`
- `services: List<Service>`
- `endorsementCount: int`
- `reviewCount: int`
- `reviews: List<Review>`
- `calendarTemplates: List<CalendarTemplate>`
- `schedules: List<Schedule>`

- specializations: List<String>
- phoneNumbers: List<String>
- address: Address

Address

- street: String
- city: String
- province: String
- postalCode: String
- country: String

Appointment

- id: String
- date: LocalDateTime
- visitType: String
- patientNotes: String
- price: double

Schedule

- week: LocalDate
- slots: List<PrenotableSlot>

CalendarTemplate

- id: String
- name: String
- slots: List<Slot>
- isDefault: boolean

Slot

- start: String
- end: String

PrenotableSlot (subclass of Slot)

- taken: boolean

Review

- name: String

- `patientId: String`
- `text: String`
- `date: LocalDate`

Service

- `service: String`
- `price: double`

PatientInfo (belong to **Appointment**)

- `id: String`
- `name: String`
- `email: String`
- `gender: String`

DoctorInfo (belong to **Appointment**)

- `id: String`
- `name: String`
- `email: String`
- `address: Address`

Chapter 4

Data Model

4.1 Document Database - MongoDB

The application utilizes MongoDB, a document-oriented NoSQL database, to store and manage data.

4.1.1 Collections

Users

Each user in the system is a registered patient with access to the platform's features such as booking appointments and interacting with doctors. Below are the typical attributes found in a user's document:

- `_id` : unique identifier of the user, stored as an `ObjectId`.
- `fiscalCode` : string representing the Italian *codice fiscale*, used as a unique tax and identity code.
- `name` : string containing the full name of the user.
- `dob` : date indicating the user's date of birth, stored in `ISODate` format.
- `gender` : string that specifies the gender of the user (e.g., "male", "female").
- `personalNumber` : string that holds the user's personal phone number.
- `email` : string containing the email address provided during registration.
- `username` : unique string chosen by the user during sign up and used to log in.
- `password` : string containing the hashed version of the password set at registration.

Doctors

Each doctor in the system is a registered healthcare professional who can provide services, receive reviews from patients, and manage their availability through calendar templates. Below are the typical attributes found in a doctor's document:

- **_id** : unique identifier of the doctor, stored as an `ObjectId`.
- **name** : string containing the full name of the doctor.
- **email** : string with the professional email address used for communication.
- **username** : unique string used for login and identification within the platform.
- **password** : string containing the hashed version of the doctor's password.
- **address** : embedded object containing **street**, **city**, **province**, **country**, and **postalCode**.
- **phoneNumbers** : array of strings listing the phone numbers associated with the doctor.
- **specializations** : array of strings indicating the medical fields in which the doctor is specialized.
- **services** : array of embedded documents, each with a **service** name and a corresponding **price**.
- **endorsementCount** : integer representing the number of professional endorsements received.
- **reviews** : array of embedded documents, each containing a **patientId**, **name**, **text** of the review, and the **date** of submission.
- **reviewsCount** : integer indicating the total number of reviews received by the doctor.
- **dob** : date representing the doctor's date of birth, stored in `ISODate` format.
- **fiscal_code** : string corresponding to the Italian fiscal code (`codice fiscale`).
- **orderRegistrationNumber** : string representing the official registration number in the professional healthcare registry.
- **calendarTemplates** : array of `ObjectIds` referencing availability templates used for scheduling appointments.

Appointments

The **Appointments** collection stores all scheduled medical visits between patients and doctors. Each document represents a single appointment and includes contextual details about the involved parties and the visit itself.

- **_id** : unique identifier of the appointment, stored as an **ObjectId**.
- **date** : **ISODate** indicating the exact date and time of the scheduled appointment.
- **doctor** : embedded object including:
 - **_id** : reference to the doctor's unique identifier.
 - **name** : full name of the doctor.
 - **address** : object containing **street**, **city**, **province**, **country**, and **postalCode**.
 - **email** : contact email address of the doctor.
- **patient** : embedded object including:
 - **_id** : reference to the patient's unique identifier.
 - **name** : full name of the patient.
 - **fiscalCode** : Italian fiscal code (codice fiscale) of the patient.
 - **email** : contact email address of the patient.
 - **gender** : gender of the patient.
- **visitType** : string describing the type of visit or consultation.
- **patientNotes** : string field where patients may leave notes prior to the appointment (optional).
- **price** : numeric value indicating the cost of the visit in euros.

CalendarTemplates

The **CalendarTemplates** collection stores predefined weekly availability templates that doctors can use to generate their actual working calendars. Each template specifies on which days and time intervals the doctor is available for appointments.

- **_id** : unique identifier of the calendar template, stored as an **ObjectId**.
- **name** : name of the template (e.g., "Standard").
- **slots** : an object mapping weekdays to arrays of available time intervals. Each interval is represented as an object with:
 - **start** : start time of the slot (format: **HH:MM**).
 - **end** : end time of the slot (format: **HH:MM**).
- **isDefault** : boolean indicating whether the template is the doctor's default schedule.

4.1.2 Document Structure

Listing 4.1: Example of a User Document

```

1 {
2   _id: ObjectId('684ada4637804916ca651761'),
3   fiscalCode: 'UGOSIG070805MA41',
4   name: 'Sig. Ugolino Franscini',
5   password: '5e884898da28047151d0e56f8dc6292773603d0d6aabbdd6
      2a11ef721d1542d8',
6   dob: ISODate('2007-08-05T00:00:00.000Z'),
7   gender: 'male',
8   personalNumber: '+39 3483861896',
9   email: 'iruberto@virgilio.it',
10  username: 'Iannelli'
11 }

```

Listing 4.2: Example of a Doctor Document

```

1 {
2   _id: ObjectId('684adad437804916ca65ba64'),
3   name: 'Barbara Montagnini',
4   email: 'barbara.montagnini@libero.it',
5   username: 'barbara_montagnini',
6   password: '5e884898da28047151d0e56f8dc6292773603d0d6aabbdd6
      2a11ef721d1542d8',
7   address: {
8     street: 'Via Mura di San Teonisto 8,',
9     city: 'Treviso',
10    province: 'TV',
11    country: 'IT',
12    postalCode: '31100'
13  },
14  phoneNumbers: [ '338 439 2066' ],
15  specializations: [ 'Psicoterapia', 'Psicologia Clinica', '
      Psicologia del Lavoro' ],
16  services: [
17    { service: 'Colloquio psicologico', price: 70 },
18    { service: 'Psicoterapia', price: 75 }
19  ],
20  endorsementCount: 24,
21  reviews: [
22    {
23      patientId: ObjectId('684ada4537804916ca644b76'),
24      name: 'Agostino Omma-Depero',

```

```

25         text: 'La Dott.ssa mi ha fatto sentire molto a mio
26             agio.\n' +
27             'E mi ha spiegato in dettaglio le dinamiche da
28             cui derivava il mio problema',
29         date: ISODate('2025-05-25T15:36:33.424Z')
30     },
31     {
32         patientId: ObjectId('684ada4537804916ca63fc97'),
33         name: 'Enzio Camuccini',
34         text: 'Mi sono rivolto per un problema di insonnia.
35             Assieme abbiamo cercato di capire l'origine e
36             mi ha fornito strategie pratiche per migliorare
37             la qualita del sonno senza farmaci.',
38         date: ISODate('2025-04-20T15:26:33.424Z')
39     },
40     {
41         patientId: ObjectId('684ada4437804916ca628d94'),
42         name: 'Diana Pacelli',
43         text: 'Ho contattato la d.ssa per stress lavorativo
44             . Mi e' servito a riflettere sulle mie emozioni,
45             in particolare sulla rabbia che non mi
46             permetteva di essere sempre obiettivo. Ero
47             arrivata con un senso di rivalsa verso la mia
48             azienda sono uscita con strumenti utili per la
49             mia vita non solo professionale. SUpEr
50             conisgliata',
51         date: ISODate('2025-04-16T14:40:45.424Z')
52     },
53     {
54         patientId: ObjectId('684ada4537804916ca646840'),
55         name: 'Giuliana Pertini',
56         text: "Completamente inaffidabile per quanto
57             riguarda appuntamenti e orari. Cambia
58             continuamente ora e giorno dell'appuntamento con
59             scuse varie...",
60         date: ISODate('2025-01-31T07:40:41.424Z')
61     }
62 ],
63 reviewsCount: 4,
64 dob: ISODate('2003-03-27T00:00:00.000Z'),
65 fiscal_code: 'MONBAR030327MKOW',
66 orderRegistrationNumber: 'TV-314952',
67 calendarTemplates: [ ObjectId('684ad9f637804916ca611d7a') ]
68 }

```

Listing 4.3: Example of an Appointment Document

```

1 {
2   _id: ObjectId('684adc6837804916ca700ad8'),
3   date: ISODate('2024-03-11T13:52:13.000Z'),
4   doctor: {
5     _id: ObjectId('684adad437804916ca662b8a'),
6     name: "Serena D'Agostino",
7     address: {
8       street: 'Viale delle Querce 8,',
9       city: 'Castrovillari',
10      province: 'CS',
11      country: 'IT',
12      postalCode: '87012'
13    },
14    email: "serena.d'agostino@yahoo.com"
15  },
16  patient: {
17    _id: ObjectId('684ada4437804916ca6287e5'),
18    name: 'Rocco Zacco',
19    fiscalCode: 'ZACROC100226M1LH',
20    email: 'chigifabrizia@poste.it',
21    gender: 'male'
22  },
23  visitType: 'Visita ginecologica',
24  patientNotes: '',
25  price: 120
26 }

```

Listing 4.4: Example of a Calendar Template Document

```

1 {
2   _id: ObjectId('684ad9f637804916ca620075'),
3   name: 'Standard',
4   slots: {
5     monday: [
6       { start: '08:30', end: '09:00' },
7       { start: '09:00', end: '09:30' },
8       { start: '09:30', end: '10:00' },
9       { start: '10:00', end: '10:30' },
10      { start: '10:30', end: '11:00' },
11      { start: '11:00', end: '11:30' },
12      { start: '11:30', end: '12:00' },
13      { start: '12:00', end: '12:30' }
14    ],
15    wednesday: [

```

```

16         { start: '14:30', end: '15:00' },
17         { start: '15:00', end: '15:30' },
18         { start: '15:30', end: '16:00' },
19         { start: '16:00', end: '16:30' },
20         { start: '16:30', end: '17:00' },
21         { start: '17:00', end: '17:30' },
22         { start: '17:30', end: '18:00' },
23         { start: '18:00', end: '18:30' }
24     ],
25     friday: [
26         { start: '10:00', end: '10:30' },
27         { start: '10:30', end: '11:00' },
28         { start: '11:00', end: '11:30' },
29         { start: '11:30', end: '12:00' },
30         { start: '16:00', end: '16:30' },
31         { start: '16:30', end: '17:00' },
32         { start: '17:00', end: '17:30' },
33         { start: '17:30', end: '18:00' }
34     ]
35 },
36 isDefault: true
37 }

```

4.2 Graph Database - Neo4j

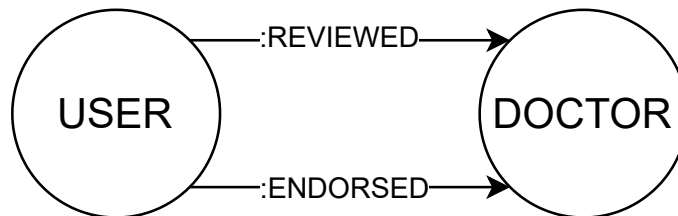


Figure 4.1: GraphDB structure for HealthHub

Our graph database (Figure 4.1) models two main entities:

- User
- Doctor

Each node retains the MongoDB identifier to maintain consistency across databases. Neo4j is used selectively, focusing on relationship patterns and core attributes needed for traversal and inference, while full entity details remain in MongoDB.

The graph database plays a central role in modeling the social structure of the platform and is exploited in two main ways:

1. As a **recommendation engine**, suggesting doctors to users based on graph proximity and shared review patterns.
2. As a **search optimization layer**, improving query efficiency by leveraging graph traversal capabilities.

Due to the lack of geographic data, we approximate location by leveraging graph proximity, under the assumption that users typically visit doctors near them. This enables us to infer locality-aware clusters and improve personalization without explicit spatial attributes.

4.2.1 Node Types

Neo4j stores only a minimal subset of each entity:

- **User** nodes: `name`, `_id`.
- **Doctor** nodes: `name`, `specializations`, `_id`.

This lean representation minimizes redundancy while supporting all graph-based operations.

4.2.2 Relationships

We define two types of directed edges:

- **:REVIEWED** — if user U has reviewed doctor D , then a **:REVIEWED** relationship is created from U to D :

$$U \xrightarrow{\text{:REVIEWED}} D$$

- **:ENDORSED** — if user U endorsed doctor D (e.g., marked them as a favorite or recommended), we store this with a **:ENDORSED** relationship:

$$U \xrightarrow{\text{:ENDORSED}} D$$

These relations enable pattern discovery and inference. Spatial behavior is indirectly modeled through interaction patterns, which often reflect geographic closeness.

Chapter 5

Implementation

5.1 Front-End

The front-end of our application has been developed as a traditional web-based interface, utilizing a combination of **JavaScript**, **HTML** and **CSS**. This technology stack enables the creation of a responsive and interactive user interface that can be accessed from any modern web browser without the need for additional installations. The application communicates asynchronously with the back-end server through **AJAX** (Asynchronous JavaScript and XML) requests, allowing data to be fetched and updated dynamically without requiring full page reloads. This results in a smoother user experience and improved performance, particularly when handling operations such as appointment booking, profile management and real-time data visualization. The separation of concerns between presentation and logic also enhances the maintainability and scalability of the application.

5.2 Front-End Structure

The **resources** directory encapsulates the front-end assets and configuration files essential to the operation and presentation of the web application. Its hierarchical organization reflects a separation of concerns, promoting maintainability and modular development.

Static resources such as CSS and JavaScript files reside in the **static** folder, structured into subdirectories for stylesheets (**css/**), scripts (**js/**) and error pages (**error/**). These support the visual and interactive components of the application.

Dynamic **HTML** views based on **Thymeleaf** are located in the **templates** folder, with dedicated files for both doctors and users. A **fragments** subfolder contains reusable layout components (e.g., headers, footers, sidebars) to ensure consistency and reusability. Application-level configuration is managed through **application.properties**.

5.2.1 Static

The **static** directory collects all static web assets that are delivered to the client without server-side processing. These assets are organized into three primary subfolders: **css/**, **js/**, and **error/**, each serving a distinct role in the client-side rendering process.

The **css** folder contains a series of modular stylesheets dedicated to specific views or components. These include both general-purpose files (e.g. `base.css`, `main.css`, `form.css`) and context-specific styles (e.g. `doctor-dashboard.css`, `user-profile.css`), which define the visual identity and layout of user and doctor interfaces.

The **js** folder hosts JavaScript files that manage client-side interactions and dynamic behavior. Each script is associated with a specific view or functionality—for instance, `login.js` handles authentication interactions, while `slot-calendar.js` governs dynamic calendar updates for appointment scheduling.

The **error** folder includes static HTML pages for standard HTTP error responses (400, 404, 500), ensuring that users receive consistent and informative feedback in case of failure scenarios.

This structure promotes modularity and clarity, allowing for straightforward maintenance and scalability of front-end assets.

5.2.2 Templates

The **templates** directory constitutes a fundamental component within the presentation layer of the web application, as it houses the markup resources responsible for rendering the user interface. This directory is hierarchically organized into subfolders, each reflecting a specific functional domain or interaction paradigm with which the user engages. The HTML documents contained therein are primarily structured using Thymeleaf, a server-side template engine designed for Java-based web environments.

Thymeleaf facilitates the synthesis of dynamic HTML views on the server side, enabling the embedding of context-sensitive data—such as session attributes, query outputs and user-submitted form content—directly into the markup prior to delivery to the client. This server-side rendering paradigm promotes a high degree of interactivity and allows for the generation of customized and context-aware web pages, thereby improving the overall responsiveness and user-centricity of the application interface.

5.3 Back-End

The back-end is implemented using **Java** due to its robustness, platform independence and strong ecosystem. Java offers a well-established set of libraries and frameworks that support rapid development of secure, scalable and maintainable web applications. Its object-oriented nature encourages clean architecture and modular design, which are essential for complex systems like healthcare platforms. Additionally, Java's widespread use in enterprise environments ensures long-term support and community-driven innovation, making it a reliable choice for production-grade applications.

5.3.1 Spring Boot

As the core of our application, we adopted the **Spring Boot** framework, which has significantly streamlined both the development and deployment processes. One of the primary

advantages of Spring Boot lies in its built-in support for an embedded **Tomcat** server, allowing the creation of a fully functional web application without requiring external configuration. This design choice has enhanced our ability to construct responsive and maintainable RESTful APIs.

In addition, to support our data management needs, we integrated **Spring Data Neo4j** and **Spring Data MongoDB**. The former facilitates interaction with the **Neo4j graph database** through a consistent repository-based abstraction, ideal for managing highly interconnected medical data such as relationships between doctors and patients. The latter offers seamless integration with **MongoDB, a document-oriented NoSQL database**, while preserving the flexibility and idiomatic programming style of the broader Spring ecosystem. Collectively, these technologies have provided a robust, scalable and modular architecture well-suited to the requirements of our application.

5.4 Back-End Structure

The application follows the layered architecture presented during lectures, organizing the codebase into clearly defined packages to promote modularity and separation of concerns.

5.4.1 `package it.unipi.healthhub.config`

The **Config** package usually contains configuration classes. In our case, we have only one class.

The **FilterConfig** class defines the registration of custom servlet filters used to manage access control within the application. Annotated with **@Configuration**, it registers filters via **FilterRegistrationBean**, associating each filter with specific URL patterns.

In detail:

- **DoctorDashboardAuthFilter**: restricts access to `/doctor/dashboard/*`;
- **PatientDashboardFilter**: applied to `/user/*`;
- **PatientApiFilter**: secures `/api/user/*`;
- **DoctorApiFilter**: secures `/api/doctor/*`;
- **LoginFilter**: manages access to the login page at `/login`.

5.4.2 `package it.unipi.healthhub.controller`

The **Controller** package manages API endpoint mappings and web page controllers for the core sections of the application, including the homepage, login, registration and doctor dashboard. All classes in this package follow the Spring MVC (Model-View-Controller) framework, ensuring a clean separation of concerns between business logic, HTTP request handling and view rendering.

More specifically:

- The **AuthController** handles authentication and registration processes for both patients (users) and doctors. Its responsibilities include rendering login, registration and password recovery views; managing user sessions; creating new user accounts; and processing password reset requests. Utility classes such as **ControllerUtil** and **HashUtil** are employed to promote modularity and code reuse.
- The **DoctorController** manages requests related to viewing a doctor's public profile. It retrieves the corresponding data via the **DoctorService** and injects it into the model for view rendering.
- The **DoctorDashboardController** supports navigation within the doctor's dashboard by mapping each HTTP GET request to a specific section such as appointments, profile, reviews, templates and weekly schedule. It retrieves the authenticated doctor's data from the session and delegates business logic to the **DoctorService**.
- The **HomeController** handles requests to the root endpoints (`/`, `/index`) and to the `/search` page. Each method ensures that session-specific data is consistently added to the model using **ControllerUtil** before returning the appropriate view.
- The **UserController** manages authenticated user interactions through endpoints under `/user`, including profile viewing, appointments and favorite doctors. It uses **ControllerUtil** to populate the model with session-based data. The class adopts a session-aware MVC design aimed at personalized user interaction.

5.4.3 package `it.unipi.healthhub.controller.api`

The `it.unipi.healthhub.controller.api` package defines a set of RESTful controllers that expose the application's core services via HTTP endpoints. These controllers are designed to return data (typically in JSON format) allowing asynchronous communication with frontend components and third-party systems.

Serving as the API layer of the system architecture, this package supports decoupled access to key resources such as users, doctors, appointments and reviews. It complements the standard MVC controllers by enabling machine-oriented interactions, in alignment with modern web application practices.

DoctorAPI

The **DoctorAPI** class defines a RESTful web interface to manage and interact with doctor-related resources within the *HealthHub* system. It is annotated with **@RestController** and is mapped to the base path `/api/doctors`. Below is a summary of the main endpoints:

- GET `/api/doctors`
Returns a list of all registered doctors.
- GET `/api/doctors/{id}`
Retrieves a specific doctor by their unique identifier.

- **POST /api/doctors**
Creates a new doctor resource.
- **PUT /api/doctors/{id}**
Updates an existing doctor identified by `id`.
- **DELETE /api/doctors/{id}**
Deletes the doctor with the specified `id`.
- **GET /api/doctors/{doctorId}/services**
Retrieves the list of services offered by the specified doctor.
- **GET /api/doctors/{doctorId}/appointments**
Returns the list of appointments associated with a doctor.
- **POST /api/doctors/{doctorId}/appointments**
Books an appointment for the authenticated patient with the given doctor.
- **GET /api/doctors/{doctorId}/templates**
Retrieves scheduling templates used by the doctor.
- **GET /api/doctors/{doctorId}/schedules/week?year=...&week=...**
Returns the weekly schedule (including prenotable slots) for a doctor, given the ISO week and year.
- **GET /api/doctors/{doctorId}/endorsements**
Retrieves the total number of endorsements for a doctor and whether the current user has endorsed them.
- **POST /api/doctors/{doctorId}/endorse**
Allows a logged-in patient to endorse a doctor.
- **POST /api/doctors/{doctorId}/unendorse**
Removes an existing endorsement made by the patient.
- **GET /api/doctors/{doctorId}/reviews**
Returns all patient reviews for the specified doctor.
- **POST /api/doctors/{doctorId}/reviews**
Submits a review for a doctor, provided the patient has had an appointment with them.
- **GET /api/doctors/{doctorId}/reviews/week/count**
Returns the number of reviews a doctor received during the current ISO week.

PrivateDoctorAPI

The `PrivateDoctorAPI` class implements a secure REST interface for authenticated doctors to manage their personal data, appointments and analytic insights. Annotated with `@RestController` and mapped to `/api/doctor`, it delegates domain logic to `AppointmentService` and `DoctorService`, enforcing session-based access control. Below is a summary of the main endpoints:

- **GET /api/doctor/appointments**
retrieves today's appointments (or for a given date).
- **DELETE /api/doctor/appointments/{id}**
deletes an appointment.
- **PUT /api/doctor/address**
updates the doctor's address.
- **PUT /api/doctor/info**
updates the doctor's profile info.
- **POST /api/doctor/phone**
adds a phone number.
- **DELETE /api/doctor/phone/{index}**
removes a phone number by a doctor.
- **POST /api/doctor/specialization**
adds a specialization.
- **DELETE /api/doctor/specialization/{index}**
removes a specialization.
- **POST /api/doctor/service**
adds a medical service.
- **GET /api/doctor/service**
retrieves all offered services.
- **PUT /api/doctor/service**
updates a medical service.
- **DELETE /api/doctor/service/{id}**
deletes a medical service.
- **POST /api/doctor/template**
creates a new availability template.
- **GET /api/doctor/template**
retrieves all availability templates.

- **PUT /api/doctor/template**
updates a template.
- **DELETE /api/doctor/template/{id}**
deletes a template.
- **PUT /api/doctor/template/default**
sets the default availability template.
- **POST /api/doctor/schedule**
creates a weekly schedule.
- **GET /api/doctor/schedule**
retrieves the weekly schedule.
- **PUT /api/doctor/schedule**
updates the weekly schedule.
- **DELETE /api/doctor/schedule**
deletes the weekly schedule.
- **GET /api/doctor/reviews**
retrieves all received reviews.
- **DELETE /api/doctor/review/{id}**
deletes a review.
- **PUT /api/doctor/password**
changes the password.
- **GET /api/doctor/stats/visits**
gets number of visits (optionally weekly).
- **GET /api/doctor/stats/earnings**
gets annual earnings.
- **GET /api/doctor/stats/patients**
gets number of new patients per month.

UserAPI

The **UserAPI** controller is a RESTful web service component. It provides endpoints for managing user entities within the HealthHub application, leveraging the **UserService** to perform CRUD (Create, Read, Update, Delete) operations on patients stored in a MongoDB database. The controller follows standard REST conventions, utilizing HTTP verbs to represent actions on resources.

- **GET /api/users**
Retrieves a list of all users registered in the system.

- **GET /api/users/{id}**
Fetches a specific user by their unique identifier. Returns HTTP 200 with the user data if found, otherwise HTTP 404.
- **POST /api/users**
Creates a new user entity based on the data provided in the request body. Returns the created user.
- **PUT /api/users/{id}**
Updates an existing user identified by the given ID with new data provided in the request body. Returns HTTP 200 with the updated user if the operation is successful; returns HTTP 404 if the user does not exist.
- **DELETE /api/users/{id}**
Deletes the user identified by the specified ID. Returns HTTP 204 No Content regardless of whether the user existed, indicating that the request has been processed.

PrivateUserAPI

The **PrivateUserAPI** controller is a RESTful web service implemented using the Spring framework in Java, designed to provide secured endpoints for user-specific operations within the HealthHub system. It leverages HTTP sessions to identify the authenticated user (typically a patient) and delegates business logic to the **UserService**. The controller exposes various endpoints to manage user details, contacts, appointments, doctor recommendation and password changes, emphasizing privacy and personalization.

- **GET /api/user/details**
Retrieves detailed personal information of the currently authenticated user by extracting the patient ID from the HTTP session. Returns user details or an error message in case of failure.
- **PUT /api/user/details**
Updates the personal details of the authenticated user. The new data is supplied in the request body as a **UserDetailsDTO** object.
- **GET /api/user/details/view?id={id}**
Retrieves contact information and demographic data for a user specified by the **id** query parameter. Additionally, if a doctor is logged in, it returns the number of visits that doctor has conducted with the user.
- **GET /api/user/contacts**
Fetches the contact information related to the authenticated user.
- **PUT /api/user/contacts**
Updates the contact information of the authenticated user using data provided in the request body as a **PatientContactsDTO**.

- **GET /api/user/appointments/upcoming**
Retrieves a list of upcoming appointments for the authenticated user.
- **GET /api/user/appointments/past**
Retrieves a list of past appointments for the authenticated user.
- **DELETE /api/user/appointments/{appointmentId}**
Cancels the appointment identified by `appointmentId`. Returns HTTP 200 on success or HTTP 400 if the appointment does not exist.
- **GET /api/user/recommendations?limit={limit}**
Returns a list of recommended doctors for the authenticated user, with an optional parameter `limit` to specify the maximum number of results (default is 3).
- **GET /api/user/doctors/reviewed**
Returns a list of doctors that the authenticated user has reviewed.
- **GET /api/user/doctors/endorsed**
Returns a list of doctors endorsed by the authenticated user.
- **PUT /api/user/password**
Allows the authenticated user to change their password. The current and new passwords are supplied in the request body as a `PasswordChangeDTO` and are hashed before validation and update.

SearchAPI

The **SearchAPI** controller is a RESTful component, responsible for executing search queries specifically targeted at medical doctors within the HealthHub system. This controller integrates data from two distinct data sources: a MongoDB database and a Neo4j graph database. The rationale behind this dual-source approach is to combine traditional document-oriented search results with personalized recommendations derived from graph-based relationships.

The primary endpoint exposed by this controller is: **GET /api/search/doctors**. This endpoint accepts a mandatory `query` parameter, representing the search string to be matched against doctor records. Additionally, it accesses the HTTP session to retrieve the identifier of the currently authenticated patient (if available), enabling personalized search refinement.

Upon receiving a valid query, the controller queries the MongoDB database via **DoctorService**, returning results as **DoctorMongoProjection** objects that represent doctor records optimized for search display. If the user is unauthenticated or the patient identifier is absent, the controller returns the top ten doctors from the MongoDB results, sorted by relevance.

For authenticated users, the controller enriches the search outcomes by querying the Neo4j graph database to retrieve personalized recommendations. These recommendations are returned as **DoctorNeo4jProjection** objects, each associated with a relevance score reflecting the strength of the doctor-patient relationship, such as prior interactions or endorsements. In the absence of Neo4j results, the system falls back to the top MongoDB results.

The integration of results from both sources is achieved by mapping doctor identifiers to Neo4j scores, then augmenting each MongoDB doctor's score accordingly. Doctors without Neo4j matches receive a default negative score to denote lack of personalized relevance. The aggregated list is sorted in descending order by combined score and the top ten entries are converted into `DoctorDTO` objects for API response consistency.

Two private helper methods, `toDto` and `toTop10Dto`, support this process by converting raw entities to DTOs and by sorting and limiting the result set. This design ensures that search responses maintain a uniform format, whether or not personalization is applied. Overall, the `SearchAPI` exemplifies a hybrid search strategy that effectively merges traditional document-based retrieval with graph-driven personalization, leveraging session data to optimize relevance and incorporating fallback mechanisms to ensure robust performance for both authenticated and anonymous users.

5.4.4 `package it.unipi.healthhub.dto`

The Data Transfer Objects (DTO) directory encapsulates a set of plain Java objects that serve as data carriers between different layers of the HealthHub application, primarily facilitating communication between the service and presentation tiers. These DTO classes are designed to aggregate and structure data relevant to specific business operations or user interactions, decoupling the internal domain model from external API contracts.

5.4.5 `package it.unipi.healthhub.exception`

The `it.unipi.healthhub.exception` package defines a set of custom unchecked exceptions, each extending `RuntimeException`, to represent specific error conditions encountered in the HealthHub application's domain logic. By encapsulating distinct failure scenarios—such as missing resources or duplicate entries—these exceptions facilitate clear error propagation and centralized handling in higher layers (e.g., controllers or global exception handlers). Moreover, they improve code readability by replacing generic error signals with semantically meaningful types.

- **`DoctorNotFoundException`**: Thrown when an operation attempts to retrieve or manipulate a doctor entity that does not exist in the system. Constructors allow for default or custom error messages and for nesting an underlying cause;
- **`ScheduleAlreadyExistsException`**: Raised to indicate an attempt to create a schedule for a doctor in a week for which a schedule is already defined. This prevents conflicting availability entries. As with other exceptions, it supports message customization and cause propagation;
- **`UserAlreadyExistsException`**: Used when a registration or user-creation process detects that a user with the same unique identifier (e.g., email or fiscal code) is already present in the database. This exception enforces uniqueness constraints at the application layer;

- **UserNotFoundException:** Thrown when a requested user resource cannot be found, for instance during authentication, profile retrieval, or update operations. Its constructors mirror those of the other exception types, enabling flexible error reporting.

5.4.6 package `it.unipi.healthhub.filter`

The `it.unipi.healthhub.filter` package provides servlet filters that centralize authentication and authorization logic for both API and web endpoints. These filters intercept incoming HTTP requests, verify session validity and user roles and either allow processing to continue, redirect to the login page, or return appropriate HTTP error codes upon access violations.

- **DoctorApiFilter:** Ensures that requests to `/api/doctor/*` originate from an authenticated doctor. If no session exists or the session attribute `role` is not `"doctor"`, the filter responds with HTTP 403 Forbidden;
- **DoctorDashboardAuthFilter:** Protects doctor dashboard pages under `/doctor/dashboard/*`. It verifies the presence of a valid session with `role = "doctor"`. If validation fails, the user is redirected to the login page;
- **LoginFilter:** Prevents authenticated users from re-accessing the login page. For existing sessions, GET requests to the login URL are redirected to `/index` (the home page), while POST attempts are rejected with HTTP 403 Forbidden;
- **PatientApiFilter:** Applies to `/api/user/*`, allowing read-only GET requests for all users but restricting state-changing methods (POST, PUT, DELETE) exclusively to authenticated patients (session attribute `role = "patient"`). Unauthorized modification attempts yield HTTP 403 Forbidden;
- **PatientDashboardFilter:** Secures patient dashboard views under `/user/*`. It checks for an active session with `role = "patient"` and redirects unauthenticated or improperly privileged users to the login page.

5.4.7 package `it.unipi.healthhub.model`

The **Model** package defines the entity classes corresponding to the database schema.

The `it.unipi.healthhub.model.mongo` package contains the core domain entities persisted in MongoDB. These classes represent the primary data models for the application and include:

- **Address.java:** encapsulates patient or doctor address fields (street, city, province, postal code and country);
- **Appointment.java:** models appointment records, linking patients and doctors with date, time and status.;
- **CalendarTemplate.java:** defines named availability templates composed of slot collections;

- **Doctor.java**: represents doctor profiles, including personal details, specializations and business rules;
- **PrenotableSlot.java**: denotes individual time slots available for booking;
- **Review.java**: captures patient-submitted reviews;
- **Schedule.java**: organizes weekly schedules as mappings from days to prenotable slots;
- **Service.java**: specifies medical services offered by a doctor, including descriptions and pricing;
- **Slot.java**: a simple time interval used within templates and schedules;
- **User.java**: models patient accounts with personal information and secure credentials.

In addition, the `it.unipi.healthhub.model.neo4j` subpackage defines lightweight DAO classes (**DoctorDAO.java**, **UserDAO.java**) used exclusively for graph queries. Since all entity attributes required by the application are already captured in the MongoDB models and relationships are managed via Neo4j projections, the MongoDB classes serve as the definitive domain models.

Their structural design has already been detailed in the Design chapter's UML class diagram section.

5.4.8 package `it.unipi.healthhub.model.mongo`

User

```

1 @Document(collection = "users")
2 public class User {
3     @Id
4     protected String id;
5
6     protected String username;
7
8     protected String name;
9     protected String password;
10    protected String fiscalCode;
11    protected LocalDate dob;
12    protected String gender;
13    private String personalNumber;
14    protected String email;
15
16    public String getId() {
17        return id;
18    }
19    public void setId(String id) { this.id = id; }
20    public String getName() {
21        return name;
22    }
23    public void setName(String name) {

```

```
24     this.name = name;
25 }
26 public String getUsername() {
27     return username;
28 }
29 public void setUsername(String username) {
30     this.username = username;
31 }
32
33 public String getFiscalCode(){ return fiscalCode; }
34 public void setFiscalCode(String fiscalCode){ this.fiscalCode =
    fiscalCode; }
35
36 public String getGender() {
37     return gender;
38 }
39 public void setGender(String gender) {
40     this.gender = gender;
41 }
42
43 public String getEmail() {
44     return email;
45 }
46 public void setEmail(String email) {
47     this.email = email;
48 }
49
50 public String getPassword() {
51     return password;
52 }
53 public void setPassword(String password) {
54     this.password = password;
55 }
56
57 public void setDob(LocalDate dob) {
58     this.dob = dob;
59 }
60 public LocalDate getDob() {
61     return dob;
62 }
63
64 public String getPersonalNumber() {
65     return personalNumber;
66 }
67
68 public void setPersonalNumber(String personalNumber) {
69     this.personalNumber = personalNumber;
70 }
71
72 public String toString() {
73     return "User{" +
74         ", username='" + username + '\',' +
75         ", password='" + password + '\',' +
76         ", dob=" + dob + "}";
```

```

77     }
78 }

```

```

1  @Document(collection = "doctors")
2  public class Doctor extends User {
3      protected String orderRegistrationNumber;
4      private List<Service> services;
5      private int endorsementCount;
6      private int reviewCount;
7      private List<Review> reviews;
8      private List<Schedule> schedules;
9      private List<String> calendarTemplates;
10     private List<String> specializations;
11     private List<String> phoneNumbers;
12     private Address address;
13
14     public Doctor() {
15         super();
16         services = new ArrayList<>();
17         reviews = new ArrayList<>();
18         schedules = new ArrayList<>();
19         calendarTemplates = new ArrayList<>();
20         specializations = new ArrayList<>();
21         phoneNumbers = new ArrayList<>();
22     }
23
24     public String getOrderRegistrationNumber() { return
25         orderRegistrationNumber; }
26
27     public void setOrderRegistrationNumber(String orderRegistrationNumber)
28         { this.orderRegistrationNumber = orderRegistrationNumber; }
29
30     public List<Service> getServices() {
31         return services;
32     }
33
34     public void setServices(List<Service> services) {
35         this.services = services;
36     }
37
38     public int getEndorsementCount() {
39         return endorsementCount;
40     }
41
42     public void setEndorsementCount(int endorsementCount) {
43         this.endorsementCount = endorsementCount;
44     }
45
46     public int getReviewCount() {
47         return reviewCount;
48     }
49
50     public void setReviewCount(int reviewCount) { this.reviewCount =
51         reviewCount; }

```



```

45
46     public List<Review> getReviews() {
47         return reviews;
48     }
49     public void setReviews(List<Review> reviews) {
50         this.reviews = reviews;
51     }
52
53     public List<Schedule> getSchedules() {
54         return schedules;
55     }
56     public void setSchedules(List<Schedule> schedules) {
57         this.schedules = schedules;
58     }
59
60     public List<String> getCalendarTemplates() {
61         return calendarTemplates;
62     }
63     public void setCalendarTemplates(List<String> calendarTemplates) {
64         this.calendarTemplates = calendarTemplates;
65     }
66
67     public List<String> getSpecializations() {
68         return specializations;
69     }
70     public void setSpecializations(List<String> specializations) {
71         this.specializations = specializations;
72     }
73
74     public List<String> getPhoneNumbers() {
75         return phoneNumbers;
76     }
77     public void setPhoneNumbers(List<String> phoneNumbers) {
78         this.phoneNumbers = phoneNumbers;
79     }
80
81     public Address getAddress() {
82         return address;
83     }
84     public void setAddress(Address address) {
85         this.address = address;
86     }
87
88     public String getId(){
89         return super.getId();
90     }
91
92     public String toString() {
93         return "User{" +
94             ", username='" + getUsername() + '\'' +
95             ", password='" + getPassword() + '\'' +
96             ", dob=" + dob + "}";
97     }
98 }

```

Appointment

```
1 @Document(collection = "appointments")
2 public class Appointment {
3     @Id
4     private String id;
5     private LocalDateTime date;
6     private DoctorInfo doctor;
7     private PatientInfo patient;
8     private String visitType;
9     private String patientNotes;
10    private double price;
11
12    // Constructors
13    public Appointment() {}
14
15    public Appointment(String id, LocalDateTime appointmentDateTime,
16                        DoctorInfo doctorInfo, PatientInfo patientInfo, String visitType,
17                        String patientNotes) {
18        this.id = id;
19        this.date = appointmentDateTime;
20        this.doctor = doctorInfo;
21        this.patient = patientInfo;
22        this.visitType = visitType;
23        this.patientNotes = patientNotes;
24    }
25
26    // Getters and Setters
27    public String getId() {
28        return id;
29    }
30
31    public void setId(String id) {
32        this.id = id;
33    }
34
35    public LocalDateTime getDate() {
36        return date;
37    }
38
39    public void setDate(LocalDateTime date) {
40        this.date = date;
41    }
42
43    public DoctorInfo getDoctor() {
44        return doctor;
45    }
46
47    public void setDoctor(DoctorInfo doctor) {
48        this.doctor = doctor;
49    }
50 }
```

```
48
49 public PatientInfo getPatient() {
50     return patient;
51 }
52
53 public void setPatient(PatientInfo patient) {
54     this.patient = patient;
55 }
56
57 public String getVisitType() {
58     return visitType;
59 }
60
61 public void setVisitType(String visitType) {
62     this.visitType = visitType;
63 }
64
65 public String getPatientNotes() {
66     return patientNotes;
67 }
68
69 public void setPatientNotes(String patientNotes) {
70     this.patientNotes = patientNotes;
71 }
72
73 public void setPrice(Double price) {
74     this.price = price;
75 }
76
77 public double getPrice() {
78     return price;
79 }
80
81 // Inner class DoctorInfo
82 public static class DoctorInfo {
83     private String id;
84     private String name;
85     private Address address;
86     private String email;
87
88     public DoctorInfo() {}
89
90     public DoctorInfo(String doctorId, String doctorName, Address address,
91         String email) {
92         this.id = doctorId;
93         this.name = doctorName;
94         this.address = address;
95         this.email = email;
96     }
97
98     public String getId() {
99         return id;
100     }
```

```
101     public void setId(String id) {
102         this.id = id;
103     }
104
105     public String getName() {
106         return name;
107     }
108
109     public void setName(String name) {
110         this.name = name;
111     }
112
113     public Address getAddress() {
114         return address;
115     }
116
117     public void setAddress(Address address) {
118         this.address = address;
119     }
120
121     public String getEmail() {
122         return email;
123     }
124
125     public void setEmail(String email) {
126         this.email = email;
127     }
128 }
129
130 // Inner class PatientInfo
131 public static class PatientInfo {
132     private String id;
133     private String name;
134     private String email;
135     private String gender;
136
137     public PatientInfo() {}
138
139     public PatientInfo(String patientId, String patientName, String email,
140         String gender) {
141         this.id = patientId;
142         this.name = patientName;
143         this.email = email;
144         this.gender = gender;
145     }
146
147     public String getId() {
148         return id;
149     }
150
151     public void setId(String id) {
152         this.id = id;
153     }
```

```
154     public String getName() {
155         return name;
156     }
157
158     public void setName(String name) {
159         this.name = name;
160     }
161
162     public String getEmail() {
163         return email;
164     }
165
166     public void setEmail(String email) {
167         this.email = email;
168     }
169
170     public String getGender() {
171         return gender;
172     }
173
174     public void setGender(String gender) {
175         this.gender = gender;
176     }
177 }
178 }
```

Address

```
1 public class Address {
2     private String street;
3     private String city;
4     private String province;
5     private String postalCode;
6     private String country;
7
8     // Constructor
9     public Address(String street, String city, String province, String
10         postalCode, String country){
11         this.street = street;
12         this.city = city;
13         this.province = province;
14         this.postalCode = postalCode;
15         this.country = country;
16     }
17
18     // Getters and Setters
19     public String getStreet() {
20         return street;
21     }
22
23     public void setStreet(String street) {
24         this.street = street;
25     }
26 }
```

```

24     }
25
26     public String getCity() {
27         return city;
28     }
29
30     public void setCity(String city) {
31         this.city = city;
32     }
33
34     public String getProvince() {
35         return province;
36     }
37
38     public void setProvince(String province) {
39         this.province = province;
40     }
41
42     public String getPostalCode() {
43         return postalCode;
44     }
45
46     public void setPostalCode(String postalCode) {
47         this.postalCode = postalCode;
48     }
49
50     public String getCountry() {
51         return country;
52     }
53
54     public void setCountry(String country) {
55         this.country = country;
56     }
57
58     public String toString() {
59         return street + ", " + city + ", " + province + ", " + postalCode + ",
60             " + country;
61     }
62
63     public String toShortString() {
64         return street + ", " + city;
65     }
66 }

```

Calendar Template

```

1     @Document(collection = "templates")
2     public class CalendarTemplate {
3         @Id
4         private String id;
5         private String name;
6         private Map<String, List<Slot>> slots;

```

```
7     private boolean isDefault;
8
9     // Constructor
10    public CalendarTemplate() {}
11
12    public CalendarTemplate(String id, String name, Map<String, List<Slot
13        >> slots, boolean isDefault) {
14        this.id = id;
15        this.name = name;
16        this.slots = slots;
17        this.isDefault = isDefault;
18    }
19
20    // Getters and Setters
21    public String getId() {
22        return id;
23    }
24
25    public void setId(String id) {
26        this.id = id;
27    }
28
29    public String getName() {
30        return name;
31    }
32
33    public void setName(String name) {
34        this.name = name;
35    }
36
37    public Map<String, List<Slot>> getSlots() {
38        return slots;
39    }
40
41    public void setSlots(Map<String, List<Slot>> slots) {
42        this.slots = slots;
43    }
44
45    public boolean isDefault(){
46        return isDefault;
47    }
48
49    public void setDefault(boolean isDefault) {
50        this.isDefault = isDefault;
51    }
52
53    public String toString() {
54        return "CalendarTemplate{" +
55            "id='" + id + '\',' +
56            ", name='" + name + '\',' +
57            ", slots=" + slots +
58            ", default=" + isDefault +
59            '}',;
```

```
60  
61 }
```

Prenotable Slot

```
1  public class PrenotableSlot extends Slot{  
2      private boolean taken;  
3  
4      public PrenotableSlot() {  
5          super();  
6      }  
7  
8      public PrenotableSlot(String start, String end, boolean taken) {  
9          super(start, end);  
10         this.taken = taken;  
11     }  
12  
13     public boolean isTaken() {  
14         return taken;  
15     }  
16  
17     public void setTaken(boolean taken) {  
18         this.taken = taken;  
19     }  
20 }
```

Review

```
1  public class Review {  
2      public String name;  
3      public String text;  
4      public LocalDate date;  
5      public String patientId;  
6  
7      public String getName() {  
8          return name;  
9      }  
10     public void setName(String name) {  
11         this.name = name;  
12     }  
13  
14     public String getText() {  
15         return text;  
16     }  
17     public void setText(String text) {  
18         this.text = text;  
19     }  
20  
21     public LocalDate getDate() {  
22         return date;  
23     }  
24 }
```



```
24     public void setDate(LocalDate date) {
25         this.date = date;
26     }
27
28     public String getPatientId() {
29         return patientId;
30     }
31     public void setPatientId(String patientId) {
32         this.patientId = patientId;
33     }
34 }
```

Schedule

```
1     public class Schedule {
2         private LocalDate week;
3         private Map<String, List<PrenotableSlot>> slots;
4
5         // Getters and Setters
6         public LocalDate getWeek() {
7             return week;
8         }
9
10        public void setWeek(LocalDate week) {
11            this.week = week;
12        }
13
14        public Map<String, List<PrenotableSlot>> getSlots() {
15            return slots;
16        }
17
18        public void setSlots(Map<String, List<PrenotableSlot>> slots) {
19            this.slots = slots;
20        }
21    }
```

Service

```
1     public class Service {
2         private String service;
3         private double price;
4
5         // Constructors
6         public Service() {}
7         public Service(String service, double price) {
8             this.service = service;
9             this.price = price;
10        }
11
12        // Getters and Setters
13        public String getService() {
```

```
14     return service;
15 }
16 public void setService(String service) {
17     this.service = service;
18 }
19
20 public double getPrice() {
21     return price;
22 }
23 public void setPrice(double price) {
24     this.price = price;
25 }
26 }
```

Slot

```
1 public class Slot {
2     protected String start;
3     protected String end;
4
5     // Constructor
6     public Slot() {}
7
8     public Slot(String start, String end) {
9         this.start = start;
10        this.end = end;
11    }
12
13    // Getters and Setters
14    public String getStart() {
15        return start;
16    }
17
18    public void setStart(String start) {
19        this.start = start;
20    }
21
22    public String getEnd() {
23        return end;
24    }
25
26    public void setEnd(String end) {
27        this.end = end;
28    }
29
30    public String toString() {
31        return "Slot{" +
32            "start='" + start + '\',' +
33            ", end='" + end + '\',' +
34            '}';
35    }
36 }
```

5.4.9 package it.unipi.healthhub.model.neo4j

User DAO

```

1  @Node("User")
2  public class UserDAO {
3      @Id
4      private String id;
5      private String name;
6
7      @Relationship(type = "ENDORSED")
8      private Set<DoctorDAO> endorsedDoctors = new HashSet<>();
9
10     @Relationship(type = "REVIEWED")
11     private Set<DoctorDAO> reviewedDoctors = new HashSet<>();
12
13     public UserDAO() {
14     }
15
16     public UserDAO(String id, String name) {
17         this.id = id;
18         this.name = name;
19     }
20
21     public String getId() {
22         return id;
23     }
24
25     public void setId(String id) {
26         this.id = id;
27     }
28
29     public String getName() {
30         return name;
31     }
32
33     public void setName(String name) {
34         this.name = name;
35     }
36
37     public Set<DoctorDAO> getEndorsedDoctors() {
38         return endorsedDoctors;
39     }
40
41     public Set<DoctorDAO> getReviewedDoctors() {
42         return reviewedDoctors;
43     }
44 }

```

Doctor DAO

```

1  @Node("Doctor")

```

```
2 public class DoctorDAO {
3     @Id
4     private String id;
5
6     private String name;
7     private List<String> specializations;
8
9     public DoctorDAO() {
10         specializations = new ArrayList<>();
11     }
12
13     public DoctorDAO(String id, String name, List<String> specializations)
14     {
15         this.id = id;
16         this.name = name;
17         this.specializations = specializations;
18     }
19
20     public String getId() {
21         return id;
22     }
23
24     public void setId(String id) {
25         this.id = id;
26     }
27
28     public String getName() {
29         return name;
30     }
31
32     public void setName(String name) {
33         this.name = name;
34     }
35
36     public List<String> getSpecializations() {
37         return specializations;
38     }
39
40     public void setSpecializations(List<String> specializations) {
41         this.specializations = specializations;
42     }
43
44     public String toString() {
45         return "DoctorDAO{" +
46             "id='" + id + '\',' +
47             ", name='" + name + '\',' +
48             ", specializations=" + specializations +
49             '}' ;
50     }
51 }
```

5.4.10 package `it.unipi.healthhub.projection`

The `it.unipi.healthhub.projection` package contains projection classes used to encapsulate doctors and their relevance scores obtained from search queries. Specifically, `DoctorMongoProjection` wraps a MongoDB `Doctor` entity with a score, while `DoctorNeo4jProjection` wraps a Neo4j `DoctorDAO` with a similar score. These projections facilitate merging and ranking results from both databases during doctor search operations, as implemented in the `SearchAPI` controller.

5.4.11 package `it.unipi.healthhub.repository`

The **Repository** package is responsible for managing data access operations by leveraging Spring Data JPA interfaces. This package abstracts the underlying persistence mechanisms, allowing for convenient CRUD operations and custom query definitions on the domain entities.

Detailed examination and implementation of specific queries within these repositories will be thoroughly discussed in the dedicated chapter on *Queries*.

5.4.12 package `it.unipi.healthhub.service`

The **Service** layer encapsulates the core business logic. This package is composed by: `AppointmentService`, `DoctorService` and `UserService`.

Appointment Service

The `AppointmentService` class provides a comprehensive set of business logic functionalities related to appointment management within the system. This service interacts directly with the `AppointmentMongoRepository` to perform CRUD operations on appointment entities stored in MongoDB.

Key aspects of `AppointmentService` include:

- **Data Sanitization:** To prevent operator injection vulnerabilities specific to MongoDB, the service sanitizes input strings by replacing prohibited characters (e.g., '\$') within appointment fields, including nested doctor and patient information such as names, emails and addresses.
- **CRUD Operations:** The service exposes methods to retrieve all appointments, fetch an appointment by its identifier, create new appointments after sanitization, update existing ones and delete appointments by ID.
- **Asynchronous Updates:** Leveraging Spring's `@Async` annotation, the service asynchronously updates related doctor and patient information within appointments, improving responsiveness and decoupling these potentially costly operations.
- **Domain-Specific Queries:** It supports domain-specific queries such as retrieving all appointments for a given doctor on a specified day, facilitating calendar and scheduling functionalities.

Doctor Service

The `DoctorService` class manages business logic related to doctor profiles, endorsements, reviews, specializations, analytics and schedule maintenance. It coordinates between MongoDB (via `doctorMongoRepository`) for persistent doctor data and Neo4j (via `doctorNeo4jRepository` and `userNeo4jRepository`) for relationship data.

Key features of `DoctorService` include:

- **Data Sanitization:** Prevents injection vulnerabilities by sanitizing input strings differently for MongoDB (e.g., replacing '\$') and Neo4j (e.g., replacing '\$' and '.'). Sanitization applies to doctor fields and nested user data.
- **Reviews Management:** Enables patients to add and remove reviews after verifying appointments, updating both doctor documents and Neo4j relationships.
- **Profile Updates:** Supports login authentication, updating doctor details such as address, personal info and phone numbers, with data sanitization and synchronization between MongoDB and Neo4j.
- **Specializations Handling:** Allows adding and removing specializations stored in MongoDB and mirrored in Neo4j, returning indices for reference.
- **Endorsements:** Manages patient endorsements for doctors by updating counts in MongoDB and maintaining endorsement relationships in Neo4j.
- **Analytics:** Provides visit and earnings statistics aggregated by type, week, month and year, leveraging the appointment repository.
- **Schedule Management:** Maintains up-to-date doctor schedules by removing outdated entries and creating new ones for missing weeks, based on default calendar templates. The process includes progress logging and exception handling for missing templates.
- **Transactional Consistency:** Critical updates are annotated with `@Transactional` to ensure atomic operations.

User Service

The `UserService` class handles all user-related business logic, managing users stored in both MongoDB and Neo4j databases. It ensures data integrity and synchronization across these stores while supporting user authentication, profile management and interactions with appointments and doctors.

Key features of `UserService` include:

- **Data Sanitization:** Prevents injection vulnerabilities by sanitizing input strings differently for MongoDB (e.g., replacing '\$') and Neo4j (e.g., replacing '\$' and '.'). Sanitization applies to user fields and nested doctor data.

- **User CRUD and Authentication:** Provides creation, retrieval, update and deletion of users with transactional integrity. Includes username/email uniqueness checks and password validation for login and password changes.
- **Profile and Contact Management:** Updates and retrieves detailed user profile information and contact data, synchronizing name changes across appointments and Neo4j relationships.
- **Appointment Integration:** Retrieves past and upcoming appointments for users, supports appointment cancellation with schedule slot management and notification via a fake mail sender.
- **Neo4j Relations and Recommendations:** Manages endorsement and review relations between users and doctors in Neo4j, fetching lists of endorsed, reviewed and recommended doctors. Recommendation logic combines personalized and popular doctors.
- **Email Operations:** Sends password reset links using a mocked mail sender component.

5.4.13 package `it.unipi.healthhub.util`

This package contains several utility classes and interfaces used within the HealthHub application to support functionalities like session handling, date manipulation, email sending, password hashing and conversion between DTOs and model objects.

- **ControllerUtil.java**
Provides utility methods for controllers, especially for setting common attributes in the Spring `Model` based on the user session. For example, it adds attributes indicating whether a user is logged in and their role (patient or doctor).
- **DateUtil.java**
Contains static methods for date calculations and conversions. Key methods include:
 - Getting the first day of a specific week in a year.
 - Checking if two dates belong to the same week.
 - Getting the next four Mondays after a given date.
 - Conversion between `java.util.Date` and `java.time.LocalDate`.
- **FakeMailSender.java**
A service implementation of the `MailSenderService` interface that simulates sending emails by printing email content to the console. It handles sending notification emails related to appointment deletions and password reset links, validating email addresses before "sending".

- **HashUtil.java**
Provides a method for hashing passwords using SHA-256. It converts the password string into its hashed hexadecimal representation, throwing a runtime exception if hashing fails.
- **MailSenderService.java**
An interface defining methods for sending various types of emails related to appointment cancellations and password resets. Implemented by classes such as **FakeMailSender**.
- **ScheduleConverter.java**
Contains static methods to convert schedule data between DTO representations and model objects. It supports transforming collections of slots and building a schedule for a given Monday date based on a calendar template.
- **TemplateConverter.java**
Provides static methods for converting template-related slot data between DTOs and model classes, including converting regular slots and "prenotable" (bookable) slots, supporting the workflow of schedule creation and manipulation.

5.4.14 Scheduler Component

The **Scheduler** class is a Spring-managed component responsible for automating weekly maintenance tasks on doctor schedules. Its **scheduleTask()** method executes every Sunday at midnight, invoking two key **DoctorService** operations:

- **cleanOldSchedules()**: Removes all schedule entries up to the current date, delegating to the repository's **cleanOldSchedules(Date)** method to purge obsolete calendar data.
- **setupNewSchedules()**: Identifies doctors lacking schedules in the next four weeks and for each missing week constructs new **Schedule** objects based on their default **CalendarTemplate**. Progress is logged in intervals to monitor throughput.

By centralizing these operations in a scheduled job, the **Scheduler** ensures that doctor availabilities remain current without manual intervention, leveraging reusable service methods and Spring's scheduling infrastructure to maintain data consistency and system reliability.

Chapter 6

Queries

In this chapter, we present only the complex queries executed by our application, without dwelling on trivial queries.

6.1 Mongo DB

6.1.1 Search Query

This query is run every time a guest user, a patient or a doctor inserts a character in the search bar. It works in the same way for all three actors. Returns a list of doctors based on the entered string. The string can be a name, a specialization, a city, a province or a mix of these.

This is the Java code:

```
1  @Override
2  public List<DoctorMongoProjection> searchDoctors(String text) {
3      // Match documents using full-text search
4      AggregationOperation match = match(Criteria.where("$text").is(new
5          Document("$search", text)));
6
7      // Add MongoDB's built-in text relevance score
8      AggregationOperation project = context -> new Document("$project",
9          new Document("doctor", "$$ROOT")
10         .append("score", new Document("$meta", "textScore")))
11     );
12
13     // Sort by relevance score
14     AggregationOperation sort = context -> new Document("$sort",
15         new Document("score", new Document("$meta", "textScore")))
16     );
17
18     // Limit to top 250 results
19     AggregationOperation limit = limit(250);
20
21     // Build the aggregation pipeline
22     Aggregation agg = newAggregation(match, project, sort, limit);
```

```

23 // Execute aggregation
24 AggregationResults<DoctorMongoProjection> results =
25 mongoTemplate.aggregate(agg, "doctors", DoctorMongoProjection.class);
26
27 return results.getMappedResults();
28 }

```

This is its one-to-one translation into a MongoDB query.

```

1 db.doctors.aggregate([
2   {
3     $match: {
4       $text: { $search: query }
5     }
6   },
7   {
8     $project: {
9       doctor: "$$ROOT",
10      score: { $meta: "textScore" }
11    }
12  },
13  {
14    $sort: {
15      score: { $meta: "textScore" }
16    }
17  },
18  {
19    $limit: 250
20  }
21 ]);

```

The above aggregation pipeline enables a full-text search over the **doctors** collection, returning the most relevant documents according to a weighted scoring scheme. Although the underlying text index will be discussed in detail in the dedicated indexing section, it is important to mention that the fields **name**, **specializations**, **address.city**, and **address.province** are indexed with different weights, thereby influencing the ranking of the search results.

The pipeline is composed of four main stages:

- **\$match**: This stage filters documents that match the full-text search criteria specified by the user input. The search is performed across all indexed fields using MongoDB's **\$text** operator;
- **\$project**: At this stage, each document is projected into a new structure where the original document is nested under the **doctor** field. Additionally, a new field **score** is introduced, which represents the text relevance score computed by MongoDB using its internal ranking algorithm based on term frequency and field weights;
- **\$sort**: Documents are then ordered in descending order of their **score**, so that the most relevant matches (according to the text index configuration) are prioritized in the output;

- `$limit`: Finally, the result set is truncated to a maximum of 250 documents, which serves both as a performance optimization and a practical upper bound for rendering results in the user interface.

6.1.2 Analytics

Earnings

This query is an analytic for the doctor's Dashboard and returns a map of monthly revenues for a given year.

```

1  @Override
2  public Map<String, Double> getEarningsByYearForDoctor(String doctorId,
3      Integer year) {
4      MatchOperation matchOperation = Aggregation.match(Criteria.where("
5          doctor.id").is(doctorId)
6          .and("date")
7          .gte(LocalDate.of(year, 1, 1).atStartOfDay())
8          .lt(LocalDate.of(year + 1, 1, 1).atStartOfDay())
9      );
10     ProjectionOperation projectMonth = Aggregation.project().andExpression(
11         ("month(date)").as("month").and("price").as("price");
12     GroupOperation groupOperation = Aggregation.group("month").sum("price")
13         .as("total");
14     ProjectionOperation projectionOperation = Aggregation.project("total")
15         .and("month").previousOperation();
16
17     Aggregation aggregation = Aggregation.newAggregation(
18         matchOperation,
19         projectMonth,
20         groupOperation,
21         projectionOperation
22     );
23
24     AggregationResults<DBObject> results = mongoTemplate.aggregate(
25         aggregation, Appointment.class, DBObject.class);
26
27     Map<String, Double> earningsByYear = new HashMap<>();
28     for (DBObject doc : results.getMappedResults()) {
29         Integer month = (Integer) doc.get("month");
30         String monthString = new DateFormatSymbols(Locale.ENGLISH).getMonths(
31             ) [month-1].toLowerCase();
32         Double total = (doc.get("total") instanceof Number) ? ((Number) doc.
33             get("total")).doubleValue() : null;
34         earningsByYear.put(monthString, total);
35     }
36
37     return earningsByYear;
38 }

```

```

1  db.appointments.aggregate([
2  {

```

```

3   $match: {
4     "doctor._id": ObjectId(doctorId),
5     date: {
6       $gte: ISODate("${year}-01-01T00:00:00Z"),
7       $lt: ISODate("${year + 1}-01-01T00:00:00Z")
8     }
9   }
10 },
11 {
12   $project: {
13     month: { $month: "$date" },
14     price: 1
15   }
16 },
17 {
18   $group: {
19     _id: "$month",
20     total: { $sum: "$price" }
21   }
22 },
23 {
24   $project: {
25     month: "$_id",
26     total: 1,
27     _id: 0
28   }
29 }
30 ]);

```

The aggregation pipeline shown above is designed to compute the monthly revenue generated by a specific doctor over the course of a given calendar year. This query supports the analytical needs of the doctor's dashboard, enabling a temporal breakdown of financial performance.

The pipeline consists of the following sequential stages:

- **\$match:** This stage performs a selective filtering of documents within the **appointments** collection. It restricts the result set to only those documents where the embedded **doctor._id** field matches the provided identifier and the **date** field falls within the specified year. The temporal filtering is implemented using the **\$gte** and **\$lt** operators to define an inclusive-exclusive date range between the first day of January and the first day of the following year;
- **\$project:** Each document is then transformed to retain only two fields: the **price** of the appointment and a computed **month** field, which is extracted from the **date** using MongoDB's **\$month** operator. This transformation facilitates the grouping operation in the subsequent stage;
- **\$group:** Documents are aggregated by the computed **month** value. For each month, the total revenue is computed by summing the values of the **price** field using the **\$sum** accumulator. The grouping key, **_id**, is implicitly set to the month number (ranging from 1 to 12);

- **\$project**: In the final projection stage, the output is restructured to explicitly include the `month` and `total` fields, while omitting the default `_id` field. This renders the result set suitable for direct transformation into a month-to-amount mapping.

The result of this pipeline is a flat collection of documents, each representing the total earnings for a specific month. In the application layer, this output is subsequently mapped into a `Map<String, Double>`, where each key corresponds to the English name of a month (e.g., `"january"`, `"february"`), and each value denotes the corresponding revenue in that month. This processed result can be directly visualized in the dashboard through line charts.

Visit Type Summary

...

Bibliography