

HealthHub: A Healthcare Data Management System

Paolo Palumbo, Francesco Panattoni, Nedal Elezaby

June 20, 2025

Contents

1	Introduction	3
2	Dataset and Web Scraping	4
2.1	Web Scraping	4
2.2	Synthetic Data Generation	4
2.3	Dataset Characteristics: Velocity and Variety	5
2.3.1	Velocity	5
2.3.2	Variety	5
2.4	Resulting Dataset	5
3	Design	6
3.1	Actors	6
3.2	Functional Requirements	6
3.2.1	Guest Users	6
3.2.2	Patients	6
3.2.3	Doctors	7
3.3	Non-Functional Requirements	7
3.3.1	Performance and Scalability	7
3.3.2	Availability and Reliability	7
3.3.3	Security and Privacy	7
3.3.4	Usability	8
3.3.5	Portability and Flexibility	8
3.3.6	Maintainability	8
3.4	Use Case Diagram	8
3.5	UML Class Diagram	8
3.5.1	Relationships between classes	8
3.5.2	Design of the classes	10
4	Data Model	13
4.1	Document Database - MongoDB	13
4.1.1	Collections	13
4.1.2	Document Structure	15
4.2	Graph Database - Neo4j	18
4.2.1	Node Types	18
4.2.2	Relationships	19
4.3	Distributed Database Design	19
4.3.1	Replica Set Deployment	19
4.3.2	Sharding (Design Perspective)	20
4.3.3	Handling Inter-database Consistency	20

5	Implementation	22
5.1	Front-End	22
5.2	Front-End Structure	22
5.2.1	Static	22
5.2.2	Templates	23
5.3	Back-End	23
5.3.1	Spring Boot	23
5.4	Back-End Structure	23
5.4.1	package it.unipi.healthhub.config	23
5.4.2	package it.unipi.healthhub.controller	24
5.4.3	package it.unipi.healthhub.controller.api	24
5.4.4	package it.unipi.healthhub.dto	29
5.4.5	package it.unipi.healthhub.exception	29
5.4.6	package it.unipi.healthhub.filter	29
5.4.7	package it.unipi.healthhub.model	30
5.4.8	package it.unipi.healthhub.model.mongo	30
5.4.9	package it.unipi.healthhub.model.neo4j	40
5.4.10	package it.unipi.healthhub.projection	42
5.4.11	package it.unipi.healthhub.repository	42
5.4.12	package it.unipi.healthhub.service	42
5.4.13	package it.unipi.healthhub.util	43
5.4.14	Scheduler Component	44
6	Queries	45
6.1	Mongo DB	45
6.1.1	Search Query	45
6.1.2	Analytics	46
6.2	Neo4j	52
6.2.1	Recommendation Queries	52
6.3	Search Query Support	56
6.4	Indexes	58
6.4.1	MongoDB	58
6.4.2	Neo4j	59
A	Original Regex-Based Query	60
B	User Manual	63
B.1	Patient	63
B.1.1	Signup and Login Process	63
B.1.2	Doctor Profile and Appointment Booking	66
B.1.3	Appointments	68
B.1.4	Recommendation System	68
B.1.5	Profile	69
B.2	Doctor	69
B.2.1	Dashboard	69
B.2.2	Appointment List	71
B.2.3	Personal Profile	71
B.2.4	Review Page	73
B.2.5	Calendar Templates	74
B.2.6	Weekly Schedule	74

Chapter 1

Introduction

Healthhub is a web-based platform designed to simplify and centralize the management of medical appointments. It offers dedicated features for both patients and healthcare professionals. Regular users can search for doctors, book visits, and leave reviews, while medical professionals can manage the services they offer and track their appointments.

- Allow users to search for medical professionals and retrieve information on their services
- Allow users to book appointments and review doctors after a visit
- Allow users to manage their appointment history and cancel or modify bookings
- Give users recommendations on doctors based on their location and past activity
- Allow doctors to specify their specializations, availability, and offered services
- Allow doctors to view and manage their upcoming appointments

Chapter 2

Dataset and Web Scraping

To populate the dataset for our application, we used a combination of web scraping, synthetic data generation, and logical inference based on real-world patterns of user behavior.

2.1 Web Scraping

The initial data was scraped from a platform containing medical reviews, MioDottore¹. This platform provides a comprehensive database of medical professionals, including their specializations, locations, and user reviews. The scraping process involved extracting the following information:

- Doctor names and specializations
- Locations (cities and regions)
- User reviews and ratings
- Contact information (where available)
- Service offerings (e.g., types of consultations, treatments)

2.2 Synthetic Data Generation

To enhance the dataset, we generated synthetic data to simulate a more realistic user base and appointment history. This involved several steps:

- **Unique Users:** We extracted all unique usernames from the scraped reviews. For each user, we generated a detailed profile including demographic and geographical information.
- **Appointments:** For every review, we created a corresponding medical appointment. These were dated in the weeks preceding the review, to simulate a realistic flow where users leave feedback shortly after being treated.
- **Likes:** To simulate engagement features such as likes on reviews, we analyzed the provinces of the doctors each user had interacted with. We then randomly selected other doctors in the same provinces and associated a number of likes comparable to the number of reviews written by each user.

¹<https://www.miodottore.it/>

2.3 Dataset Characteristics: Velocity and Variety

2.3.1 Velocity

The application context is inherently dynamic, and our dataset reflects this through the following characteristics:

- **High review frequency:** The dataset reflects an average of at least 100 new reviews per day, simulating the continuous flow of patient feedback that a live system would experience.
- **Growing doctor base:** Based on trends from medical boards, we estimate around 400 to 500 new doctors would register on the platform annually, contributing to the system’s dynamism and evolving content.

2.3.2 Variety

The dataset incorporates diverse data types, which contribute to its heterogeneity:

- **Multi-format records:** Structured user and doctor profiles, unstructured textual reviews, and timestamped interaction logs (likes, appointments) are all represented.
- **Diverse entities:** The inclusion of users, doctors, reviews, likes, and appointments enables multi-relational analysis and feature richness.

2.4 Resulting Dataset

The original scraped dataset, stored in `scraped.json` (265 MB), contains the raw information collected from the web. This dataset served as the foundation for the data generation process, providing:

- 699,987 reviews
- 214,682 unique reviewers
- 87,632 doctors

Building upon this, the final dataset is stored in JSON format, totaling approximately 960 MB, and includes both real and synthetic data entries. Its main components are:

- **Doctors** (`doctors.json`, 289 MB) Contains 87,632 healthcare professionals, each with profile data and linked to 699,987 reviews. Reviews include timestamps and patient feedback.
- **Users** (`users.json`, 66 MB) Comprises 214,682 unique users extracted from the original dataset and extended with synthetic profiles.
- **Appointments** (`appointments.json`, 422 MB) Each review is connected to a synthetic appointment, scheduled in the weeks preceding the review date. Appointments simulate realistic scheduling and clinic visit patterns.
- **Templates** (`templates.json`, 162 MB) Stores template structures for appointment scheduling logic, such as available time slots, weekdays, and timing constraints.
- **User Likes** (`user_likes.json`, 22 MB) Represents user interactions with doctors, generated according to the geographic distribution of the doctors reviewed by each user.

Overall, this dataset offers a realistic simulation of user and doctor activity on a healthcare review platform, reflecting both volume and behavioral complexity.

Chapter 3

Design

3.1 Actors

The application involves three primary categories of actors, each with distinct roles and permissions:

- *Non-Authenticated User (Guest User)*: refers to an anonymous individual who accesses the application without logging in. This actor is allowed to register, authenticate using existing credentials, or explore the platform by searching for doctors and viewing their public profiles.
- *Patient (User)*: represents the end-user of the service. Once authenticated, the patient can book appointments with doctors, confirm or cancel them, and subsequently provide feedback through reviews.
- *Doctor*: a professional figure who offers medical appointments. The doctor can manage their availability schedule, making it accessible to patients, and consult an analytics dashboard displaying aggregated data such as earnings, reviews, and the number of visits performed.

3.2 Functional Requirements

The following section outlines the functional requirements.

3.2.1 Guest Users

Guest users are allowed to:

- Register for an account within the application;
- Log in using their credentials;
- Initiate the password recovery process in case of forgotten credentials;
- Search for doctors and access their public profiles.

3.2.2 Patients

Authenticated users identified as **patients** are granted the following capabilities:

- Manage and update their personal profiles;
- Search for doctors and access their public profiles;
- Book appointments and optionally include a note for the doctor;
- Endorse a doctor as a form of support or recommendation;

- Submit reviews regarding their medical experience;
- View their appointments, including past, current, and upcoming ones;
- Search for doctors through the recommendation feature provided by the system.

3.2.3 Doctors

Authenticated users identified as **doctors** are provided with the following functionalities:

- Manage and update their personal and professional profiles;
- Search for other doctors and access their public profiles;
- Define and configure the types of visits offered, along with the associated pricing;
- Oversee appointments by identifying scheduled patients and appointments for the current day;
- Configure and manage appointment availability through customizable scheduling templates;
- Monitor and respond to patient reviews;
- Access analytics, including visual representations of revenue, newly acquired patients, recent reviews, and a summary of completed visits.

3.3 Non-Functional Requirements

The following non-functional requirements address the quality attributes of the system, ensuring its performance, reliability, security, scalability, and maintainability.

3.3.1 Performance and Scalability

- The system shall maintain **acceptable response times** for common operations such as viewing medical records and scheduling appointments;
- The system shall efficiently **handle increased workload** during peak usage periods without degradation of service.

3.3.2 Availability and Reliability

- The application shall be **available** to all users **24/7**, minimizing downtime through redundancy and failover mechanisms;
- The system shall implement **backup and recovery procedures** to preserve data integrity and support rapid restoration after failures.
- The system shall tolerate occasional data staleness in non-critical views, ensuring high availability even under degraded conditions.

3.3.3 Security and Privacy

- The system shall enforce **secure, authenticated access for all users**, employing strong password policies and session management;
- All sensitive data shall be encrypted both in transit and at rest;
- The system provides defenses against injections.

3.3.4 Usability

- The user interface shall be intuitive and **user-friendly**, enabling users to perform tasks with minimal learning curve.
- The application shall exhibit **low latency** in user interactions to maintain a responsive experience.

3.3.5 Portability and Flexibility

- The application shall be **deployable on multiple operating systems** (e.g., Windows, macOS, Linux) without requiring behavioral changes.
- The system's architecture shall support the addition of new attributes or modules (e.g., new appointment types) with minimal code modification.

3.3.6 Maintainability

- The codebase shall follow **Object-Oriented design** principles, promoting modularity and ease of comprehension.
- The system shall **minimize single points of failure** through component decoupling and redundancy.
- The application shall include comprehensive documentation and code comments to facilitate future enhancements and debugging.

3.4 Use Case Diagram

The use case diagram illustrates the core functionalities of the *HealthHub* application, categorized by the actors involved: *Guest*, *User (Patient)*, and *Doctor*. The Guest actor can access basic features such as registration, login, and browsing doctors. The authenticated user (Patient) can view and edit their profile, search for doctors, view doctor reviews, book and cancel appointments, endorse or review doctors, and receive personalized recommendations. The Doctor actor has access to a personal dashboard from which they can manage their schedule, define and edit calendar templates, view appointments, delete reviews, and handle their availability. Use cases are modeled with <<include>> and <<extend>> relationships to reflect functional modularity and reuse of common interactions like viewing profiles or appointments.

3.5 UML Class Diagram

3.5.1 Relationships between classes

Here we briefly summarise the relation between each class.

- **Doctor** is a specialization of **User**;
- A **Doctor** may offer zero or more **Service** instances and each **Service** is provided by exactly one **Doctor**;
- A **Doctor** may have zero or more **Review** entries and each **Review** refers to exactly one **Doctor**;
- A **Doctor** may define zero or more **CalendarTemplates** and each **CalendarTemplate** belongs to exactly one **Doctor**;
- A **Doctor** may publish zero or more **Schedules** and each **Schedule** is associated with exactly one **Doctor**;
- Each **Schedule** comprises zero or more **PrenotableSlots** and each **PrenotableSlot** is part of exactly one **Schedule**;

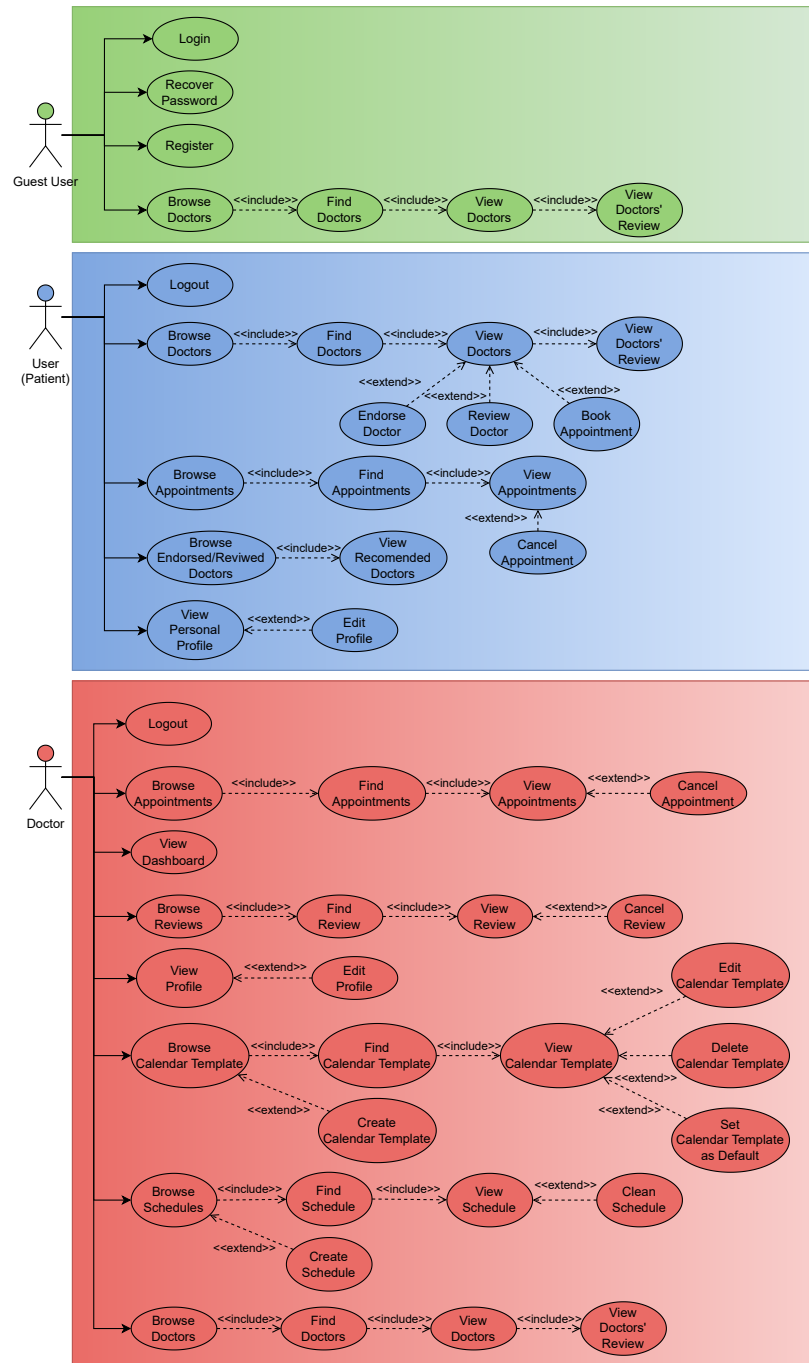


Figure 3.1: Use Case Diagram for HealthHub

- Each **CalendarTemplate** comprises zero or more **Slots** and each **Slot** belongs to exactly one **CalendarTemplate**;
- A **Doctor** has exactly one **Address** and each **Address** is linked to exactly one **Doctor** (and similarly to **DoctorInfo**);

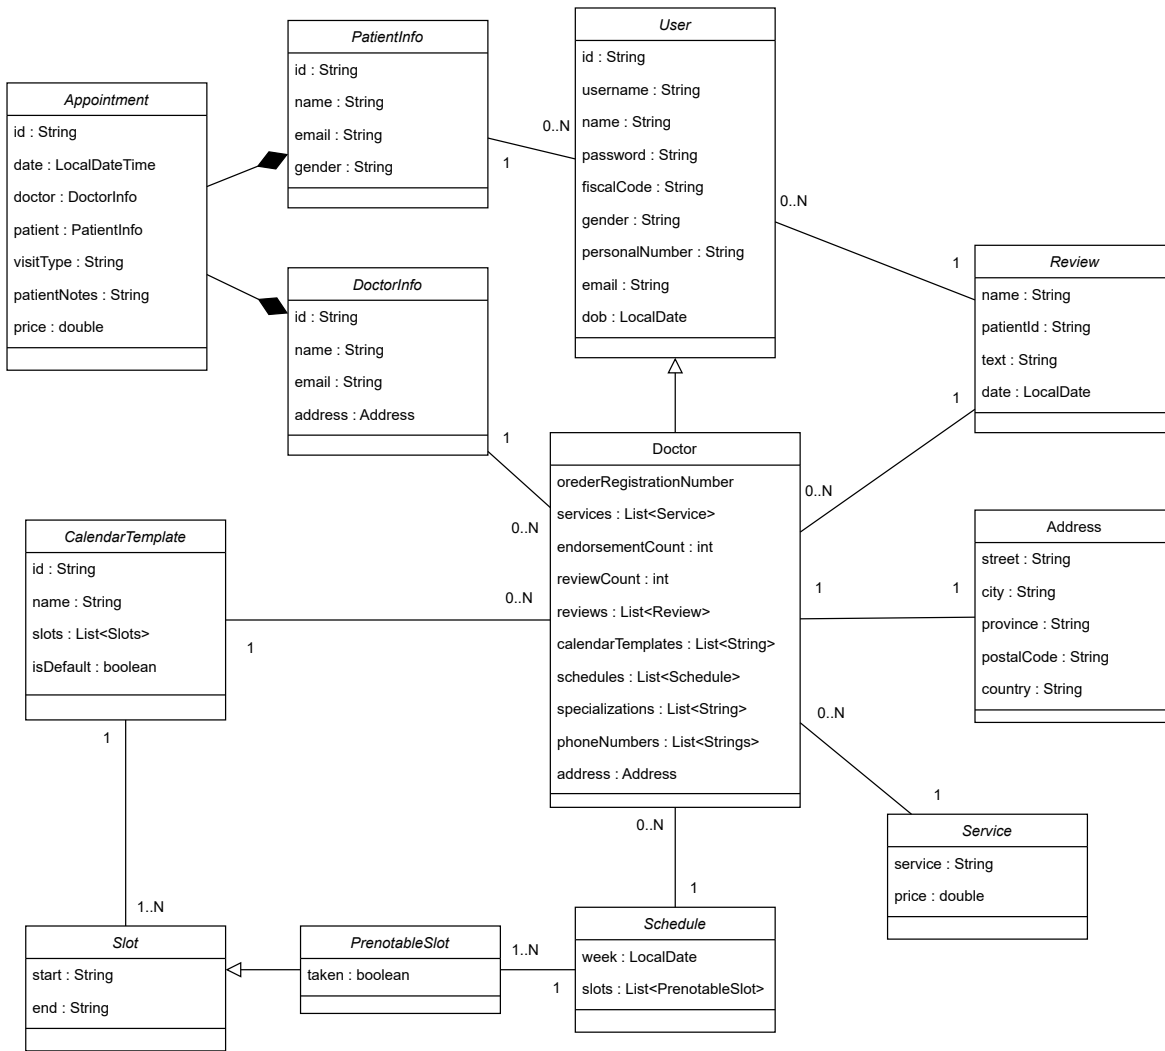


Figure 3.2: UML Class Diagram for HealthHub

- An **Appointment** is associated with exactly one **DoctorInfo** and one **PatientInfo** and each **DoctorInfo** or **PatientInfo** may have zero or more **Appointments**.

3.5.2 Design of the classes

Here we briefly summarise the primary attributes of each entity.

User

- `id`: `String` – unique identifier
- `username`: `String`
- `name`: `String`
- `password`: `String` – stored as hash
- `fiscalCode`: `String`
- `gender`: `String`

- `personalNumber: String`
- `email: String`
- `dob: LocalDate`

Doctor (specialization of **User**)

- `orderRegistrationNumber: String`
- `services: List<Service>`
- `endorsementCount: int`
- `reviewCount: int`
- `reviews: List<Review>`
- `calendarTemplates: List<CalendarTemplate>`
- `schedules: List<Schedule>`
- `specializations: List<String>`
- `phoneNumbers: List<String>`
- `address: Address`

Address

- `street: String`
- `city: String`
- `province: String`
- `postalCode: String`
- `country: String`

Appointment

- `id: String`
- `date: LocalDateTime`
- `visitType: String`
- `patientNotes: String`
- `price: double`

Schedule

- `week: LocalDate`
- `slots: List<PrenotableSlot>`

CalendarTemplate

- `id: String`
- `name: String`
- `slots: List<Slot>`
- `isDefault: boolean`

Slot

- `start: String`
- `end: String`

PrenotableSlot (subclass of **Slot**)

- taken: boolean

Review

- name: String
- patientId: String
- text: String
- date: LocalDate

Service

- service: String
- price: double

PatientInfo (belong to **Appointment**)

- id: String
- name: String
- email: String
- gender: String

DoctorInfo (belong to **Appointment**)

- id: String
- name: String
- email: String
- address: Address

Chapter 4

Data Model

4.1 Document Database - MongoDB

The application utilizes MongoDB, a document-oriented NoSQL database, to store and manage data.

4.1.1 Collections

Users

Each user in the system is a registered patient with access to the platform's features such as booking appointments and interacting with doctors. Below are the typical attributes found in a user's document:

- `_id` : unique identifier of the user, stored as an `ObjectId`.
- `fiscalCode` : string representing the Italian *codice fiscale*, used as a unique tax and identity code.
- `name` : string containing the full name of the user.
- `dob` : date indicating the user's date of birth, stored in `ISODate` format.
- `gender` : string that specifies the gender of the user (e.g., "male", "female").
- `personalNumber` : string that holds the user's personal phone number.
- `email` : string containing the email address provided during registration.
- `username` : unique string chosen by the user during sign up and used to log in.
- `password` : string containing the hashed version of the password set at registration.

Doctors

Each doctor in the system is a registered healthcare professional who can provide services, receive reviews from patients, and manage their availability through calendar templates. Below are the typical attributes found in a doctor's document:

- `_id` : unique identifier of the doctor, stored as an `ObjectId`.
- `name` : string containing the full name of the doctor.
- `email` : string with the professional email address used for communication.
- `username` : unique string used for login and identification within the platform.
- `password` : string containing the hashed version of the doctor's password.

- **address** : embedded object containing **street**, **city**, **province**, **country**, and **postalCode**.
- **phoneNumbers** : array of strings listing the phone numbers associated with the doctor.
- **specializations** : array of strings indicating the medical fields in which the doctor is specialized.
- **services** : array of embedded documents, each with a **service** name and a corresponding **price**.
- **endorsementCount** : integer representing the number of professional endorsements received.
- **reviews** : array of embedded documents, each containing a **patientId**, **name**, **text** of the review, and the **date** of submission.
- **reviewsCount** : integer indicating the total number of reviews received by the doctor.
- **dob** : date representing the doctor's date of birth, stored in ISODate format.
- **fiscal_code** : string corresponding to the Italian fiscal code (codice fiscale).
- **orderRegistrationNumber** : string representing the official registration number in the professional healthcare registry.
- **calendarTemplates** : array of **ObjectIds** referencing availability templates used for scheduling appointments.

Appointments

The **Appointments** collection stores all scheduled medical visits between patients and doctors. Each document represents a single appointment and includes contextual details about the involved parties and the visit itself.

- **_id** : unique identifier of the appointment, stored as an **ObjectId**.
- **date** : ISODate indicating the exact date and time of the scheduled appointment.
- **doctor** : embedded object including:
 - **_id** : reference to the doctor's unique identifier.
 - **name** : full name of the doctor.
 - **address** : object containing **street**, **city**, **province**, **country**, and **postalCode**.
 - **email** : contact email address of the doctor.
- **patient** : embedded object including:
 - **_id** : reference to the patient's unique identifier.
 - **name** : full name of the patient.
 - **fiscalCode** : Italian fiscal code (codice fiscale) of the patient.
 - **email** : contact email address of the patient.
 - **gender** : gender of the patient.
- **visitType** : string describing the type of visit or consultation.
- **patientNotes** : string field where patients may leave notes prior to the appointment (optional).
- **price** : numeric value indicating the cost of the visit in euros.

CalendarTemplates

The `CalendarTemplates` collection stores predefined weekly availability templates that doctors can use to generate their actual working calendars. Each template specifies on which days and time intervals the doctor is available for appointments.

- `_id` : unique identifier of the calendar template, stored as an `ObjectId`.
- `name` : name of the template (e.g., "Standard").
- `slots` : an object mapping weekdays to arrays of available time intervals. Each interval is represented as an object with:
 - `start` : start time of the slot (format: HH:MM).
 - `end` : end time of the slot (format: HH:MM).
- `isDefault` : boolean indicating whether the template is the doctor's default schedule.

4.1.2 Document Structure

Listing 4.1: Example of a User Document

```

1 {
2   _id: ObjectId('684ada4637804916ca651761'),
3   fiscalCode: 'UGOSIG070805MA41',
4   name: 'Sig. Ugolino Franscini',
5   password: '5e884898da28047151d0e56f8dc6292773603d0d6aabbdd62a11ef721d1
6     542d8',
7   dob: ISODate('2007-08-05T00:00:00.000Z'),
8   gender: 'male',
9   personalNumber: '+39 3483861896',
10  email: 'iruberto@virgilio.it',
11  username: 'Iannelli'
12 }
```

Listing 4.2: Example of a Doctor Document

```

1 {
2   _id: ObjectId('684adad437804916ca65ba64'),
3   name: 'Barbara Montagnini',
4   email: 'barbara.montagnini@libero.it',
5   username: 'barbara_montagnini',
6   password: '5e884898da28047151d0e56f8dc6292773603d0d6aabbdd62a11ef721d1
7     542d8',
8   address: {
9     street: 'Via Mura di San Teonisto 8,',
10    city: 'Treviso',
11    province: 'TV',
12    country: 'IT',
13    postalCode: '31100'
14  },
15  phoneNumbers: [ '338 439 2066' ],
16  specializations: [ 'Psicoterapia', 'Psicologia Clinica', 'Psicologia
17    del Lavoro' ],
18  services: [
19    { service: 'Colloquio psicologico', price: 70 },
20  ]
21 }
```



```

18     { service: 'Psicoterapia', price: 75 }
19 ],
20 endorsementCount: 24,
21 reviews: [
22     {
23         patientId: ObjectId('684ada4537804916ca644b76'),
24         name: 'Agostino Omma-Depero',
25         text: 'La Dott.ssa mi ha fatto sentire molto a mio agio.\n' +
26             'E mi ha spiegato in dettaglio le dinamiche da cui
27             derivava il mio problema',
28         date: ISODate('2025-05-25T15:36:33.424Z')
29     },
30     {
31         patientId: ObjectId('684ada4537804916ca63fc97'),
32         name: 'Enzio Camuccini',
33         text: 'Mi sono rivolto per un problema di insonnia. Assieme
34             abbiamo cercato di capire l'origine e mi ha fornito
35             strategie pratiche per migliorare la qualita del sonno
36             senza farmaci.',
37         date: ISODate('2025-04-20T15:26:33.424Z')
38     },
39     {
40         patientId: ObjectId('684ada4437804916ca628d94'),
41         name: 'Diana Pacelli',
42         text: 'Ho contattato la d.ssa per stress lavorativo. Mi e'
43             servito a riflettere sulle mie emozioni, in particolare
44             sulla rabbia che non mi permetteva di essere sempre
45             obiettivo. Ero arrivata con un senso di rivalsa verso la
46             mia azienda sono uscita con strumenti utili per la mia vita
47             non solo professionale. SUpEr conisgliata',
48         date: ISODate('2025-04-16T14:40:45.424Z')
49     },
50     {
51         patientId: ObjectId('684ada4537804916ca646840'),
52         name: 'Giuliana Pertini',
53         text: "Completamente inaffidabile per quanto riguarda
54             appuntamenti e orari. Cambia continuamente ora e giorno
55             dell'appuntamento con scuse varie...",
56         date: ISODate('2025-01-31T07:40:41.424Z')
57     }
58 ],
59 reviewsCount: 4,
60 dob: ISODate('2003-03-27T00:00:00.000Z'),
61 fiscal_code: 'MONBAR030327MKOW',
62 orderRegistrationNumber: 'TV-314952',
63 calendarTemplates: [ ObjectId('684ad9f637804916ca611d7a') ]
64 }

```

Listing 4.3: Example of an Appointment Document

```

1 {
2     _id: ObjectId('684adc6837804916ca700ad8'),
3     date: ISODate('2024-03-11T13:52:13.000Z'),
4     doctor: {

```

```

5      _id: ObjectId('684adad437804916ca662b8a'),
6      name: "Serena D'Agostino",
7      address: {
8          street: 'Viale delle Querce 8,',
9          city: 'Castrovillari',
10         province: 'CS',
11         country: 'IT',
12         postalCode: '87012'
13     },
14     email: "serena.d'agostino@yahoo.com"
15 },
16 patient: {
17     _id: ObjectId('684ada4437804916ca6287e5'),
18     name: 'Rocco Zacco',
19     fiscalCode: 'ZACROC100226M1LH',
20     email: 'chigifabrizia@poste.it',
21     gender: 'male'
22 },
23 visitType: 'Visita ginecologica',
24 patientNotes: '',
25 price: 120
26 }

```

Listing 4.4: Example of a Calendar Template Document

```

1 {
2     _id: ObjectId('684ad9f637804916ca620075'),
3     name: 'Standard',
4     slots: {
5         monday: [
6             { start: '08:30', end: '09:00' },
7             { start: '09:00', end: '09:30' },
8             { start: '09:30', end: '10:00' },
9             { start: '10:00', end: '10:30' },
10            { start: '10:30', end: '11:00' },
11            { start: '11:00', end: '11:30' },
12            { start: '11:30', end: '12:00' },
13            { start: '12:00', end: '12:30' }
14        ],
15        wednesday: [
16            { start: '14:30', end: '15:00' },
17            { start: '15:00', end: '15:30' },
18            { start: '15:30', end: '16:00' },
19            { start: '16:00', end: '16:30' },
20            { start: '16:30', end: '17:00' },
21            { start: '17:00', end: '17:30' },
22            { start: '17:30', end: '18:00' },
23            { start: '18:00', end: '18:30' }
24        ],
25        friday: [
26            { start: '10:00', end: '10:30' },
27            { start: '10:30', end: '11:00' },
28            { start: '11:00', end: '11:30' },
29            { start: '11:30', end: '12:00' },
30            { start: '16:00', end: '16:30' },

```

```

31         { start: '16:30', end: '17:00' },
32         { start: '17:00', end: '17:30' },
33         { start: '17:30', end: '18:00' }
34     ],
35 },
36 isDefault: true
37 }

```

4.2 Graph Database - Neo4j




Figure 4.1: GraphDB structure for HealthHub

Our graph database (Figure 4.1) models two main entities:

- User
- Doctor

Each node retains the MongoDB identifier to maintain consistency across databases. Neo4j is used selectively, focusing on relationship patterns and core attributes needed for traversal and inference, while full entity details remain in MongoDB.

The graph database plays a central role in modeling the social structure of the platform and is exploited in two main ways:

1. As a **recommendation engine**, suggesting doctors to users based on graph proximity and shared review patterns.
2. As a **search optimization layer**, improving query efficiency by leveraging graph traversal capabilities.

Due to the lack of geographic data, we approximate location by leveraging graph proximity, under the assumption that users typically visit doctors near them. This enables us to infer locality-aware clusters and improve personalization without explicit spatial attributes.

4.2.1 Node Types

Neo4j stores only a minimal subset of each entity:

- **User** nodes: `name`, `_id`.
- **Doctor** nodes: `name`, `specializations`, `_id`.

This lean representation minimizes redundancy while supporting all graph-based operations.

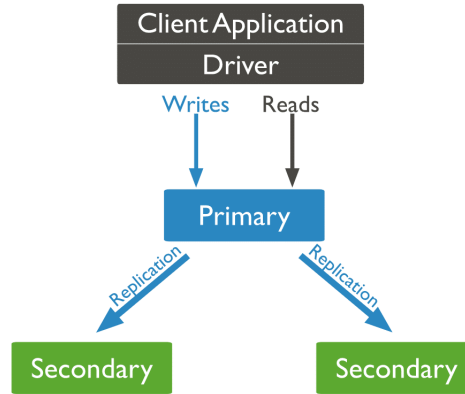


Figure 4.2: MongoDB Replica Set Structure

4.2.2 Relationships

We define two types of directed edges:

- **:REVIEWED** — if user U has reviewed doctor D , then a **:REVIEWED** relationship is created from U to D :

$$U \xrightarrow{\text{:REVIEWED}} D$$

- **:ENDORSED** — if user U endorsed doctor D (e.g., marked them as a favorite or recommended), we store this with a **:ENDORSED** relationship:

$$U \xrightarrow{\text{:ENDORSED}} D$$

These relations enable pattern discovery and inference. Spatial behavior is indirectly modeled through interaction patterns, which often reflect geographic closeness.

4.3 Distributed Database Design

4.3.1 Replica Set Deployment

We were given access to a cluster composed of three nodes. On each node, we deployed a MongoDB instance as part of a replica set. Neo4j was deployed in standalone mode, as replication requires the Enterprise edition, which was not available.

In the MongoDB replica set, only the **primary node** handles write operations. Each write is logged in the primary's *oplog* (operation log), which is then replicated to the **secondary nodes**. These secondaries apply the operations in the same order to maintain a consistent dataset. While all members of the replica set can serve read operations, by default, the application is configured to read from the primary node to preserve strong consistency.

Consistency Level Consistency across replicas is determined by the acknowledgment settings used during read and write operations. Our system adopts a hybrid consistency strategy:

- The **write concern** is set to **w:1**, meaning that a write is acknowledged as soon as it has been applied by the primary node.
- The **read concern** is generally set to **local (r:1)**, allowing reads from the local state of a node, which may not reflect the latest writes.
- However, for critical data such as a doctor's availability schedule, we explicitly enforce **readConcernLevel=majority** to ensure that only the most up-to-date and replicated information is retrieved.

This approach enables us to maintain **eventual consistency** where acceptable — improving performance and reducing latency — while enforcing **strong consistency** only in the specific cases that require it. It provides a good balance between responsiveness and data reliability, especially in a system that emphasizes availability and user experience.

CAP Theorem Considerations Our architecture emphasizes the **Availability (A)** and **Partition Tolerance (P)** aspects of the CAP theorem. The system is designed to:

- Remain accessible and operational even during node failures or network partitions.
- Avoid sacrificing user responsiveness, ensuring that essential features—like searching doctors or leaving endorsements—remain available under most failure scenarios.

4.3.2 Sharding (Design Perspective)

Sharding has not been implemented in the current system but was considered as a viable strategy to scale the document database in case of high demand.

Horizontal Partitioning Strategy Sharding would be employed solely on the MongoDB side, as distributing a graph database like Neo4j across nodes introduces significant traversal overhead that defeats its performance benefits.

We proposed the following sharding design:

- **Doctors Collection:**
 - **Shard Key:** `address.province`
 - **Sharding Mechanism:** hash-based

This approach would evenly distribute documents geographically, avoiding load concentration in specific provinces.

- **Appointments Collection:**
 - **Shard Key:** `appointmentDateTime`
 - **Sharding Mechanism:** range-based

Optimized for time-range queries, this configuration would support efficient access to recent or upcoming appointments.

These design options would ensure balanced workloads and scalable query processing if user and appointment volumes significantly increase in the future.

4.3.3 Handling Inter-database Consistency

Given that we use both a document-oriented database (MongoDB) and a graph database (Neo4j), some redundancy exists across the systems. Certain information—such as user endorsements—must be stored and updated in both databases.

To ensure consistency, we follow a two-phase update pattern:

- A write operation is first performed on MongoDB.
- Only if the MongoDB operation is successful, the corresponding write on Neo4j is triggered.
- If the Neo4j update fails, a rollback is initiated to restore the system to a consistent state.

This logic is managed using Spring’s `@Transactional` annotation, ensuring atomicity across operations. Additionally, asynchronous execution is enabled via `@Async`, allowing long-running updates to Neo4j to execute in separate threads, enhancing system availability.




Figure 4.3: Handling consistency between MongoDB and Neo4j

Eventual Consistency and Responsiveness The system is designed to tolerate brief inconsistencies between MongoDB and Neo4j. This decision improves performance and responsiveness—key aspects for a user-focused social system. Consistency is restored asynchronously and transparently, without impacting the user experience.

Chapter 5

Implementation

5.1 Front-End

The front-end of our application has been developed as a traditional web-based interface, utilizing a combination of **JavaScript**, **HTML** and **CSS**. This technology stack enables the creation of a responsive and interactive user interface that can be accessed from any modern web browser without the need for additional installations. The application communicates asynchronously with the back-end server through **AJAX** (Asynchronous JavaScript and XML) requests, allowing data to be fetched and updated dynamically without requiring full page reloads. This results in a smoother user experience and improved performance, particularly when handling operations such as appointment booking, profile management and real-time data visualization. The separation of concerns between presentation and logic also enhances the maintainability and scalability of the application.

5.2 Front-End Structure

The **resources** directory encapsulates the front-end assets and configuration files essential to the operation and presentation of the web application. Its hierarchical organization reflects a separation of concerns, promoting maintainability and modular development.

Static resources such as CSS and JavaScript files reside in the **static** folder, structured into subdirectories for stylesheets (**css/**), scripts (**js/**) and error pages (**error/**). These support the visual and interactive components of the application.

Dynamic **HTML** views based on **Thymeleaf** are located in the **templates** folder, with dedicated files for both doctors and users. A **fragments** subfolder contains reusable layout components (e.g., headers, footers, sidebars) to ensure consistency and reusability. Application-level configuration is managed through **application.properties**.

5.2.1 Static

The **static** directory collects all static web assets that are delivered to the client without server-side processing. These assets are organized into three primary subfolders: **css/**, **js/**, and **error/**, each serving a distinct role in the client-side rendering process.

The **css** folder contains a series of modular stylesheets dedicated to specific views or components. These include both general-purpose files (e.g. **base.css**, **main.css**, **form.css**) and context-specific styles (e.g. **doctor-dashboard.css**, **user-profile.css**), which define the visual identity and layout of user and doctor interfaces.

The **js** folder hosts JavaScript files that manage client-side interactions and dynamic behavior. Each script is associated with a specific view or functionality—for instance, **login.js** handles authentication interactions, while **slot-calendar.js** governs dynamic calendar updates for appointment scheduling.

The **error** folder includes static HTML pages for standard HTTP error responses (400, 404, 500), ensuring that users receive consistent and informative feedback in case of failure scenarios.

This structure promotes modularity and clarity, allowing for straightforward maintenance and scalability of front-end assets.

5.2.2 Templates

The **templates** directory constitutes a fundamental component within the presentation layer of the web application, as it houses the markup resources responsible for rendering the user interface. This directory is hierarchically organized into subfolders, each reflecting a specific functional domain or interaction paradigm with which the user engages. The HTML documents contained therein are primarily structured using Thymeleaf, a server-side template engine designed for Java-based web environments.

Thymeleaf facilitates the synthesis of dynamic HTML views on the server side, enabling the embedding of context-sensitive data—such as session attributes, query outputs and user-submitted form content—directly into the markup prior to delivery to the client. This server-side rendering paradigm promotes a high degree of interactivity and allows for the generation of customized and context-aware web pages, thereby improving the overall responsiveness and user-centricity of the application interface.

5.3 Back-End

The back-end is implemented using **Java** due to its robustness, platform independence and strong ecosystem. Java offers a well-established set of libraries and frameworks that support rapid development of secure, scalable and maintainable web applications. Its object-oriented nature encourages clean architecture and modular design, which are essential for complex systems like healthcare platforms. Additionally, Java's widespread use in enterprise environments ensures long-term support and community-driven innovation, making it a reliable choice for production-grade applications.

5.3.1 Spring Boot

As the core of our application, we adopted the **Spring Boot** framework, which has significantly streamlined both the development and deployment processes. One of the primary advantages of Spring Boot lies in its built-in support for an embedded **Tomcat** server, allowing the creation of a fully functional web application without requiring external configuration. This design choice has enhanced our ability to construct responsive and maintainable RESTful APIs.

In addition, to support our data management needs, we integrated **Spring Data Neo4j** and **Spring Data MongoDB**. The former facilitates interaction with the **Neo4j graph database** through a consistent repository-based abstraction, ideal for managing highly interconnected medical data such as relationships between doctors and patients. The latter offers seamless integration with **MongoDB, a document-oriented NoSQL database**, while preserving the flexibility and idiomatic programming style of the broader Spring ecosystem. Collectively, these technologies have provided a robust, scalable and modular architecture well-suited to the requirements of our application.

5.4 Back-End Structure

The application follows the layered architecture presented during lectures, organizing the codebase into clearly defined packages to promote modularity and separation of concerns.

5.4.1 package `it.unipi.healthhub.config`

The **Config** package usually contains configuration classes. In our case, we have only one class.

The **FilterConfig** class defines the registration of custom servlet filters used to manage access control within the application. Annotated with **@Configuration**, it registers filters via **FilterRegistrationBean**, associating each filter with specific URL patterns.

In detail:

- **DoctorDashboardAuthFilter**: restricts access to `/doctor/dashboard/*`;
- **PatientDashboardFilter**: applied to `/user/*`;
- **PatientApiFilter**: secures `/api/user/*`;
- **DoctorApiFilter**: secures `/api/doctor/*`;
- **LoginFilter**: manages access to the login page at `/login`.

5.4.2 package `it.unipi.healthhub.controller`

The **Controller** package manages API endpoint mappings and web page controllers for the core sections of the application, including the homepage, login, registration and doctor dashboard. All classes in this package follow the Spring MVC (Model-View-Controller) framework, ensuring a clean separation of concerns between business logic, HTTP request handling and view rendering.

More specifically:

- The **AuthController** handles authentication and registration processes for both patients (users) and doctors. Its responsibilities include rendering login, registration and password recovery views; managing user sessions; creating new user accounts; and processing password reset requests. Utility classes such as **ControllerUtil** and **HashUtil** are employed to promote modularity and code reuse.
- The **DoctorController** manages requests related to viewing a doctor's public profile. It retrieves the corresponding data via the **DoctorService** and injects it into the model for view rendering.
- The **DoctorDashboardController** supports navigation within the doctor's dashboard by mapping each HTTP GET request to a specific section such as appointments, profile, reviews, templates and weekly schedule. It retrieves the authenticated doctor's data from the session and delegates business logic to the **DoctorService**.
- The **HomeController** handles requests to the root endpoints (`/`, `/index`) and to the `/search` page. Each method ensures that session-specific data is consistently added to the model using **ControllerUtil** before returning the appropriate view.
- The **UserController** manages authenticated user interactions through endpoints under `/user`, including profile viewing, appointments and favorite doctors. It uses **ControllerUtil** to populate the model with session-based data. The class adopts a session-aware MVC design aimed at personalized user interaction.

5.4.3 package `it.unipi.healthhub.controller.api`

The `it.unipi.healthhub.controller.api` package defines a set of RESTful controllers that expose the application's core services via HTTP endpoints. These controllers are designed to return data (typically in JSON format) allowing asynchronous communication with frontend components and third-party systems.

Serving as the API layer of the system architecture, this package supports decoupled access to key resources such as users, doctors, appointments and reviews. It complements the standard MVC controllers by enabling machine-oriented interactions, in alignment with modern web application practices.

DoctorAPI

The **DoctorAPI** class defines a RESTful web interface to manage and interact with doctor-related resources within the *HealthHub* system. It is annotated with **@RestController** and is mapped to the base path `/api/doctors`. Below is a summary of the main endpoints:

- **GET `/api/doctors`**
Returns a list of all registered doctors.

- **GET /api/doctors/{id}**
Retrieves a specific doctor by their unique identifier.
- **POST /api/doctors**
Creates a new doctor resource.
- **PUT /api/doctors/{id}**
Updates an existing doctor identified by `id`.
- **DELETE /api/doctors/{id}**
Deletes the doctor with the specified `id`.
- **GET /api/doctors/{doctorId}/services**
Retrieves the list of services offered by the specified doctor.
- **GET /api/doctors/{doctorId}/appointments**
Returns the list of appointments associated with a doctor.
- **POST /api/doctors/{doctorId}/appointments**
Books an appointment for the authenticated patient with the given doctor.
- **GET /api/doctors/{doctorId}/templates**
Retrieves scheduling templates used by the doctor.
- **GET /api/doctors/{doctorId}/schedules/week?year=...&week=...**
Returns the weekly schedule (including prenotable slots) for a doctor, given the ISO week and year.
- **GET /api/doctors/{doctorId}/endorsements**
Retrieves the total number of endorsements for a doctor and whether the current user has endorsed them.
- **POST /api/doctors/{doctorId}/endorse**
Allows a logged-in patient to endorse a doctor.
- **POST /api/doctors/{doctorId}/unendorse**
Removes an existing endorsement made by the patient.
- **GET /api/doctors/{doctorId}/reviews**
Returns all patient reviews for the specified doctor.
- **POST /api/doctors/{doctorId}/reviews**
Submits a review for a doctor, provided the patient has had an appointment with them.
- **GET /api/doctors/{doctorId}/reviews/week/count**
Returns the number of reviews a doctor received during the current ISO week.

PrivateDoctorAPI

The `PrivateDoctorAPI` class implements a secure REST interface for authenticated doctors to manage their personal data, appointments and analytic insights. Annotated with `@RestController` and mapped to `/api/doctor`, it delegates domain logic to `AppointmentService` and `DoctorService`, enforcing session-based access control. Below is a summary of the main endpoints:

- **GET /api/doctor/appointments**
retrieves today's appointments (or for a given date).
- **DELETE /api/doctor/appointments/{id}**
deletes an appointment.
- **PUT /api/doctor/address**
updates the doctor's address.

- **PUT /api/doctor/info**
updates the doctor's profile info.
- **POST /api/doctor/phone**
adds a phone number.
- **DELETE /api/doctor/phone/{index}**
removes a phone number by a doctor.
- **POST /api/doctor/specialization**
adds a specialization.
- **DELETE /api/doctor/specialization/{index}**
removes a specialization.
- **POST /api/doctor/service**
adds a medical service.
- **GET /api/doctor/service**
retrieves all offered services.
- **PUT /api/doctor/service**
updates a medical service.
- **DELETE /api/doctor/service/{id}**
deletes a medical service.
- **POST /api/doctor/template**
creates a new availability template.
- **GET /api/doctor/template**
retrieves all availability templates.
- **PUT /api/doctor/template**
updates a template.
- **DELETE /api/doctor/template/{id}**
deletes a template.
- **PUT /api/doctor/template/default**
sets the default availability template.
- **POST /api/doctor/schedule**
creates a weekly schedule.
- **GET /api/doctor/schedule**
retrieves the weekly schedule.
- **PUT /api/doctor/schedule**
updates the weekly schedule.
- **DELETE /api/doctor/schedule**
deletes the weekly schedule.
- **GET /api/doctor/reviews**
retrieves all received reviews.
- **DELETE /api/doctor/review/{id}**
deletes a review.
- **PUT /api/doctor/password**
changes the password.

- **GET /api/doctor/stats/visits**
gets number of visits (optionally weekly).
- **GET /api/doctor/stats/earnings**
gets annual earnings.
- **GET /api/doctor/stats/patients**
gets number of new patients per month.

UserAPI

The **UserAPI** controller is a RESTful web service component. It provides endpoints for managing user entities within the HealthHub application, leveraging the **UserService** to perform CRUD (Create, Read, Update, Delete) operations on patients stored in a MongoDB database. The controller follows standard REST conventions, utilizing HTTP verbs to represent actions on resources.

- **GET /api/users**
Retrieves a list of all users registered in the system.
- **GET /api/users/{id}**
Fetches a specific user by their unique identifier. Returns HTTP 200 with the user data if found, otherwise HTTP 404.
- **POST /api/users**
Creates a new user entity based on the data provided in the request body. Returns the created user.
- **PUT /api/users/{id}**
Updates an existing user identified by the given ID with new data provided in the request body. Returns HTTP 200 with the updated user if the operation is successful; returns HTTP 404 if the user does not exist.
- **DELETE /api/users/{id}**
Deletes the user identified by the specified ID. Returns HTTP 204 No Content regardless of whether the user existed, indicating that the request has been processed.

PrivateUserAPI

The **PrivateUserAPI** controller is a RESTful web service implemented using the Spring framework in Java, designed to provide secured endpoints for user-specific operations within the HealthHub system. It leverages HTTP sessions to identify the authenticated user (typically a patient) and delegates business logic to the **UserService**. The controller exposes various endpoints to manage user details, contacts, appointments, doctor recommendations and password changes, emphasizing privacy and personalization.

- **GET /api/user/details**
Retrieves detailed personal information of the currently authenticated user by extracting the patient ID from the HTTP session. Returns user details or an error message in case of failure.
- **PUT /api/user/details**
Updates the personal details of the authenticated user. The new data is supplied in the request body as a **UserDetailsDTO** object.
- **GET /api/user/details/view?id={id}**
Retrieves contact information and demographic data for a user specified by the `id` query parameter. Additionally, if a doctor is logged in, it returns the number of visits that doctor has conducted with the user.
- **GET /api/user/contacts**
Fetches the contact information related to the authenticated user.

- **PUT /api/user/contacts**
Updates the contact information of the authenticated user using data provided in the request body as a `PatientContactsDTO`.
- **GET /api/user/appointments/upcoming**
Retrieves a list of upcoming appointments for the authenticated user.
- **GET /api/user/appointments/past**
Retrieves a list of past appointments for the authenticated user.
- **DELETE /api/user/appointments/{appointmentId}**
Cancels the appointment identified by `appointmentId`. Returns HTTP 200 on success or HTTP 400 if the appointment does not exist.
- **GET /api/user/recommendations?limit={limit}**
Returns a list of recommended doctors for the authenticated user, with an optional parameter `limit` to specify the maximum number of results (default is 3).
- **GET /api/user/doctors/reviewed**
Returns a list of doctors that the authenticated user has reviewed.
- **GET /api/user/doctors/endorsed**
Returns a list of doctors endorsed by the authenticated user.
- **PUT /api/user/password**
Allows the authenticated user to change their password. The current and new passwords are supplied in the request body as a `PasswordChangeDTO` and are hashed before validation and update.

SearchAPI

The **SearchAPI** controller is a RESTful component, responsible for executing search queries specifically targeted at medical doctors within the HealthHub system. This controller integrates data from two distinct data sources: a MongoDB database and a Neo4j graph database. The rationale behind this dual-source approach is to combine traditional document-oriented search results with personalized recommendations derived from graph-based relationships.

The primary endpoint exposed by this controller is: **GET /api/search/doctors**. This endpoint accepts a mandatory `query` parameter, representing the search string to be matched against doctor records. Additionally, it accesses the HTTP session to retrieve the identifier of the currently authenticated patient (if available), enabling personalized search refinement.

Upon receiving a valid query, the controller queries the MongoDB database via **DoctorService**, returning results as **DoctorMongoProjection** objects that represent doctor records optimized for search display. If the user is unauthenticated or the patient identifier is absent, the controller returns the top ten doctors from the MongoDB results, sorted by relevance.

For authenticated users, the controller enriches the search outcomes by querying the Neo4j graph database to retrieve personalized recommendations. These recommendations are returned as **DoctorNeo4jProjection** objects, each associated with a relevance score reflecting the strength of the doctor-patient relationship, such as prior interactions or endorsements. In the absence of Neo4j results, the system falls back to the top MongoDB results.

The integration of results from both sources is achieved by mapping doctor identifiers to Neo4j scores, then augmenting each MongoDB doctor's score accordingly. Doctors without Neo4j matches receive a default negative score to denote lack of personalized relevance. The aggregated list is sorted in descending order by combined score and the top ten entries are converted into **DoctorDTO** objects for API response consistency.

Two private helper methods, `toDto` and `toTop10Dto`, support this process by converting raw entities to DTOs and by sorting and limiting the result set. This design ensures that search responses maintain a uniform format, whether or not personalization is applied. Overall, the **SearchAPI** exemplifies a hybrid search strategy that effectively merges traditional document-based retrieval with graph-driven personalization, leveraging session data to optimize relevance and incorporating fallback mechanisms to ensure robust performance for both authenticated and anonymous users.

5.4.4 package `it.unipi.healthhub.dto`

The Data Transfer Objects (DTO) directory encapsulates a set of plain Java objects that serve as data carriers between different layers of the HealthHub application, primarily facilitating communication between the service and presentation tiers. These DTO classes are designed to aggregate and structure data relevant to specific business operations or user interactions, decoupling the internal domain model from external API contracts.

5.4.5 package `it.unipi.healthhub.exception`

The `it.unipi.healthhub.exception` package defines a set of custom unchecked exceptions, each extending `RuntimeException`, to represent specific error conditions encountered in the HealthHub application's domain logic. By encapsulating distinct failure scenarios—such as missing resources or duplicate entries—these exceptions facilitate clear error propagation and centralized handling in higher layers (e.g., controllers or global exception handlers). Moreover, they improve code readability by replacing generic error signals with semantically meaningful types.

- **DoctorNotFoundException:** Thrown when an operation attempts to retrieve or manipulate a doctor entity that does not exist in the system. Constructors allow for default or custom error messages and for nesting an underlying cause;
- **ScheduleAlreadyExistsException:** Raised to indicate an attempt to create a schedule for a doctor in a week for which a schedule is already defined. This prevents conflicting availability entries. As with other exceptions, it supports message customization and cause propagation;
- **UserAlreadyExistsException:** Used when a registration or user-creation process detects that a user with the same unique identifier (e.g., email or fiscal code) is already present in the database. This exception enforces uniqueness constraints at the application layer;
- **UserNotFoundException:** Thrown when a requested user resource cannot be found, for instance during authentication, profile retrieval, or update operations. Its constructors mirror those of the other exception types, enabling flexible error reporting.

5.4.6 package `it.unipi.healthhub.filter`

The `it.unipi.healthhub.filter` package provides servlet filters that centralize authentication and authorization logic for both API and web endpoints. These filters intercept incoming HTTP requests, verify session validity and user roles and either allow processing to continue, redirect to the login page, or return appropriate HTTP error codes upon access violations.

- **DoctorApiFilter:** Ensures that requests to `/api/doctor/*` originate from an authenticated doctor. If no session exists or the session attribute `role` is not `"doctor"`, the filter responds with HTTP 403 Forbidden;
- **DoctorDashboardAuthFilter:** Protects doctor dashboard pages under `/doctor/dashboard/*`. It verifies the presence of a valid session with `role = "doctor"`. If validation fails, the user is redirected to the login page;
- **LoginFilter:** Prevents authenticated users from re-accessing the login page. For existing sessions, GET requests to the login URL are redirected to `/index` (the home page), while POST attempts are rejected with HTTP 403 Forbidden;
- **PatientApiFilter:** Applies to `/api/user/*`, allowing read-only GET requests for all users but restricting state-changing methods (POST, PUT, DELETE) exclusively to authenticated patients (session attribute `role = "patient"`). Unauthorized modification attempts yield HTTP 403 Forbidden;
- **PatientDashboardFilter:** Secures patient dashboard views under `/user/*`. It checks for an active session with `role = "patient"` and redirects unauthenticated or improperly privileged users to the login page.

5.4.7 package `it.unipi.healthhub.model`

The **Model** package defines the entity classes corresponding to the database schema.

The `it.unipi.healthhub.model.mongo` package contains the core domain entities persisted in MongoDB. These classes represent the primary data models for the application and include:

- **Address.java**: encapsulates patient or doctor address fields (street, city, province, postal code and country);
- **Appointment.java**: models appointment records, linking patients and doctors with date, time and status;
- **CalendarTemplate.java**: defines named availability templates composed of slot collections;
- **Doctor.java**: represents doctor profiles, including personal details, specializations and business rules;
- **PrenotableSlot.java**: denotes individual time slots available for booking;
- **Review.java**: captures patient-submitted reviews;
- **Schedule.java**: organizes weekly schedules as mappings from days to prenotable slots;
- **Service.java**: specifies medical services offered by a doctor, including descriptions and pricing;
- **Slot.java**: a simple time interval used within templates and schedules;
- **User.java**: models patient accounts with personal information and secure credentials.

In addition, the `it.unipi.healthhub.model.neo4j` subpackage defines lightweight DAO classes (**DoctorDAO.java**, **UserDAO.java**) used exclusively for graph queries. Since all entity attributes required by the application are already captured in the MongoDB models and relationships are managed via Neo4j projections, the MongoDB classes serve as the definitive domain models.

Their structural design has already been detailed in the Design chapter's UML class diagram section.

5.4.8 package `it.unipi.healthhub.model.mongo`

User

```

1 @Document(collection = "users")
2 public class User {
3     @Id
4     protected String id;
5
6     protected String username;
7
8     protected String name;
9     protected String password;
10    protected String fiscalCode;
11    protected LocalDate dob;
12    protected String gender;
13    private String personalNumber;
14    protected String email;
15
16    public String getId() {
17        return id;
18    }
19    public void setId(String id) { this.id = id; }
20    public String getName() {
21        return name;
22    }
23    public void setName(String name) {
24        this.name = name;
25    }

```

```

26 public String getUsername() {
27     return username;
28 }
29 public void setUsername(String username) {
30     this.username = username;
31 }
32
33 public String getFiscalCode(){ return fiscalCode; }
34 public void setFiscalCode(String fiscalCode){ this.fiscalCode = fiscalCode; }
35
36 public String getGender() {
37     return gender;
38 }
39 public void setGender(String gender) {
40     this.gender = gender;
41 }
42
43 public String getEmail() {
44     return email;
45 }
46 public void setEmail(String email) {
47     this.email = email;
48 }
49
50 public String getPassword() {
51     return password;
52 }
53 public void setPassword(String password) {
54     this.password = password;
55 }
56
57 public void setDob(LocalDate dob) {
58     this.dob = dob;
59 }
60 public LocalDate getDob() {
61     return dob;
62 }
63
64 public String getPersonalNumber() {
65     return personalNumber;
66 }
67
68 public void setPersonalNumber(String personalNumber) {
69     this.personalNumber = personalNumber;
70 }
71
72 public String toString() {
73     return "User{" +
74         ", username='" + username + '\'' +
75         ", password='" + password + '\'' +
76         ", dob=" + dob + "}";
77 }
78 }

```

```

1 @Document(collection = "doctors")
2 public class Doctor extends User {
3     protected String orderRegistrationNumber;
4     private List<Service> services;
5     private int endorsementCount;
6     private int reviewCount;
7     private List<Review> reviews;
8     private List<Schedule> schedules;
9     private List<String> calendarTemplates;
10    private List<String> specializations;

```



```

11 private List<String> phoneNumbers;
12 private Address address;
13
14 public Doctor() {
15     super();
16     services = new ArrayList<>();
17     reviews = new ArrayList<>();
18     schedules = new ArrayList<>();
19     calendarTemplates = new ArrayList<>();
20     specializations = new ArrayList<>();
21     phoneNumbers = new ArrayList<>();
22 }
23
24 public String getOrderRegistrationNumber() { return orderRegistrationNumber; }
25 public void setOrderRegistrationNumber(String orderRegistrationNumber) { this.
    orderRegistrationNumber = orderRegistrationNumber; }
26
27 public List<Service> getServices() {
28     return services;
29 }
30 public void setServices(List<Service> services) {
31     this.services = services;
32 }
33
34 public int getEndorsementCount() {
35     return endorsementCount;
36 }
37 public void setEndorsementCount(int endorsementCount) {
38     this.endorsementCount = endorsementCount;
39 }
40
41 public int getReviewCount() {
42     return reviewCount;
43 }
44 public void setReviewCount(int reviewCount) { this.reviewCount = reviewCount; }
45
46 public List<Review> getReviews() {
47     return reviews;
48 }
49 public void setReviews(List<Review> reviews) {
50     this.reviews = reviews;
51 }
52
53 public List<Schedule> getSchedules() {
54     return schedules;
55 }
56 public void setSchedules(List<Schedule> schedules) {
57     this.schedules = schedules;
58 }
59
60 public List<String> getCalendarTemplates() {
61     return calendarTemplates;
62 }
63 public void setCalendarTemplates(List<String> calendarTemplates) {
64     this.calendarTemplates = calendarTemplates;
65 }
66
67 public List<String> getSpecializations() {
68     return specializations;
69 }
70 public void setSpecializations(List<String> specializations) {
71     this.specializations = specializations;
72 }
73
74 public List<String> getPhoneNumbers() {
75     return phoneNumbers;
76 }
77 public void setPhoneNumbers(List<String> phoneNumbers) {

```

```

78     this.phoneNumbers = phoneNumbers;
79 }
80
81 public Address getAddress() {
82     return address;
83 }
84 public void setAddress(Address address) {
85     this.address = address;
86 }
87
88 public String getId(){
89     return super.getId();
90 }
91
92 public String toString() {
93     return "User{" +
94         ", username='" + getUsername() + '\',' +
95         ", password='" + getPassword() + '\',' +
96         ", dob=" + dob + "}";
97 }
98 }

```

Appointment

```

1  @Document(collection = "appointments")
2  public class Appointment {
3      @Id
4      private String id;
5      private LocalDateTime date;
6      private DoctorInfo doctor;
7      private PatientInfo patient;
8      private String visitType;
9      private String patientNotes;
10     private double price;
11
12     // Constructors
13     public Appointment() {}
14
15     public Appointment(String id, LocalDateTime appointmentDateTime, DoctorInfo doctorInfo,
16         PatientInfo patientInfo, String visitType, String patientNotes) {
17         this.id = id;
18         this.date = appointmentDateTime;
19         this.doctor = doctorInfo;
20         this.patient = patientInfo;
21         this.visitType = visitType;
22         this.patientNotes = patientNotes;
23     }
24
25     // Getters and Setters
26     public String getId() {
27         return id;
28     }
29
30     public void setId(String id) {
31         this.id = id;
32     }
33
34     public LocalDateTime getDate() {
35         return date;
36     }
37
38     public void setDate(LocalDateTime date) {
39         this.date = date;
40     }
41
42     public DoctorInfo getDoctor() {

```

```
42     return doctor;
43 }
44
45 public void setDoctor(DoctorInfo doctor) {
46     this.doctor = doctor;
47 }
48
49 public PatientInfo getPatient() {
50     return patient;
51 }
52
53 public void setPatient(PatientInfo patient) {
54     this.patient = patient;
55 }
56
57 public String getVisitType() {
58     return visitType;
59 }
60
61 public void setVisitType(String visitType) {
62     this.visitType = visitType;
63 }
64
65 public String getPatientNotes() {
66     return patientNotes;
67 }
68
69 public void setPatientNotes(String patientNotes) {
70     this.patientNotes = patientNotes;
71 }
72
73 public void setPrice(Double price) {
74     this.price = price;
75 }
76
77 public double getPrice() {
78     return price;
79 }
80
81 // Inner class DoctorInfo
82 public static class DoctorInfo {
83     private String id;
84     private String name;
85     private Address address;
86     private String email;
87
88     public DoctorInfo() {}
89
90     public DoctorInfo(String doctorId, String doctorName, Address address, String email) {
91         this.id = doctorId;
92         this.name = doctorName;
93         this.address = address;
94         this.email = email;
95     }
96
97     public String getId() {
98         return id;
99     }
100
101     public void setId(String id) {
102         this.id = id;
103     }
104
105     public String getName() {
106         return name;
107     }
108
109     public void setName(String name) {
```

```
110     this.name = name;
111 }
112
113 public Address getAddress() {
114     return address;
115 }
116
117 public void setAddress(Address address) {
118     this.address = address;
119 }
120
121 public String getEmail() {
122     return email;
123 }
124
125 public void setEmail(String email) {
126     this.email = email;
127 }
128 }
129
130 // Inner class PatientInfo
131 public static class PatientInfo {
132     private String id;
133     private String name;
134     private String email;
135     private String gender;
136
137     public PatientInfo() {}
138
139     public PatientInfo(String patientId, String patientName, String email, String gender) {
140         this.id = patientId;
141         this.name = patientName;
142         this.email = email;
143         this.gender = gender;
144     }
145
146     public String getId() {
147         return id;
148     }
149
150     public void setId(String id) {
151         this.id = id;
152     }
153
154     public String getName() {
155         return name;
156     }
157
158     public void setName(String name) {
159         this.name = name;
160     }
161
162     public String getEmail() {
163         return email;
164     }
165
166     public void setEmail(String email) {
167         this.email = email;
168     }
169
170     public String getGender() {
171         return gender;
172     }
173
174     public void setGender(String gender) {
175         this.gender = gender;
176     }
177 }
```

178 }

Address

```
1 public class Address {
2     private String street;
3     private String city;
4     private String province;
5     private String postalCode;
6     private String country;
7
8     // Constructor
9     public Address(String street, String city, String province, String postalCode, String
        country){
10         this.street = street;
11         this.city = city;
12         this.province = province;
13         this.postalCode = postalCode;
14         this.country = country;
15     }
16
17     // Getters and Setters
18     public String getStreet() {
19         return street;
20     }
21
22     public void setStreet(String street) {
23         this.street = street;
24     }
25
26     public String getCity() {
27         return city;
28     }
29
30     public void setCity(String city) {
31         this.city = city;
32     }
33
34     public String getProvince() {
35         return province;
36     }
37
38     public void setProvince(String province) {
39         this.province = province;
40     }
41
42     public String getPostalCode() {
43         return postalCode;
44     }
45
46     public void setPostalCode(String postalCode) {
47         this.postalCode = postalCode;
48     }
49
50     public String getCountry() {
51         return country;
52     }
53
54     public void setCountry(String country) {
55         this.country = country;
56     }
57
58     public String toString() {
59         return street + ", " + city + ", " + province + ", " + postalCode + ", " + country;
60     }
61 }
```

```

62 public String toShortString() {
63     return street + ", " + city;
64 }
65 }

```

Calendar Template

```

1  @Document(collection = "templates")
2  public class CalendarTemplate {
3      @Id
4      private String id;
5      private String name;
6      private Map<String, List<Slot>> slots;
7      private boolean isDefault;
8
9      // Constructor
10     public CalendarTemplate() {}
11
12     public CalendarTemplate(String id, String name, Map<String, List<Slot>> slots, boolean
13         isDefault) {
14         this.id = id;
15         this.name = name;
16         this.slots = slots;
17         this.isDefault = isDefault;
18     }
19
20     // Getters and Setters
21     public String getId() {
22         return id;
23     }
24
25     public void setId(String id) {
26         this.id = id;
27     }
28
29     public String getName() {
30         return name;
31     }
32
33     public void setName(String name) {
34         this.name = name;
35     }
36
37     public Map<String, List<Slot>> getSlots() {
38         return slots;
39     }
40
41     public void setSlots(Map<String, List<Slot>> slots) {
42         this.slots = slots;
43     }
44
45     public boolean isDefault(){
46         return isDefault;
47     }
48
49     public void setDefault(boolean isDefault) {
50         this.isDefault = isDefault;
51     }
52
53     public String toString() {
54         return "CalendarTemplate{" +
55             "id='" + id + '\'' +
56             ", name='" + name + '\'' +
57             ", slots=" + slots +
58             ", default=" + isDefault +
59             '}';

```

```
59     }  
60  
61 }
```

Prenotable Slot

```
1  public class PrenotableSlot extends Slot{  
2      private boolean taken;  
3  
4      public PrenotableSlot() {  
5          super();  
6      }  
7  
8      public PrenotableSlot(String start, String end, boolean taken) {  
9          super(start, end);  
10         this.taken = taken;  
11     }  
12  
13     public boolean isTaken() {  
14         return taken;  
15     }  
16  
17     public void setTaken(boolean taken) {  
18         this.taken = taken;  
19     }  
20 }
```

Review

```
1  public class Review {  
2      public String name;  
3      public String text;  
4      public LocalDate date;  
5      public String patientId;  
6  
7      public String getName() {  
8          return name;  
9      }  
10     public void setName(String name) {  
11         this.name = name;  
12     }  
13  
14     public String getText() {  
15         return text;  
16     }  
17     public void setText(String text) {  
18         this.text = text;  
19     }  
20  
21     public LocalDate getDate() {  
22         return date;  
23     }  
24     public void setDate(LocalDate date) {  
25         this.date = date;  
26     }  
27  
28     public String getPatientId() {  
29         return patientId;  
30     }  
31     public void setPatientId(String patientId) {  
32         this.patientId = patientId;  
33     }  
34 }
```

Schedule

```
1 public class Schedule {
2     private LocalDate week;
3     private Map<String, List<PrenotableSlot>> slots;
4
5     // Getters and Setters
6     public LocalDate getWeek() {
7         return week;
8     }
9
10    public void setWeek(LocalDate week) {
11        this.week = week;
12    }
13
14    public Map<String, List<PrenotableSlot>> getSlots() {
15        return slots;
16    }
17
18    public void setSlots(Map<String, List<PrenotableSlot>> slots) {
19        this.slots = slots;
20    }
21 }
```

Service

```
1 public class Service {
2     private String service;
3     private double price;
4
5     // Constructors
6     public Service() {}
7     public Service(String service, double price) {
8         this.service = service;
9         this.price = price;
10    }
11
12    // Getters and Setters
13    public String getService() {
14        return service;
15    }
16    public void setService(String service) {
17        this.service = service;
18    }
19
20    public double getPrice() {
21        return price;
22    }
23    public void setPrice(double price) {
24        this.price = price;
25    }
26 }
```

Slot

```
1 public class Slot {
2     protected String start;
3     protected String end;
4
5     // Constructor
6     public Slot() {}
7
8     public Slot(String start, String end) {
9         this.start = start;
```



```

10     this.end = end;
11 }
12
13 // Getters and Setters
14 public String getStart() {
15     return start;
16 }
17
18 public void setStart(String start) {
19     this.start = start;
20 }
21
22 public String getEnd() {
23     return end;
24 }
25
26 public void setEnd(String end) {
27     this.end = end;
28 }
29
30 public String toString() {
31     return "Slot{" +
32         "start='" + start + '\',' +
33         ", end='" + end + '\',' +
34         '}';
35 }
36 }

```

5.4.9 package it.unipi.healthhub.model.neo4j

User DAO

```

1  @Node("User")
2  public class UserDAO {
3      @Id
4      private String id;
5      private String name;
6
7      @Relationship(type = "ENDORSED")
8      private Set<DoctorDAO> endorsedDoctors = new HashSet<>();
9
10     @Relationship(type = "REVIEWED")
11     private Set<DoctorDAO> reviewedDoctors = new HashSet<>();
12
13     public UserDAO() {
14     }
15
16     public UserDAO(String id, String name) {
17         this.id = id;
18         this.name = name;
19     }
20
21     public String getId() {
22         return id;
23     }
24
25     public void setId(String id) {
26         this.id = id;
27     }
28
29     public String getName() {
30         return name;
31     }
32
33     public void setName(String name) {
34         this.name = name;

```

```

35     }
36
37     public Set<DoctorDAO> getEndorsedDoctors() {
38         return endorsedDoctors;
39     }
40
41     public Set<DoctorDAO> getReviewedDoctors() {
42         return reviewedDoctors;
43     }
44 }

```

Doctor DAO

```

1  @Node("Doctor")
2  public class DoctorDAO {
3      @Id
4      private String id;
5
6      private String name;
7      private List<String> specializations;
8
9      public DoctorDAO() {
10         specializations = new ArrayList<>();
11     }
12
13     public DoctorDAO(String id, String name, List<String> specializations) {
14         this.id = id;
15         this.name = name;
16         this.specializations = specializations;
17     }
18
19     public String getId() {
20         return id;
21     }
22
23     public void setId(String id) {
24         this.id = id;
25     }
26
27     public String getName() {
28         return name;
29     }
30
31     public void setName(String name) {
32         this.name = name;
33     }
34
35     public List<String> getSpecializations() {
36         return specializations;
37     }
38
39     public void setSpecializations(List<String> specializations) {
40         this.specializations = specializations;
41     }
42
43     public String toString() {
44         return "DoctorDAO{" +
45             "id='" + id + '\'' +
46             ", name='" + name + '\'' +
47             ", specializations=" + specializations +
48             '}';
49     }
50 }

```

5.4.10 package `it.unipi.healthhub.projection`

The `it.unipi.healthhub.projection` package contains projection classes used to encapsulate doctors and their relevance scores obtained from search queries. Specifically, `DoctorMongoProjection` wraps a MongoDB `Doctor` entity with a score, while `DoctorNeo4jProjection` wraps a Neo4j `DoctorDAO` with a similar score. These projections facilitate merging and ranking results from both databases during doctor search operations, as implemented in the `SearchAPI` controller.

5.4.11 package `it.unipi.healthhub.repository`

The **Repository** package is responsible for managing data access operations by leveraging Spring Data JPA interfaces. This package abstracts the underlying persistence mechanisms, allowing for convenient CRUD operations and custom query definitions on the domain entities.

Detailed examination and implementation of specific queries within these repositories will be thoroughly discussed in the dedicated chapter on *Queries*.

5.4.12 package `it.unipi.healthhub.service`

The **Service** layer encapsulates the core business logic. This package is composed by: `AppointmentService`, `DoctorService` and `UserService`.

Appointment Service

The `AppointmentService` class provides a comprehensive set of business logic functionalities related to appointment management within the system. This service interacts directly with the `AppointmentMongoRepository` to perform CRUD operations on appointment entities stored in MongoDB.

Key aspects of `AppointmentService` include:

- **Data Sanitization:** To prevent operator injection vulnerabilities specific to MongoDB, the service sanitizes input strings by replacing prohibited characters (e.g., '\$') within appointment fields, including nested doctor and patient information such as names, emails and addresses.
- **CRUD Operations:** The service exposes methods to retrieve all appointments, fetch an appointment by its identifier, create new appointments after sanitization, update existing ones and delete appointments by ID.
- **Asynchronous Updates:** Leveraging Spring's `@Async` annotation, the service asynchronously updates related doctor and patient information within appointments, improving responsiveness and decoupling these potentially costly operations.
- **Domain-Specific Queries:** It supports domain-specific queries such as retrieving all appointments for a given doctor on a specified day, facilitating calendar and scheduling functionalities.

Doctor Service

The `DoctorService` class manages business logic related to doctor profiles, endorsements, reviews, specializations, analytics and schedule maintenance. It coordinates between MongoDB (via `doctorMongoRepository`) for persistent doctor data and Neo4j (via `doctorNeo4jRepository` and `userNeo4jRepository`) for relationship data.

Key features of `DoctorService` include:

- **Data Sanitization:** Prevents injection vulnerabilities by sanitizing input strings differently for MongoDB (e.g., replacing '\$') and Neo4j (e.g., replacing '\$' and '.'). Sanitization applies to doctor fields and nested user data.
- **Doctor CRUD and Authentication:** Provides creation, retrieval, update and deletion of doctors with transactional integrity. Includes username/email uniqueness checks and password validation for login and password changes.

- **Reviews Management:** Enables patients to add and remove reviews after verifying appointments, updating both doctor documents and Neo4j relationships.
- **Profile Updates:** Supports login authentication, updating doctor details such as address, personal info and phone numbers, with data sanitization and synchronization between MongoDB and Neo4j.
- **Specializations Handling:** Allows adding and removing specializations stored in MongoDB and mirrored in Neo4j, returning indices for reference.
- **Endorsements:** Manages patient endorsements for doctors by updating counts in MongoDB and maintaining endorsement relationships in Neo4j.
- **Analytics:** Provides visit and earnings statistics aggregated by type, week, month and year, leveraging the appointment repository.
- **Schedule Management:** Maintains up-to-date doctor schedules by removing outdated entries and creating new ones for missing weeks, based on default calendar templates. The process includes progress logging and exception handling for missing templates.
- **Transactional Consistency:** Critical updates are annotated with `@Transactional` to ensure atomic operations.

User Service

The `UserService` class handles all user-related business logic, managing users stored in both MongoDB and Neo4j databases. It ensures data integrity and synchronization across these stores while supporting user authentication, profile management and interactions with appointments and doctors.

Key features of `UserService` include:

- **Data Sanitization:** Prevents injection vulnerabilities by sanitizing input strings differently for MongoDB (e.g., replacing '\$') and Neo4j (e.g., replacing '\$' and '.'). Sanitization applies to user fields and nested doctor data.
- **User CRUD and Authentication:** Provides creation, retrieval, update and deletion of users with transactional integrity. Includes username/email uniqueness checks and password validation for login and password changes.
- **Profile and Contact Management:** Updates and retrieves detailed user profile information and contact data, synchronizing name changes across appointments and Neo4j relationships.
- **Appointment Integration:** Retrieves past and upcoming appointments for users, supports appointment cancellation with schedule slot management and notification via a fake mail sender.
- **Neo4j Relations and Recommendations:** Manages endorsement and review relations between users and doctors in Neo4j, fetching lists of endorsed, reviewed and recommended doctors. Recommendation logic combines personalized and popular doctors.
- **Email Operations:** Sends password reset links using a mocked mail sender component.

5.4.13 package `it.unipi.healthhub.util`

This package contains several utility classes and interfaces used within the HealthHub application to support functionalities like session handling, date manipulation, email sending, password hashing and conversion between DTOs and model objects.

- **ControllerUtil.java**
Provides utility methods for controllers, especially for setting common attributes in the Spring `Model` based on the user session. For example, it adds attributes indicating whether a user is logged in and their role (patient or doctor).

- **DateUtil.java**
Contains static methods for date calculations and conversions. Key methods include:
 - Getting the first day of a specific week in a year.
 - Checking if two dates belong to the same week.
 - Getting the next four Mondays after a given date.
 - Conversion between `java.util.Date` and `java.time.LocalDate`.
- **FakeMailSender.java**
A service implementation of the `MailSenderService` interface that simulates sending emails by printing email content to the console. It handles sending notification emails related to appointment deletions and password reset links, validating email addresses before "sending".
- **HashUtil.java**
Provides a method for hashing passwords using SHA-256. It converts the password string into its hashed hexadecimal representation, throwing a runtime exception if hashing fails.
- **MailSenderService.java**
An interface defining methods for sending various types of emails related to appointment cancellations and password resets. Implemented by classes such as `FakeMailSender`.
- **ScheduleConverter.java**
Contains static methods to convert schedule data between DTO representations and model objects. It supports transforming collections of slots and building a schedule for a given Monday date based on a calendar template.
- **TemplateConverter.java**
Provides static methods for converting template-related slot data between DTOs and model classes, including converting regular slots and "prenotable" (bookable) slots, supporting the workflow of schedule creation and manipulation.

5.4.14 Scheduler Component

The `Scheduler` class is a Spring-managed component responsible for automating weekly maintenance tasks on doctor schedules. Its `scheduleTask()` method executes every Sunday at midnight — a time slot deliberately chosen due to the low system load — invoking two key `DoctorService` operations:

- **cleanOldSchedules():** Removes all schedule entries up to the current date, delegating to the repository's `cleanOldSchedules(Date)` method to purge obsolete calendar data.
- **setupNewSchedules():** Identifies doctors lacking schedules in the next four weeks and for each missing week constructs new `Schedule` objects based on their default `CalendarTemplate`.

By centralizing these operations in a scheduled job, the `Scheduler` ensures that doctor availabilities remain current without manual intervention, leveraging reusable service methods and Spring's scheduling infrastructure to maintain data consistency and system reliability.

Chapter 6

Queries

In this chapter, we present only the complex queries executed by our application, without dwelling on trivial queries.

6.1 Mongo DB

6.1.1 Search Query

This query is run every time a guest user, a patient or a doctor inserts a character in the search bar. It works in the same way for all three actors. Returns a list of doctors based on the entered string. The string can be a name, a specialization, a city, a province or a mix of these.

```
1  @Override
2  public List<DoctorMongoProjection> searchDoctors(String text) {
3      // Match documents using full-text search
4      AggregationOperation match = match(Criteria.where("$text").is(new Document("$search",
5          text)));
6
7      // Add MongoDB's built-in text relevance score
8      AggregationOperation project = context -> new Document("$project",
9          new Document("doctor", "$$ROOT")
10         .append("score", new Document("$meta", "textScore"))
11         );
12
13      // Sort by relevance score
14      AggregationOperation sort = context -> new Document("$sort",
15          new Document("score", new Document("$meta", "textScore"))
16          );
17
18      // Limit to top 250 results
19      AggregationOperation limit = limit(250);
20
21      // Build the aggregation pipeline
22      Aggregation agg = newAggregation(match, project, sort, limit);
23
24      // Execute aggregation
25      AggregationResults<DoctorMongoProjection> results =
26      mongoTemplate.aggregate(agg, "doctors", DoctorMongoProjection.class);
27
28      return results.getMappedResults();
29  }
```

Listing 6.1: Java code for the MongoDB Search Query

```
1  db.doctors.aggregate([
2      {
3          $match: {
```

```

4      $text: { $search: query }
5    }
6  },
7  {
8    $project: {
9      doctor: "$$ROOT",
10     score: { $meta: "textScore" }
11   }
12 },
13 {
14   $sort: {
15     score: { $meta: "textScore" }
16   }
17 },
18 {
19   $limit: 250
20 }
21 });

```

Listing 6.2: Equivalent MongoDB Aggregation Pipeline for the Search Query

The above aggregation pipeline enables a full-text search over the `doctors` collection, returning the most relevant documents according to a weighted scoring scheme. Although the underlying text index will be discussed in detail in the dedicated indexing section, it is important to mention that the fields `name`, `specializations`, `address.city`, and `address.province` are indexed with different weights, thereby influencing the ranking of the search results.

The pipeline is composed of four main stages:

- **\$match**: This stage filters documents that match the full-text search criteria specified by the user input. The search is performed across all indexed fields using MongoDB's `$text` operator;
- **\$project**: At this stage, each document is projected into a new structure where the original document is nested under the `doctor` field. Additionally, a new field `score` is introduced, which represents the text relevance score computed by MongoDB using its internal ranking algorithm based on term frequency and field weights;
- **\$sort**: Documents are then ordered in descending order of their `score`, so that the most relevant matches (according to the text index configuration) are prioritized in the output;
- **\$limit**: Finally, the result set is truncated to a maximum of 250 documents, which serves both as a performance optimization and a practical upper bound for rendering results in the user interface.

6.1.2 Analytics

Earnings

This query is an analytic for the doctor's Dashboard and returns a map of monthly revenues for a given year.

```

1  @Override
2  public Map<String, Double> getEarningsByYearForDoctor(String doctorId, Integer year) {
3
4      MatchOperation matchOperation = Aggregation.match(Criteria.where("doctor.id").is(
5          doctorId)
6          .and("date")
7          .gte(LocalDate.of(year, 1, 1).atStartOfDay())
8          .lt(LocalDate.of(year + 1, 1, 1).atStartOfDay())
9      );
10     ProjectionOperation projectMonth = Aggregation.project().andExpression("month(date)").as(
11         "month").and("price").as("price");
12     GroupOperation groupOperation = Aggregation.group("month").sum("price").as("total");
13     ProjectionOperation projectionOperation = Aggregation.project("total").and("month").
14         previousOperation();

```

```

13     Aggregation aggregation = Aggregation.newAggregation(
14         matchOperation,
15         projectMonth,
16         groupOperation,
17         projectionOperation
18     );
19
20     AggregationResults<DBObject> results = mongoTemplate.aggregate(aggregation, Appointment.
21         class, DBObject.class);
22
23     Map<String, Double> earningsByYear = new HashMap<>();
24     for (DBObject doc : results.getMappedResults()) {
25         Integer month = (Integer) doc.get("month");
26         String monthString = new DateFormatSymbols(Locale.ENGLISH).getMonths()[month-1].
27             toLowerCase();
28         Double total = (doc.get("total") instanceof Number) ? ((Number) doc.get("total")).
29             doubleValue() : null;
30         earningsByYear.put(monthString, total);
31     }
32
33     return earningsByYear;
34 }

```

Listing 6.3: Java code for the Earning analytic

```

1  db.appointments.aggregate([
2  {
3      $match: {
4          "doctor._id": ObjectId(doctorId),
5          date: {
6              $gte: ISODate("${year}-01-01T00:00:00Z"),
7              $lt: ISODate("${year + 1}-01-01T00:00:00Z")
8          }
9      }
10 },
11 {
12     $project: {
13         month: { $month: "$date" },
14         price: 1
15     }
16 },
17 {
18     $group: {
19         _id: "$month",
20         total: { $sum: "$price" }
21     }
22 },
23 {
24     $project: {
25         month: "$_id",
26         total: 1,
27         _id: 0
28     }
29 }
30 ]);

```

Listing 6.4: Equivalent MongoDB Aggregation Pipeline for the Earning analytic

The aggregation pipeline shown above is designed to compute the monthly revenue generated by a specific doctor over the course of a given calendar year. This query supports the analytical needs of the doctor's dashboard, enabling a temporal breakdown of financial performance.

The pipeline consists of the following sequential stages:

- **\$match:** This stage performs a selective filtering of documents within the `appointments` collection. It restricts the result set to only those documents where the embedded `doctor._id` field matches the provided identifier and the `date` field falls within the specified year. The temporal filtering is

implemented using the `$gte` and `$lt` operators to define an inclusive-exclusive date range between the first day of January and the first day of the following year;

- `$project`: Each document is then transformed to retain only two fields: the `price` of the appointment and a computed `month` field, which is extracted from the `date` using MongoDB's `$month` operator. This transformation facilitates the grouping operation in the subsequent stage;
- `$group`: Documents are aggregated by the computed `month` value. For each month, the total revenue is computed by summing the values of the `price` field using the `$sum` accumulator. The grouping key, `_id`, is implicitly set to the month number (ranging from 1 to 12);
- `$project`: In the final projection stage, the output is restructured to explicitly include the `month` and `total` fields, while omitting the default `_id` field. This renders the result set suitable for direct transformation into a month-to-amount mapping.

The result of this pipeline is a flat collection of documents, each representing the total earnings for a specific month. In the application layer, this output is subsequently mapped into a `Map<String, Double>`, where each key corresponds to the English name of a month (e.g., "january", "february"), and each value denotes the corresponding revenue in that month. This processed result can be directly visualized in the dashboard through line charts.

Visit Type Summary

This query is an analytic for the doctor's Dashboard and is designed to produce a statistical summary of the different types of visits conducted by a particular doctor, identified uniquely by their `doctorId`. It serves as a key analytic metric within the doctor's dashboard, providing insight into the distribution of appointment types.

```

1  @Override
2  public Map<String, Integer> getVisitsCountByTypeForDoctor(String doctorId){
3      MatchOperation matchOperation = Aggregation.match(Criteria.where("doctor.id").is(
4          doctorId));
5      GroupOperation groupOperation = Aggregation.group("visitType").count().as("total");
6      ProjectionOperation projectionOperation = Aggregation.project("total").and("visitType").
7          previousOperation();
8
9      Aggregation aggregation = Aggregation.newAggregation(
10         matchOperation,
11         groupOperation,
12         projectionOperation
13     );
14
15     AggregationResults<DBObject> results = mongoTemplate.aggregate(aggregation, Appointment.
16         class, DBObject.class);
17
18     Map<String, Integer> visitsCountByType = new HashMap<>();
19     for (DBObject doc : results.getMappedResults()) {
20         String visitType = (String) doc.get("visitType");
21         Integer total = (Integer) doc.get("total");
22         visitsCountByType.put(visitType, total);
23     }
24
25     return visitsCountByType;
26 }

```

Listing 6.5: Java code for the Visit Type Summary analytic

```

1  db.appointments.aggregate([
2      { $match: { "doctor._id": ObjectId(doctorId) } },
3      { $group: { _id: "$visitType", total: { $sum: 1 } } },
4      { $project: { _id: 0, visitType: "$_id", total: 1 } }
5  ])

```

Listing 6.6: Equivalent MongoDB Aggregation Pipeline for the Visit Type Summary analytic

The pipeline consists of three primary stages:

- **\$match**: This initial stage filters the documents in the `appointments` collection to include only those records where the embedded field `doctor._id` matches the specified `doctorId`. This restriction ensures that subsequent aggregations consider exclusively the appointments associated with the target medical professional;
- **\$group**: After filtering, the documents are grouped according to their `visitType` attribute. The grouping key, represented by `_id`, is assigned the value of the `visitType` field. For each distinct visit type, the pipeline computes the total count of corresponding appointments by incrementing a counter via the `$sum` accumulator initialized to 1 for every document;
- **\$project**: In the final stage, the pipeline reformats the output documents to enhance readability and usability. It excludes the default `_id` field and introduces two explicit fields: `visitType`, mapped from the grouping key `_id`, and `total`, representing the count of appointments for that visit type.

The resulting data structure is a set of key-value pairs, where each key corresponds to a visit type (e.g., consultation, follow-up, diagnostic) and each value denotes the frequency of that visit type performed by the doctor. This aggregation output is subsequently transformed within the application layer into a `Map<String, Integer>`, facilitating efficient access and visualization of the visit type distribution in the user interface, with a pie graph in our case.

New Patients of the Month

This query is aimed at identifying the number of new patients a specific doctor has seen for the first time within a given month and year. From a clinical and administrative perspective, this information is crucial for assessing patient acquisition trends, measuring outreach effectiveness, and planning capacity for onboarding procedures. The logic is implemented through a multi-stage MongoDB aggregation pipeline that systematically filters, groups, and analyzes the data within the `appointments` collection.

```

1  @Override
2  public Integer findNewPatientsVisitedByDoctorInCurrentMonth(String doctorId, Integer year,
3      Integer month){
4      LocalDateTime startOfMonth = LocalDate.of(year, month, 1).atStartOfDay();
5      LocalDateTime endOfMonth = startOfMonth.plusMonths(1);
6
7      // 1. Filter all appointments for this doctor
8      MatchOperation matchDoctor = Aggregation.match(
9          Criteria.where("doctor.id").is(doctorId)
10     );
11
12     // 2. Group by patient ID and compute the date of their first visit
13     GroupOperation groupByPatient = Aggregation.group("patient.id")
14         .min("date").as("firstVisitDate");
15
16     // 3. Keep only those patients whose first visit falls in the current month
17     MatchOperation matchFirstVisitInMonth = Aggregation.match(
18         Criteria.where("firstVisitDate").gte(startOfMonth).lt(endOfMonth)
19     );
20
21     // 4. Count how many unique new patients remain
22     CountOperation countNewPatients = Aggregation.count().as("newPatientsCount");
23
24     Aggregation aggregation = Aggregation.newAggregation(
25         matchDoctor,
26         groupByPatient,
27         matchFirstVisitInMonth,
28         countNewPatients
29     );
30
31     // 5. Execute the aggregation pipeline
32     AggregationResults<org.bson.Document> results =
33         mongoTemplate.aggregate(aggregation, Appointment.class, org.bson.Document.class);

```

```

33
34 List<org.bson.Document> mappedResults = results.getMappedResults();
35 if(mappedResults.isEmpty()){
36     return 0;
37 }
38 else{
39     return mappedResults.get(0).getInteger("newPatientsCount", 0);
40 }
41 }

```

Listing 6.7: Java code for the New Patients of the Month analytic

```

1 db.appointments.aggregate([
2     {
3         $match: {
4             "doctor._id": doctorId
5         }
6     },
7     {
8         $group: {
9             _id: "$patient._id",
10            firstVisitDate: { $min: "$date" }
11        }
12    },
13    {
14        $match: {
15            firstVisitDate: {
16                $gte: startOfMonth,
17                $lt: endOfMonth
18            }
19        }
20    },
21    {
22        $count: "newPatientsCount"
23    }
24 ]);

```

Listing 6.8: Equivalent MongoDB Aggregation Pipeline for the New Patients of the Month analytic

The pipeline performs the following operations:

- **\$match**: This initial stage filters the documents to include only those appointments where the embedded field `doctor._id` matches the specified `doctorId`. This ensures the analysis is scoped exclusively to the appointments associated with the targeted physician;
- **\$group**: At this stage, documents are grouped by the unique identifier of each patient, i.e., `patient._id`. For each patient, the aggregation computes the earliest recorded visit date using the accumulator `$min: "$date"`. The resulting field, `firstVisitDate`, represents the patient's first interaction with the doctor in question;
- **\$match** (second occurrence): This filtering stage retains only those patient groups whose `firstVisitDate` falls within the specified month and year. This is achieved by comparing `firstVisitDate` to a temporal interval delimited by `startOfMonth` (inclusive) and `endOfMonth` (exclusive), both computed using Java's `LocalDateTime` API. Patients whose initial interaction with the doctor predates the selected month are excluded from further consideration;
- **\$count**: The final stage performs a scalar aggregation to compute the total number of new patients that passed all previous filters. The result is a single document containing the field `newPatientsCount`, which holds the computed cardinality.

In the application layer, the result of the aggregation is retrieved using the Spring Data MongoDB `mongoTemplate.aggregate` method. If no patients match the criteria, the method safely returns zero; otherwise, it extracts the value associated with the `newPatientsCount` field.

From a technical standpoint, this aggregation ensures uniqueness at the patient level and temporal precision by leveraging both document-level filtering and grouped date evaluation. It is efficient, interpretable, and aligned with the goal of producing actionable monthly cohort data for clinical decision-making.

Number of Visits for a Week

This query returns a count of the number of visits for each day of the week (Sunday through Saturday) for a given doctor, in a specific week and year. It is designed to provide a temporal distribution of clinical activity for a specific doctor over the course of a given week, identified by a year and ISO week number.

```

1  @Override
2  public Map<String, Integer> getVisitsCountByDayForDoctorWeek(String doctorId, Integer week
    , Integer year) {
3      LocalDateTime startOfWeek = DateUtil.getFirstDayOfWeek(week, year).atStartOfDay();
4      LocalDateTime endOfWeek = startOfWeek.plusDays(7);
5
6      MatchOperation matchOperation = Aggregation.match(Criteria.where("doctor.id").is(
    doctorId)
7      .and("date")
8      .gte(startOfWeek)
9      .lt(endOfWeek)
10     );
11     ProjectionOperation projectDay = Aggregation.project().andExpression("dayOfWeek(date)").
    as("day");
12     GroupOperation groupOperation = Aggregation.group("day").count().as("total");
13
14     Aggregation aggregation = Aggregation.newAggregation(
15     matchOperation,
16     projectDay,
17     groupOperation
18     );
19
20     AggregationResults<DBObject> results = mongoTemplate.aggregate(aggregation, Appointment.
    class, DBObject.class);
21
22     Map<String, Integer> visitsCountByDay = new HashMap<>();
23     for (DBObject doc : results.getMappedResults()) {
24         Integer day = (Integer) doc.get("_id");
25         String dayString = new DateFormatSymbols(Locale.ENGLISH).getWeekdays()[day].
    toLowerCase();
26         Integer total = (Integer) doc.get("total");
27         visitsCountByDay.put(dayString, total);
28     }
29
30     return visitsCountByDay;
31 }

```

Listing 6.9: Java code for the Number of Visits in a Week analytic

```

1  db.appointments.aggregate([
2  {
3      $match: {
4          "doctor._id": doctorId,
5          date: {
6              $gte: startOfWeek,
7              $lt: endOfWeek
8          }
9      }
10 },
11 {
12     $project: {
13         day: { $dayOfWeek: "$date" }
14     }
15 },
16 {
17     $group: {

```

```

18     _id: "$day",
19     total: { $sum: 1 }
20   }
21 },
22 {
23   $project: {
24     _id: 0,
25     day: "$_id",
26     total: 1
27   }
28 }
29 ])
```

Listing 6.10: Equivalent MongoDB Aggregation Pipeline for the Number of Visits in a Week analytic

The query follows a multi-stage aggregation process applied to the `appointments` collection, and operates as follows:

- **\$match:** In the first stage, the aggregation filters documents to include only those appointments that match two criteria: the embedded field `doctor._id` equals the specified `doctorId` AND the `date` field falls within the temporal window delimited by `startOfWeek` (inclusive) and `endOfWeek` (exclusive);
- **\$project:** The second stage projects a derived field named `day`, which is computed via the MongoDB expression `$dayOfWeek`. This operator returns an integer in the range `[1, 7]`, representing the day of the week (where 1 corresponds to Sunday and 7 to Saturday), extracted from the `date` field of each matching document;
- **\$group:** In this stage, the documents are grouped according to the previously computed `day` field. For each day of the week, the total number of appointments is calculated using the accumulator `$sum: 1`, effectively counting the number of occurrences per day;
- **\$project:** The final stage restructures the output by renaming the grouping key `_id` to the more semantically meaningful field `day`, while removing the default `_id` field from the output. The resulting documents contain only two fields: `day` (an integer representing the weekday) and `total` (the count of appointments).

On the application side, the output of this aggregation is mapped into a `Map<String, Integer>`, where each key corresponds to the English name of the day of the week (converted to lowercase), and each value indicates the number of visits on that day. The mapping leverages the `DateFormatSymbols` utility class to convert day indices into human-readable weekday names, ensuring correct localization.

6.2 Neo4j

6.2.1 Recommendation Queries

First Query (Collaborative Filtering)

This method recommends a list of doctors to a specific patient (`userId`) based on a collaborative filtering algorithm implemented in Neo4j via Cypher. The goal is to suggest doctors that “similar” users have liked, but that the target user has not yet rated or endorsed.

```

1  /**
2  * Recommends doctors to a specific user based on collaborative filtering:
3  * 1) Finds "similar" users who have at least 3 doctors in common via endorsement/review.
4  * 2) Gets their favorite doctors, excluding those already endorsed/reviewed by the target
   user.
5  * 3) Sorts the doctors by aggregated score and returns the top-N results.
6  *
7  * Complexity: O(E_shared + E_rec), where E_shared is the number of shared relationships
8  * and E_rec is the number of endorsement/review relationships of the similar users.
9  * The query is parameterized and uses indexes.
10 *
```

```

11  * @param limit Maximum number of doctors to recommend.
12  * @return List of recommended doctors with their details.
13  */
14  @Override
15  public List<DoctorDAO> recommendDoctorsForUser(String userId, int limit){
16      try(Session session = driver.session()){
17          return session.readTransaction(tx -> {
18              String cypher =
19                  "MATCH (me:User {id:$uid})-[:ENDORSED|REVIEWED]->(d:Doctor)<-[:ENDORSED|REVIEWED]-(
20                      other:User) " +
21                  "WITH me, other, count(d) AS sharedCount " +
22                  "WHERE sharedCount >= 3 AND me <> other " +
23                  "LIMIT 250 " +
24                  // Get other doctors from 'other' that 'me' has not yet endorsed/reviewed
25                  "MATCH (other)-[r:ENDORSED|REVIEWED*1..3]->(rec:Doctor) " +
26                  "WHERE NOT (me)-[:ENDORSED|REVIEWED]->(rec) " +
27                  // Aggregate a score: more relationships and more similar users increase the score
28                  "WITH rec, sum(sharedCount) AS score, count(r) AS endorsements " +
29                  "ORDER BY score DESC, endorsements DESC " +
30                  "RETURN rec.id AS id, " +
31                  "rec.name AS name, " +
32                  "rec.specializations AS specializations " +
33                  "LIMIT $limit";
34
35              Result result = tx.run(cypher,
36                  Values.parameters("uid", userId, "limit", limit));
37              List<DoctorDAO> out = new ArrayList<>(limit);
38              while (result.hasNext()) {
39                  Record rec = result.next();
40                  out.add(new DoctorDAO(
41                      rec.get("id").asString(),
42                      rec.get("name").asString(),
43                      rec.get("specializations").asList(Value::asString)
44                  ));
45              }
46              return out;
47          });
48      }

```

Listing 6.11: Java code for the Recommendations with Collaborative Filtering

```

1  MATCH (me:User {id: uid})-[:ENDORSED|REVIEWED]->(d:Doctor)<-[:ENDORSED|REVIEWED]-(other:
    User)
2  WHERE me <> other
3  WITH me, other, count(d) AS sharedCount
4  WHERE sharedCount >= 3
5  LIMIT 250
6
7  MATCH (other)-[r:ENDORSED|REVIEWED*1..3]->(rec:Doctor)
8  WHERE NOT (me)-[:ENDORSED|REVIEWED]->(rec)
9
10 WITH rec, sum(sharedCount) AS score, count(r) AS endorsements
11 ORDER BY score DESC, endorsements DESC
12 RETURN rec.id AS id, rec.name AS name, rec.specializations AS specializations
13 LIMIT limit

```

Listing 6.12: Equivalent Cypher query for Recommendations with Collaborative Filtering

The procedure can be formally decomposed into the following stages:

- **Step 1: Identification of Similar Users.** The query first retrieves all users who share at least three doctors in common with the target user via either `ENDORSED` or `REVIEWED` relationships. This is achieved by performing a bi-directional match of the form:

(me:User)-[:ENDORSED|REVIEWED]->(d:Doctor)<-[:ENDORSED|REVIEWED]-(other:User)

followed by an aggregation using `count(d)` to compute the number of shared doctors. The threshold `sharedCount ≥ 3` ensures that only users with a significant overlap are considered similar. A cap of 250 similar users is imposed to ensure query tractability.

- **Step 2: Discovery of Candidate Doctors.** For each similar user, the query searches for additional doctors to whom they are connected through endorsement or review paths up to a length of 3 hops:

`(other)-[r:ENDORSED|REVIEWED*1..3]->(rec:Doctor)`

The constraint `WHERE NOT (me)-[:ENDORSED|REVIEWED]->(rec)` filters out doctors already known (and evaluated) by the target user, ensuring novelty in the recommendations.

- **Step 3: Scoring and Ranking.** Each candidate doctor `rec` is scored based on two aggregated metrics:

1. `score = sum(sharedCount)`: Reflects the cumulative similarity strength from all similar users who have recommended this doctor.
2. `endorsements = count(r)`: Captures the total number of paths (endorsements/reviews) that link similar users to the candidate doctor, indicating popularity.

The final list is sorted in descending order of `score`, and in case of tie, by `endorsements`. The top- N doctors are then returned as personalized recommendations.

From a computational perspective, the algorithm has a complexity of $\mathcal{O}(E_{\text{shared}} + E_{\text{rec}})$, where E_{shared} denotes the number of shared relationships used to determine similarity, and E_{rec} denotes the number of endorsement or review relationships used to evaluate and score recommendations.

This approach effectively exploits Neo4j's native graph traversal capabilities, using path-based semantics to identify non-trivial, yet meaningful associations. It balances locality (similar users) and popularity (frequent endorsements), aligning well with established practices in collaborative filtering literature.

Second Query (Most Popular)

If the previous query does not return an acceptable number of results, this query is used which simply returns the most popular doctors and adds them to the result.

```

1  /**
2  * Recommends popular doctors based on overall endorsements or reviews.
3  * Retrieves doctors ordered by their popularity score, which reflects the total number
4  * of endorsements or reviews received from all users.
5  *
6  * @param limit Maximum number of popular doctors to recommend.
7  * @return List of popular doctors with their details.
8  */
9  @Override
10 public List<DoctorDAO> recommendPopularDoctors(int limit){
11     try(Session session = driver.session()){
12         return session.readTransaction(tx -> {
13             Result result = tx.run(
14                 // Match all doctors with non-null specializations
15                 "MATCH (d:Doctor) WHERE d.specializations IS NOT NULL " +
16                 // Optionally match endorsements or reviews to compute popularity
17                 "OPTIONAL MATCH (u:User)-[r:ENDORSED|REVIEWED]->(d) " +
18                 "WITH d, count(r) AS popularityScore " +
19                 // Return popular doctors ordered by popularity
20                 "RETURN d.id AS id, d.name AS name, d.specializations AS specializations " +
21                 "ORDER BY popularityScore DESC " +
22                 "LIMIT $limit",
23                 Values.parameters("limit", limit)
24             );
25             List<DoctorDAO> doctors = new ArrayList<>();
26             while(result.hasNext()){
27                 Record record = result.next();
28                 Value specValue = record.get("specializations");

```

```

29         List<String> specializations = specValue.isNull()
30         ? new ArrayList<>()
31         : specValue.asList(Value::asString);
32         doctors.add(new DoctorDAO(
33             record.get("id").asString(),
34             record.get("name").asString(),
35             specializations
36         ));
37     }
38     return doctors;
39 });
40 }
41 }

```

Listing 6.13: Java code for the Recommendations with the most popular doctors

```

1  MATCH (d:Doctor)
2  WHERE d.specializations IS NOT NULL
3  OPTIONAL MATCH (u:User)-[r:ENDORSED|REVIEWED]->(d)
4  WITH d, count(r) AS popularityScore
5  RETURN d.id AS id, d.name AS name, d.specializations AS specializations
6  ORDER BY popularityScore DESC
7  LIMIT limit

```

Listing 6.14: Equivalent Cypher query for Recommendations with the most popular doctors

This secondary recommendation query acts as a fallback strategy, intended to ensure the system provides a meaningful output even when the primary collaborative filtering query yields an insufficient number of doctor recommendations.

The query is structured as follows:

- **Step 1: Doctor Selection.** It begins by selecting all nodes of label `Doctor` for which the `specializations` attribute is not null:

```
MATCH (d:Doctor) WHERE d.specializations IS NOT NULL
```

This constraint ensures that only doctors with defined specialization information are considered, likely to increase the relevance and utility of the recommendation.

- **Step 2: Popularity Computation.** An optional match is then performed to collect all `ENDORSED` and `REVIEWED` relationships originating from any `User` and pointing to the selected doctors:

```
OPTIONAL MATCH (u:User)-[r:ENDORSED|REVIEWED]->(d)
```

The keyword `OPTIONAL` is crucial here: it ensures that doctors with zero endorsements or reviews are still included in the result set, albeit with a popularity score of zero. The query then aggregates the total number of such relationships per doctor using:

```
WITH d, count(r) AS popularityScore
```

The resulting `popularityScore` reflects the doctor's visibility or trustworthiness as perceived by the user community.

- **Step 3: Ranking and Limiting.** The doctors are subsequently ordered in descending order of their computed `popularityScore`, and the top-*N* results (as specified by the input parameter `limit`) are returned:

```
ORDER BY popularityScore DESC LIMIT limit
```

The query finally returns the doctor's unique identifier (`id`), name (`name`), and areas of expertise (`specializations`), which are then mapped into `DoctorDAO` objects by the application logic.

From a technical perspective, this query is highly efficient, operating in linear time relative to the number of `ENDORSED` and `REVIEWED` edges in the graph. It is well-suited for cold-start scenarios, such as when a new user lacks sufficient interaction history to enable collaborative filtering. While it lacks personalization, it serves as a robust fallback to preserve the responsiveness and coverage of the recommendation system.

Random Sampling

```

1  @Transactional
2  public List<DoctorDAO> getRecommendedDoctors(String userId, int limit) {
3      List<DoctorDAO> recommendedDoctors = userNeo4jRepository.recommendDoctorsForUser(userId,
4          limit * 10);
5
6      if (recommendedDoctors.size() < limit * 10) {
7          int remaining = limit - recommendedDoctors.size();
8          Set<String> alreadyAddedIds = recommendedDoctors.stream()
9              .map(DoctorDAO::getId)
10             .collect(Collectors.toSet());
11
12         List<DoctorDAO> popularDoctors = userNeo4jRepository.recommendPopularDoctors(remaining
13             * 2);
14         if (!popularDoctors.isEmpty()) {
15             for (DoctorDAO doctor : popularDoctors) {
16                 if (!alreadyAddedIds.contains(doctor.getId())) {
17                     recommendedDoctors.add(doctor);
18                     alreadyAddedIds.add(doctor.getId());
19                 }
20             }
21         }
22
23         // Random sampling
24         Collections.shuffle(recommendedDoctors);
25         return recommendedDoctors.stream()
26             .limit(limit)
27             .collect(Collectors.toList());
28     }

```

Listing 6.15: Java code for the Recommendations

In practice:

- Requires up to $\text{limit} * 10$ personalized recommendations based on similar users (collaborative filtering), to have enough choices before filtering and shuffling;
- If there are few personalized recommendations, it retrieves the most popular doctors and gradually adds them;
- Randomly shuffles (**Random Sampling**) the combined list and returns only the top limits.

6.3 Search Query Support

The `SearchAPI` class implements a RESTful endpoint for doctor search, combining document-based retrieval from MongoDB with graph-based personalization from Neo4j. This hybrid approach enables the system to deliver both relevant and personalized results while maintaining scalability and robustness. This feature is used only if the patient is logged.

This dual-database strategy leverages the strengths of both MongoDB and Neo4j: MongoDB provides efficient document retrieval based on textual relevance, while Neo4j introduces personalized ranking grounded in social and behavioral context. The use of Neo4j complements MongoDB by enabling context-aware recommendations, thus enhancing the overall search experience, especially for returning or registered users.

Neo4j performs a text search of doctors within the patient's social graph (User) based on the connection via *REVIEWED* or *ENDORSED* relations, and returns doctors closest to the user based on the minimum number of "steps" in the graph.

```

1  @Query("MATCH path = (u:User {id:$uid})-[:REVIEWED|ENDORSED*1..3]-(d:Doctor) " +
2  "WHERE toLower(d.name) CONTAINS toLower($search) " +
3  " OR any(spec IN d.specializations WHERE toLower(spec) CONTAINS toLower($search)) " +
4  "WITH d, min(length(path)) AS steps " +
5  "RETURN d as doctor, (5-steps) as score " +

```

```

6  "ORDER BY steps " +
7  "LIMIT 250")
8  List<DoctorNeo4jProjection> findConnectedDoctorsBySteps(@Param("uid") String patientId,
    @Param("search") String search);

```

Listing 6.16: Java code for the Neo4j Search Query

```

1  MATCH path = (u:User {id: uid})-[:REVIEWED|ENDORSED*1..3]-(d:Doctor)
2  WHERE toLower(d.name) CONTAINS toLower(search)
3  OR any(spec IN d.specializations WHERE toLower(spec) CONTAINS toLower(search))
4  WITH d, min(length(path)) AS steps
5  RETURN d AS doctor, (5 - steps) AS score
6  ORDER BY steps
7  LIMIT 250

```

Listing 6.17: Equivalent Cypher query for Neo4j Search Query

The query is structured as follows:

1. **Pattern Matching:** The query initiates a graph traversal by matching paths (`path`) between a given user node (`u:User`)—identified via the parameter `$uid`—and doctor nodes (`d:Doctor`). The traversal is constrained to edges labeled `REVIEWED` or `ENDORSED`, with a maximum traversal depth of 3 hops:

```
MATCH path = (u:User {id: $uid})-[:REVIEWED|ENDORSED*1..3]-(d:Doctor)
```

This clause effectively constructs the patient’s *social neighborhood* within the healthcare interaction network.

2. **Textual Filtering:** The query applies a text-based filter to select only those doctor nodes whose names or specializations partially match the input search string (`$search`). Case-insensitive matching is achieved via the `toLower` function:

```
WHERE toLower(d.name) CONTAINS toLower($search)
OR any(spec IN d.specializations WHERE toLower(spec) CONTAINS toLower($search))
```

This ensures that only semantically relevant doctors are considered, even within the personalized neighborhood.

3. **Proximity-Based Scoring:** For each doctor node, the query computes the shortest path length (in terms of hops) from the patient:

```
WITH d, min(length(path)) AS steps
```

A score is then assigned to each doctor by subtracting the number of steps from a constant (5), producing higher scores for doctors that are more closely connected to the user:

```
RETURN d AS doctor, (5 - steps) AS score
```

4. **Ranking and Limiting Results:** The resulting doctor nodes are sorted by ascending path length (`ORDER BY steps`), prioritizing closer relationships. The result set is capped at 250 entries to prevent computational overload and ensure responsiveness:

```
LIMIT 250
```

From a system architecture perspective, this approach allows the recommendation engine to exploit Neo4j’s strengths in relationship traversal and path evaluation, complementing MongoDB’s capabilities in scalable document storage and retrieval. The fusion of these two models supports a robust, intelligent and user-aware search experience.

6.4 Indexes

6.4.1 MongoDB

MongoDB supports a variety of index types, each designed to optimize query performance for specific use cases. The following sections detail the indexes implemented in our MongoDB collections, which are crucial for enhancing the efficiency of our application's data retrieval operations.

Appointments Collection - `idx_patient_id`

Table 6.1: Performance comparison with and without the `idx_patient_id` index

Metric	With index	Without index	% Change
Execution time (ms)	333.20	892.80	-62.68%
Documents examined (avg)	4,634.60	699,989.00	-99.34%
Index keys examined (avg)	4,634.60	0.00	N/A
Documents returned (avg)	4,634.60	4,634.60	+0.00%

The `idx_patient_id` index significantly improves query performance when retrieving appointments for patients with a high number of visits. The average execution time decreased by more than 60%, and the number of documents examined dropped by over 99%. This demonstrates the importance of indexing frequently queried fields.

Appointments Collection - `idx_doctor_id`

Table 6.2: Performance comparison with and without the `idx_doctor_id` index

Metric	With index	Without index	% Change
Execution time (ms)	13.70	566.80	-97.58%
Documents examined (avg)	1,076.00	699,989.00	-99.85%
Index keys examined (avg)	1,076.00	0.00	N/A
Documents returned (avg)	5.90	5.90	+0.00%

Similarly, the `idx_doctor_id` index is essential for queries involving doctor-related appointment data, such as revenue analytics. The index reduced execution time by nearly 98% and minimized the documents examined from nearly 700,000 to just over 1,000, highlighting a dramatic gain in efficiency.

Doctors Collection - `DoctorsTextIndex`

Table 6.3: Performance comparison between regex-based and text index-based search

Metric	Text Search	Regex Search	% Change
Execution time (ms)	108.29	274.14	-60.50%
Documents examined (avg)	14,081.00	87,632.00	-83.93%
Index keys examined (avg)	31,073.71	0.00	N/A
Documents returned (avg)	14,081.00	9,587.71	+46.87%

We compared regex-based search with full-text indexed search using `DoctorsTextIndex`. The indexed search not only executed over 60% faster but also returned 47% more documents on average, with significantly fewer documents examined. This underlines the advantage of using text indexes for flexible and performant search capabilities (see Appendix A for the original regex-based query).

6.4.2 Neo4j

During the initial setup of the Neo4j database, the insertion of relationships between **User** (patients) and **Doctor** nodes was significantly hindered by the lack of indexing. The script responsible for importing the data failed to complete within a reasonable time. By introducing indexes on the `id` field of both node types—which duplicates the original MongoDB `_id`—relationship creation was drastically accelerated. After indexing, the setup process, which previously stalled indefinitely, was completed in just a few seconds. This confirms the critical role of indexes even in graph databases when handling large-scale relationship insertions.

Appendix A

Original Regex-Based Query

This implementation defines a custom scoring logic using `$regex` and `$cond` operators to prioritize documents based on the match location. Although flexible, it does not leverage indexes and leads to full collection scans, making it inefficient for large datasets.

```
1 public List<DoctorMongoProjection> searchDoctors(String text) {
2     AggregationOperation match = match(new Criteria().orOperator(
3         Criteria.where("name").regex(text, "i"),
4         Criteria.where("specializations").regex(text, "i"),
5         Criteria.where("address.city").regex(text, "i"),
6         Criteria.where("address.province").regex(text, "i")
7     ));
8
9     Document nameCond = new Document("$cond", Arrays.asList(
10         new Document("$regexMatch", new Document("input", "$name").append("regex",
11             text).append("options", "i")),
12         0,
13     ));
14     Document specFilter = new Document("$filter", new Document("input", "
15         $specializations")
16         .append("as", "s")
17         .append("cond", new Document("$regexMatch",
18             new Document("input", "$$s").append("regex", text).append("options",
19                 "i"))));
20
21     Document specCond = new Document("$cond", Arrays.asList(
22         new Document("$gt", Arrays.asList(new Document("$size", specFilter), 0)),
23         1,
24     ));
25     Document cityCond = new Document("$cond", Arrays.asList(
26         new Document("$regexMatch", new Document("input", "$address.city").append("
27             regex", text).append("options", "i")),
28         2,
29     ));
30     Document provCond = new Document("$cond", Arrays.asList(
31         new Document("$regexMatch", new Document("input", "$address.province").
32             append("regex", text).append("options", "i")),
33         3,
34     ));
35     Document scoreExpr = new Document("$min", Arrays.asList(nameCond, specCond, cityCond
36         , provCond));
37
38     AggregationOperation project = ctx -> new Document("$project",
39         new Document("doctor", "$$ROOT")
40             .append("score", scoreExpr)
41     );
```

```

40
41     AggregationOperation sort = sort(Sort.Direction.ASC, "score");
42
43     AggregationOperation limit = limit(250);
44
45     Aggregation agg = newAggregation(match, project, sort, limit);
46
47     AggregationResults<DoctorMongoProjection> results =
48         mongoTemplate.aggregate(agg, "doctors", DoctorMongoProjection.class);
49
50     return results.getMappedResults();
51 }

```

Listing A.1: Java code for searching doctors with regex

```

1 db.doctors.aggregate([
2     {
3         "$match": {
4             "$or": [
5                 {"name": {"$regex": term, "$options": "i"}},
6                 {"specializations": {"$regex": term, "$options": "i"}},
7                 {"address.city": {"$regex": term, "$options": "i"}},
8                 {"address.province": {"$regex": term, "$options": "i"}},
9             ]
10        },
11    },
12    {
13        "$addFields": {
14            "score": {
15                "$min": [
16                    {
17                        "$cond": [
18                            {"$regexMatch": {"input": "$name", "regex": term, "options": "i"}},
19                            0,
20                            5
21                        ],
22                    },
23                    {
24                        "$cond": [
25                            {
26                                "$gt": [
27                                    {
28                                        "$size": {
29                                            "$filter": {
30                                                "input": "$specializations",
31                                                "as": "s",
32                                                "cond": {
33                                                    "$regexMatch": {
34                                                        "input": "$$s",
35                                                        "regex": term,
36                                                        "options": "i"
37                                                    }
38                                                }
39                                            }
40                                        },
41                                        0
42                                    ]
43                                },
44                                1,
45                                5
46                            ]
47                        },
48                    },
49                    {
50                        "$cond": [
51                            {"$regexMatch": {"input": "$address.city", "regex": term, "options": "i"}},

```

```
52         2,  
53         5  
54     ]  
55 },  
56 {  
57     "$cond": [  
58         {"$regexMatch": {"input": "$address.province", "regex": term  
59         , "options": "i"}},  
60         3,  
61         5  
62     ]  
63 }  
64 ]  
65 }  
66 },  
67 {"$sort": {"score": 1}},  
68 {"$limit": 250}  
69 ])
```

Listing A.2: MongoDB aggregation pipeline for searching doctors with regex

Appendix B

User Manual

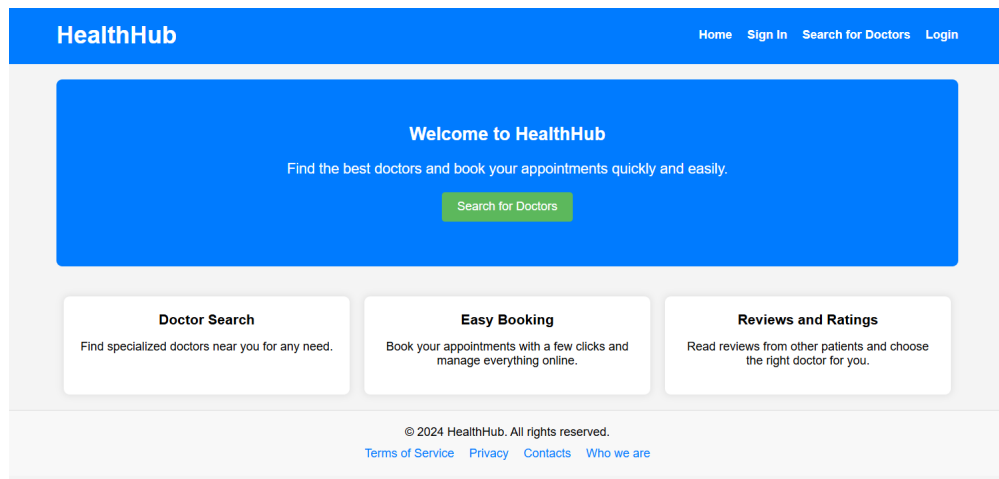



Figure B.1: Home Page of the HealthHub Application

B.1 Patient

B.1.1 Signup and Login Process


From the Home Page, users can initiate the authentication process by selecting the *"Sign In"* option located at the top-right corner of the interface.



The screenshot shows the 'HealthHub' website with a blue header containing 'Home', 'Sign In', 'Search for Doctors', and 'Login'. The main content area features a 'Register as a Patient' form. The form includes fields for Username (filled with 'junior'), Name (filled with 'Giacomo Verdi'), Fiscal Code (filled with 'GCMVRD84D04A703A'), Email (filled with 'giacomo.verdi@gmail.com'), Password (filled with '*****'), Confirm Password (filled with '*****'), Date of Birth (filled with '05/04/1995' and a calendar icon), Gender (a dropdown menu with 'Male' selected), and Personal Phone Number (filled with '0123456789'). A green 'Register' button is at the bottom of the form. Below the button are two links: 'Already have an account? Log in' and 'Are you a doctor? Sign up to use our services'.

Figure B.2: Patient Registration Form

Upon accessing the registration interface, a patient may create a new account. Alternatively, users who already possess valid credentials can proceed to the login form.



The screenshot shows the 'HealthHub' website with a blue header containing 'Home', 'Sign In', 'Search for Doctors', and 'Login'. The main content area features a 'Login' form. The form includes fields for 'Email or Username' (filled with 'junior') and 'Password' (filled with '*****'). A green 'Login' button is at the bottom of the form. Below the button are two links: 'Register' and 'Forgot your password?'. At the bottom of the page, there is a footer with the text '© 2024 HealthHub. All rights reserved.' and four links: 'Terms of Service', 'Privacy', 'Contacts', and 'Who we are'.

Figure B.3: Patient Login Form

The login interface also provides auxiliary navigation: by selecting the *Register* button, users are redirected to the registration form. Additionally, by selecting the *Forgot Password* option, users can initiate the password recovery process. This procedure simulates the dispatch of a recovery link to the registered e-mail address.

The screenshot shows the 'Recover Your Password' form on the HealthHub website. The form is centered on a light gray background. It has a title 'Recover Your Password' and a message: 'A password reset link has been sent to your email address.' Below this, there is a label 'Enter your email address' and a text input field containing 'giacomo.verdi@gmail.com'. A green button labeled 'Send Reset Link' is positioned below the input field. At the bottom of the form, there are two links: 'Back to Login' and 'Register'. The footer of the page includes the copyright notice '© 2024 HealthHub. All rights reserved.' and links for 'Terms of Service', 'Privacy', 'Contacts', and 'Who we are'.

Figure B.4: Password Recovery Request Form

The screenshot shows an email confirmation message. The header indicates the email is being sent to 'giacomo.verdi@gmail.com' with the subject 'HealthHub - Password Reset'. The body of the email starts with 'Hello,' followed by a message: 'We received a request to reset your password. Please click the link below to choose a new password:'. A blue hyperlink is provided: 'http://localhost:8080/forgot-password/6854886fae4d184d32e6981e'. Below the link, it says 'If you did not request a password reset, you can safely ignore this email.' The email is signed off with 'Best regards, The HealthHub Team'. The footer of the email states 'Email sent successfully.'

Figure B.5: Confirmation of Email Dispatch

The screenshot shows the 'Set a New Password' form on the HealthHub website. The form is centered on a light gray background. It has a title 'Set a New Password' and two input fields: 'New Password' and 'Confirm Password'. Both fields have masked characters (dots) and are outlined with a green border. A green button labeled 'Reset Password' is positioned below the 'Confirm Password' field. At the bottom of the form, there is a link 'Back to Login'. The footer of the page includes the copyright notice '© 2024 HealthHub. All rights reserved.' and links for 'Terms of Service', 'Privacy', 'Contacts', and 'Who we are'.

Figure B.6: Password Successfully Updated

Once authenticated, the user is redirected to the personalized Patient Home Page. This dashboard provides access to core functionalities including doctor search, profile management, appointment tracking, and the display of favoredited doctors—specifically those who have been reviewed or endorsed by the user. A recommendation system enhances this section by suggesting relevant doctors based on the patient’s interaction history.

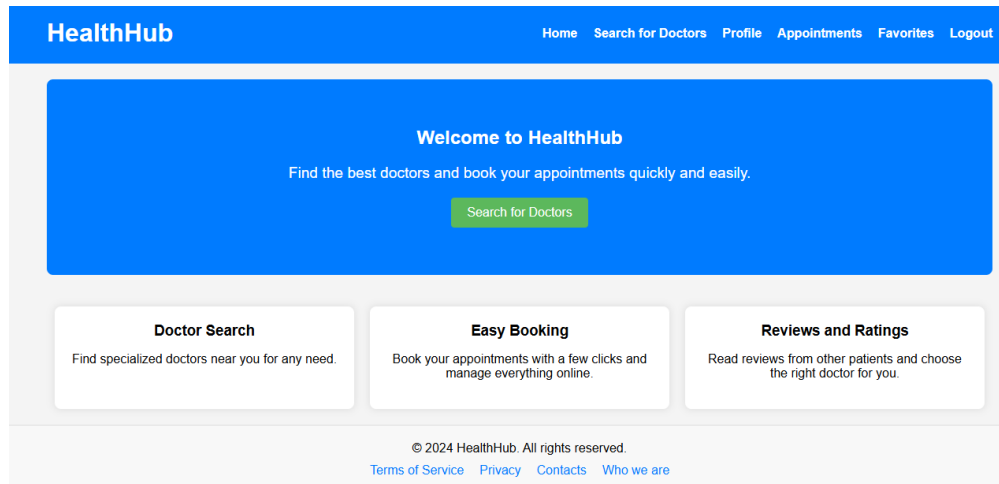


Figure B.7: Patient Dashboard After Login

B.1.2 Doctor Profile and Appointment Booking

By selecting the *"Search for Doctors"* option, users can enter a query string that may consist of a doctor's name, medical specialization, city, province, or a combination of these criteria.




Figure B.8: Doctor Search Interface

After performing a search, the user can select a doctor from the resulting list to access the corresponding public profile page.

The doctor's profile page displays information regarding the available types of visits and the corresponding fees. From this page, the user can book an appointment by selecting a specific visit type. An optional note can be attached to the appointment request to communicate additional information to the doctor.

HealthHub Home Search for Doctors Profile Appointments Favorites Logout

Dr. Gianmarco Torrisi Endorsements: 39
Specializations: Dermatologia e Venereologia, Tricologia, Medicina Estetica
[Endorse](#)

Contact Information
Address: Viale Odorico da Pordenone, 25, Catania, CT, 95128, IT
Phone: 0931 196 2573, 095 951 7591
Email: gianmarco.torrisi@hotmail.it
[Show map](#)

Services and Prices

- Visita dermatologica - €80.0
- Epilazione laser - €50.0
- Laser - €50.0
- Mappatura nevi - €100.0
- Visita di controllo - €45.0
- Visita tricologica - €80.0
- Mappatura nei - €100.0
- Filler acido ialuronico - €300.0
- Filler labbra - €300.0
- Laser CO2 - €100.0
- Patch test - €100.0
- Prima visita dermatologica - €150.0
- Prima visita tricologica - €150.0
- Visita dermatologica + mappatura nei - €150.0
- Visita dermatologica di controllo - €109.0
- Visita venereologica - €150.0

Book an Appointment
Office Visit
Viale Odorico da Pordenone, 25, Catania
Choose the type of visit
Laser CO2
Patient Notes
I might be a little late.
Appointment Date
Jun 16-22, 2025

Mon	Tue	Wed	Thu	Fri	Sat
08:30	14:30	10:00			
09:00	15:00	10:30			
09:30	15:30	11:00			
10:00	16:00	11:30			
10:30	16:30	16:00			
11:00	17:00	16:30			
11:30	17:30	17:00			
12:00	18:00	17:30			

[Book](#)

Figure B.9: Appointment Booking Interface

Furthermore, users have the option to endorse the doctor or leave a review. However, reviews are only permitted if the patient has attended at least one appointment with the doctor in question. This ensures the authenticity and reliability of the feedback provided.

• Laser CO2 - €100.0
• Patch test - €100.0
• Prima visita dermatologica - €150.0
• Prima visita tricologica - €150.0
• Visita dermatologica + mappatura nei - €150.0
• Visita dermatologica di controllo - €109.0
• Visita venereologica - €150.0
• Epiluminescenza - €100.0
• Esame tricologico - €100.0

Reviews

Leave a review...

[Submit Review](#)

Rolando Modiano
2025-06-07
Dottore davvero molto ma molto preparato, attento e disponibile Non posso ancora dire se la cura prescrittami è stata risolutiva perché richiede tempo, ma mi sento di esprimere un parere davvero altamente positivo sul Dott. Re Torrisi.

Concetta Dossi
2025-06-06
Chiara ed esauritiva nelle spiegazioni, approccio preciso e professionale, strumentazioni ottime e studio pulito.

Sig. Delfino Ruberto
2025-06-06
Disponibile ed empatico con la piccola di 3 anni che aveva un po' paura..!

Figure B.10: Doctor Review Interface (Note: Reviewing is allowed only after completing at least one appointment with the doctor)

B.1.3 Appointments

The *Appointments* section can be accessed via the right-hand sidebar. Within this interface, the user's scheduled appointments are displayed and automatically sorted in chronological order, from the nearest to the most distant date.

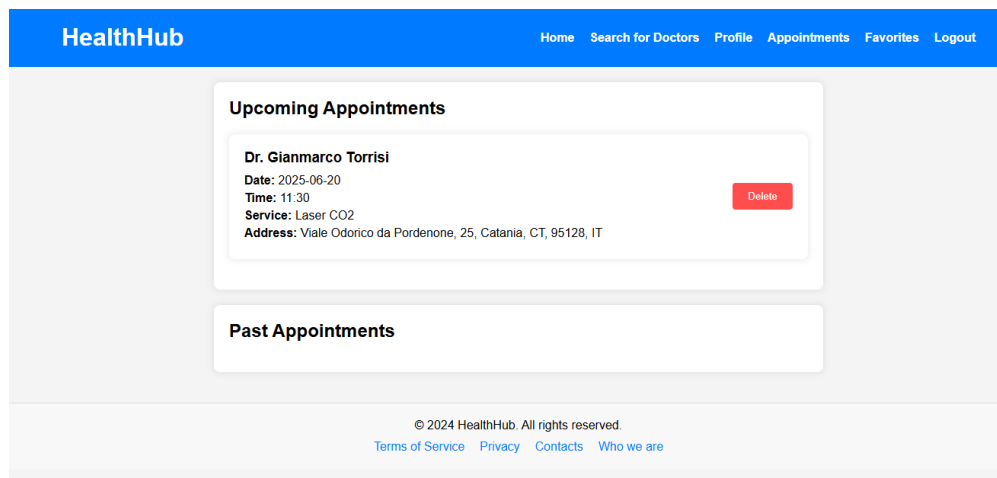


Figure B.11: Appointment Overview Interface

Appointments are dynamically updated based on the current date and time. Once an appointment's scheduled time has passed, it is automatically transitioned from the *Upcoming Appointments* section to the *Past Appointments* section, which is sorted in reverse chronological order.

Only the *Upcoming Appointments* can be deleted.




Figure B.12: Appointment moved from Upcoming Appointments to Past Appointments

B.1.4 Recommendation System

The *Favorites* section is accessible via the right-hand sidebar. This interface displays the list of doctors that the patient has either endorsed or reviewed.

In addition, a list of recommended doctors is presented at the top of the page. These recommendations are generated dynamically through a Neo4j-based graph query, as previously described in the Neo4j section of the *Queries* chapter.

By clicking on the name of any doctor listed, the user can directly navigate to the corresponding public profile page.

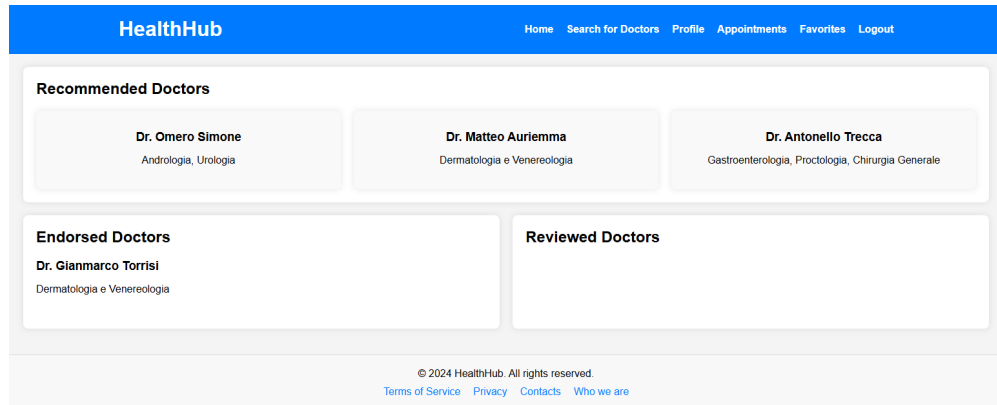


Figure B.13: Favorites and Recommendations Interface

B.1.5 Profile

The user can also access and manage their personal profile through the dedicated *Profile* section. This area allows for the visualization and editing of personal information and contact details.

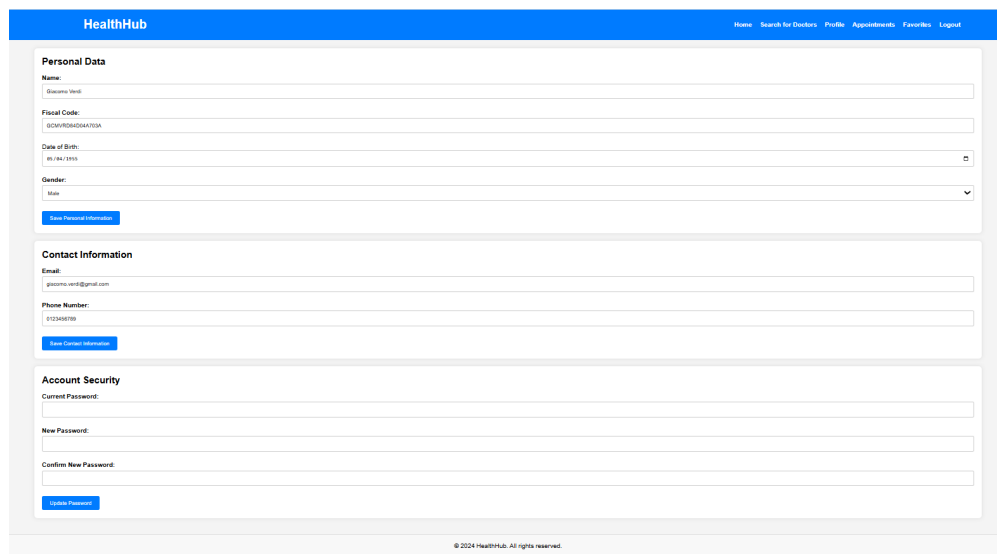


Figure B.14: Patient Profile Management Interface

B.2 Doctor

Doctors have access to a dedicated interface that allows them to manage their appointments, review patient feedback, and edit their personal and professional information. They can also monitor key analytics related to their activity and configure their weekly availability through calendar templates.

B.2.1 Dashboard

The dashboard provides doctors with an overview of their current performance. It includes:

- A line chart displaying total earnings for the current year.
- A pie chart illustrating the distribution of services provided.

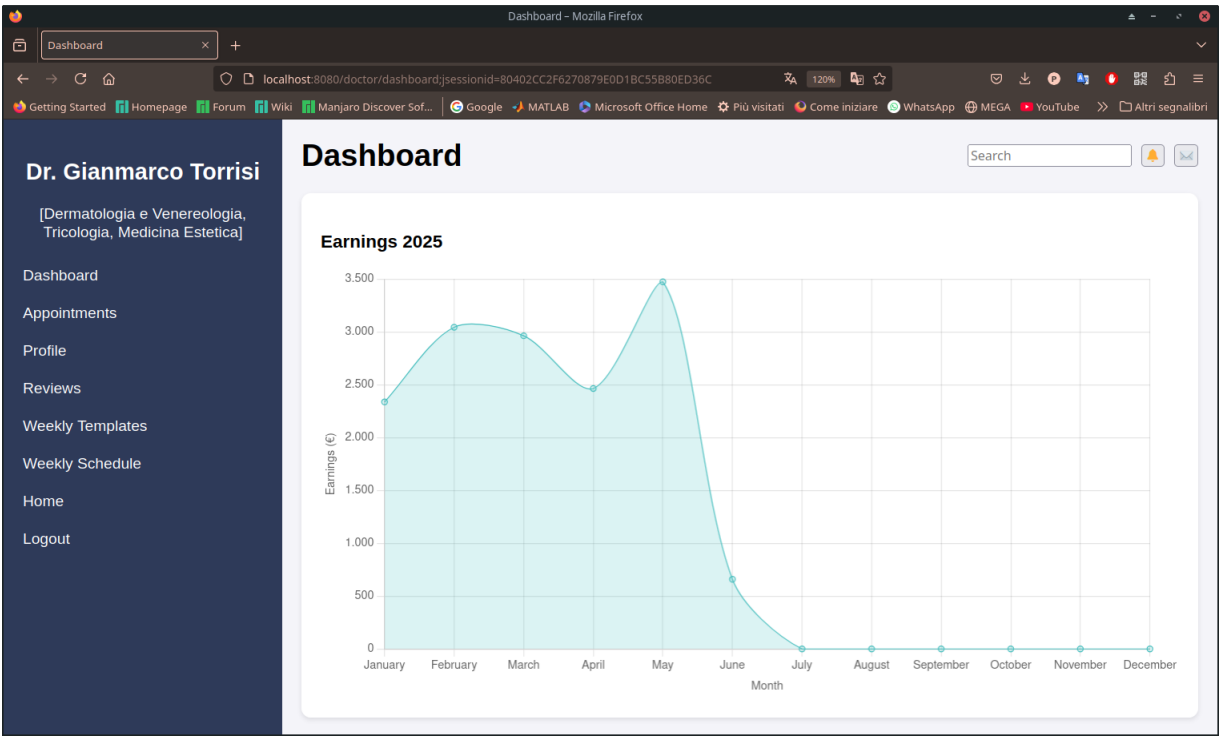


Figure B.15: Line chart displaying earnings by month for the current year.

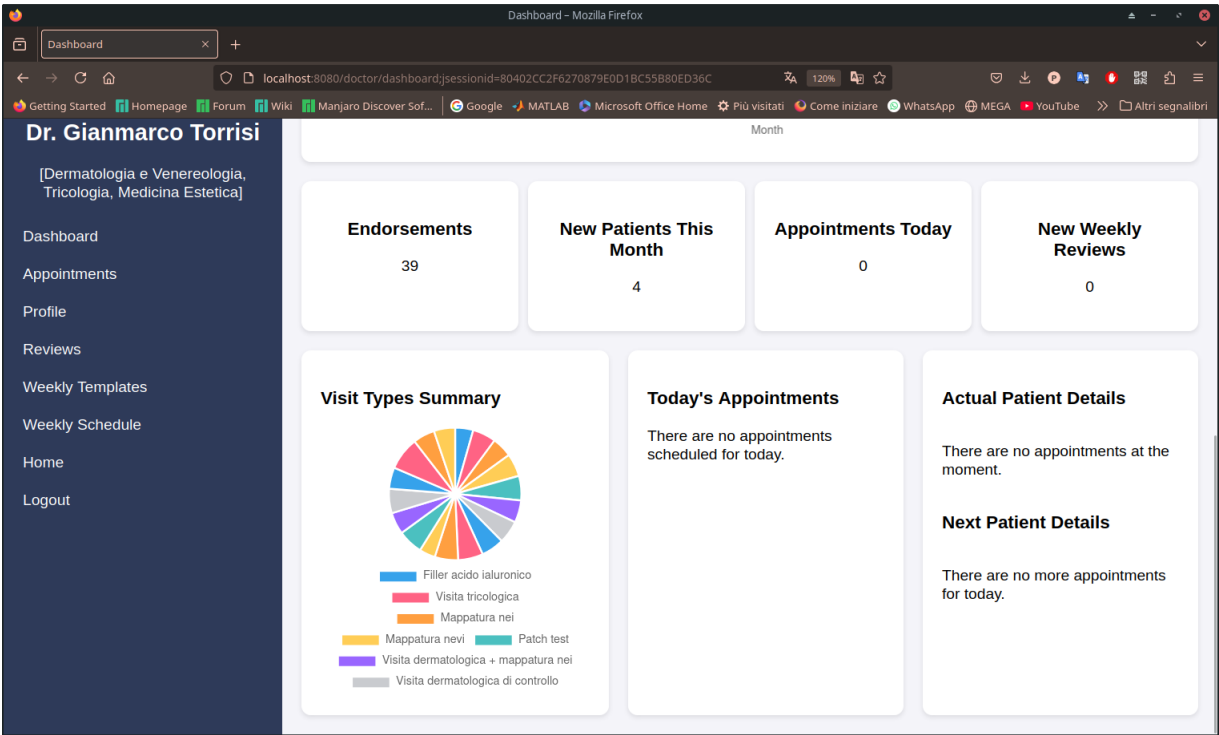


Figure B.16: Pie chart showing the distribution of services provided by the doctor.

- A summary of today's appointments, with access to detailed information.

- A preview of the next scheduled patient.
- Key performance indicators, including:
 - **Endorsements**
 - **New Patients This Month**
 - **Appointments Today**
 - **New Weekly Reviews**

B.2.2 Appointment List

This section allows the doctor to view a comprehensive list of appointments. The list can be filtered by date, and each entry provides access to detailed information or cancellation options.




Figure B.17: Appointment list with filtering options.

B.2.3 Personal Profile

In the profile section, doctors can view and edit their personal details and professional information. Specifically, they can:

- Update their office address, phone numbers, and email address.
- Provide a list of services they offer, including the name and cost of each.
- Add or edit their medical specialties.
- Change their account password.

Dr. Gianmarco Torrissi
[Dermatologia e Venereologia, Tricologia, Medicina Estetica]

Dashboard
Appointments
Profile
Reviews
Weekly Templates
Weekly Schedule
Home
Logout

Doctor Profile

Personal Details

Full Name:
Gianmarco Torrissi

Order Registration Number:
CT-200155

Fiscal Code:
Fiscal Code

Email:
gianmarco.torrissi@hotmail.it

Date of Birth:
18 / 08 / 1980

Gender:
Male

Figure B.18: Doctor's profile page with editable fields.

Dr. Gianmarco Torrissi
[Dermatologia e Venereologia, Tricologia, Medicina Estetica]

Dashboard
Appointments
Profile
Reviews
Weekly Templates
Weekly Schedule
Home
Logout

Office Address

Street:
Viale Odorico da Pordenone, 25

Office City:
Catania

Office Province:
CT

ZIP Code:
95128

State:
IT

Save

Phone Numbers

0931 196 2573 x 095 951 7591 x

Figure B.19: Doctor's office address and contact information.

The screenshot shows the 'Doctor Profile' page for Dr. Gianmarco Torrissi. The left sidebar contains navigation links: Dashboard, Appointments, Profile, Reviews, Weekly Templates, Weekly Schedule, Home, and Logout. The main content area is titled 'Dr. Gianmarco Torrissi' and lists his specialties: [Dermatologia e Venereologia, Tricologia, Medicina Estetica]. Below this, there are sections for 'State' (IT), 'Phone Numbers' (0931 196 2573, 095 951 7591), 'Specializations' (Dermatologia e Venereologia, Tricologia, Medicina Estetica), and 'Types of Visits and Prices'.

Figure B.20: Medical specialties that the doctor can add or edit.

The screenshot shows the 'Doctor Profile' page for Dr. Gianmarco Torrissi, focusing on the 'Types of Visits and Prices' section. The left sidebar is the same as in Figure B.20. The main content area lists the following services and their costs:

Service	Cost	Save	Remove
Visita dermatologica	80.0	Save	Remove
Epilazione laser	50.0	Save	Remove
Laser	50.0	Save	Remove
Mappatura nevi	100.0	Save	Remove
Visita di controllo	45.0	Save	Remove
Visita tricologica	80.0	Save	Remove
Mappatura nei	100.0	Save	Remove

Figure B.21: Services offered by the doctor, including name and cost.

B.2.4 Review Page

Doctors can browse the reviews left by their patients. If necessary, inappropriate or outdated reviews can be deleted directly from this section.

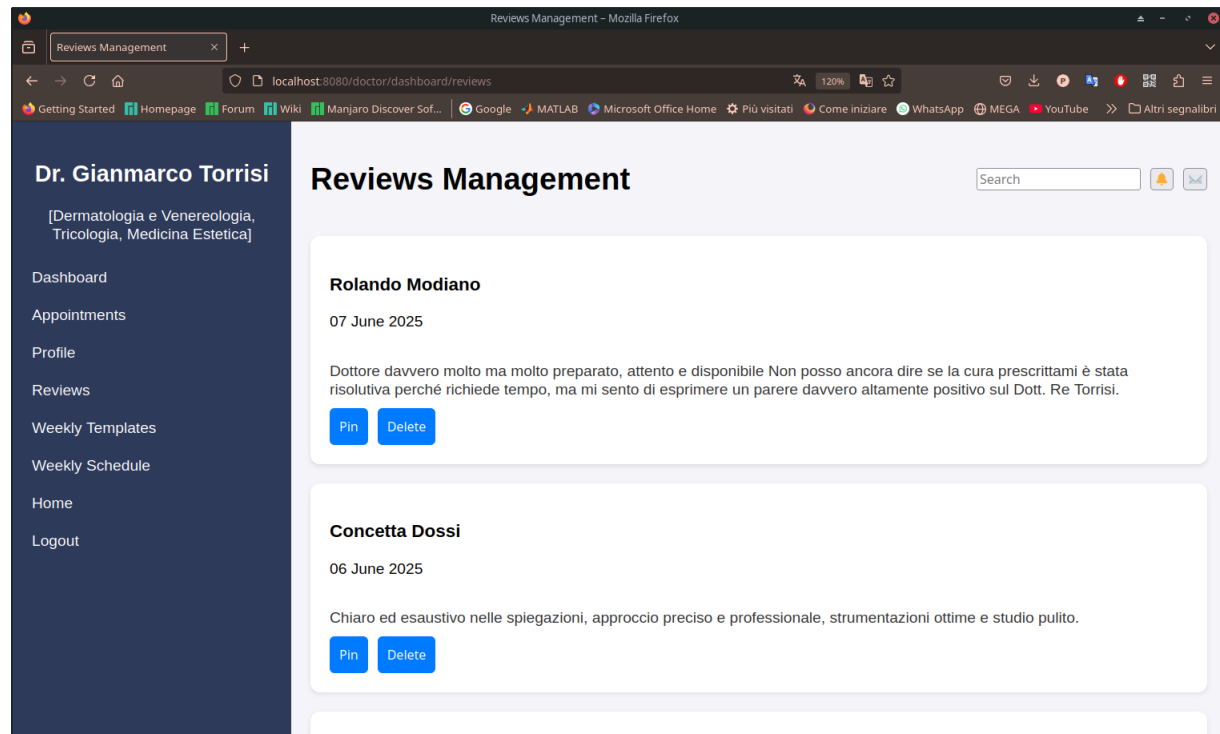


Figure B.22: Patient reviews with options to delete or view details.

B.2.5 Calendar Templates

In this area, doctors can create and manage calendar templates that define their weekly availability. Each template includes:

- A custom name for the template.
- For each day of the week, a list of time slots indicating availability.

Templates can be edited or deleted at any time.

B.2.6 Weekly Schedule

The schedule page allows doctors to view and manage their weekly appointment calendar. Key features include:

- Viewing existing appointments.
- Assigning a predefined calendar template to a specific week.
- Deleting the calendar for a selected week.

The calendar visually distinguishes between booked slots (highlighted in red) and available time slots (in blue), allowing for quick assessment and management.

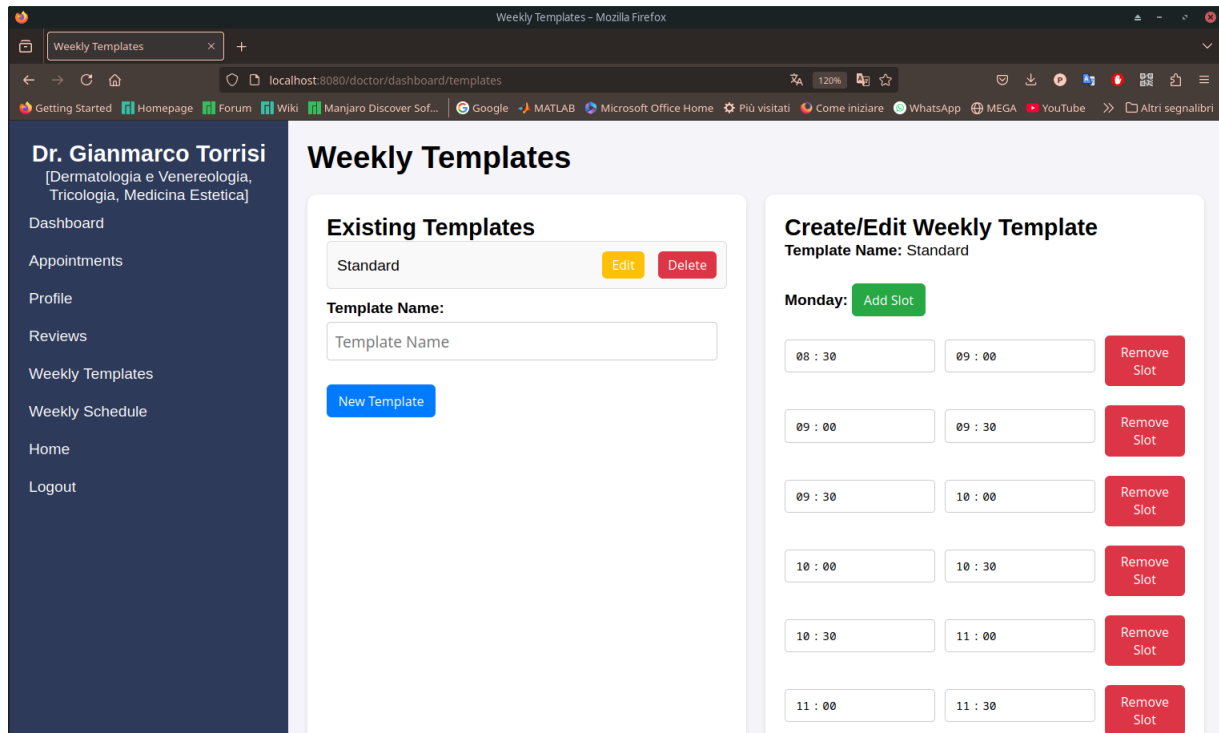


Figure B.23: Calendar templates for managing weekly availability.

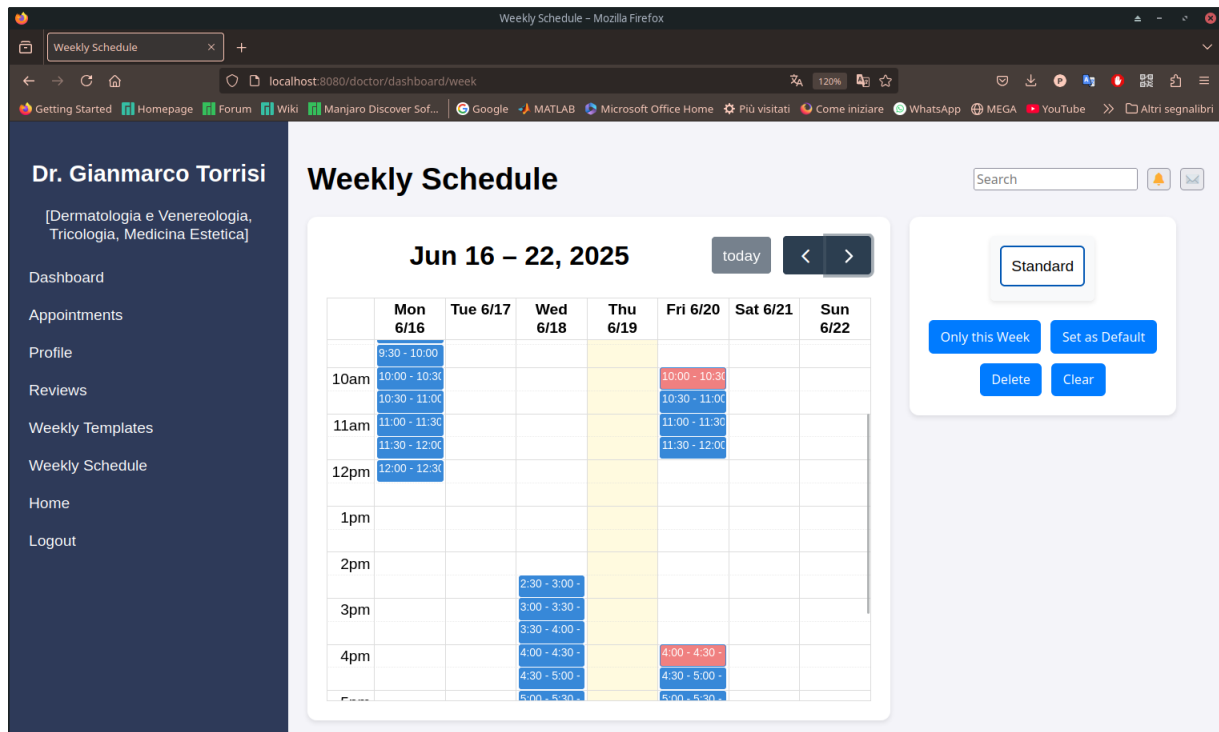


Figure B.24: Weekly schedule with booked and available time slots.